


Multi Tier Filesystem



Thèse de Bachelor (Annexes) présentée par

Monsieur David WITTWER

pour l'obtention du titre Bachelor of Science HES-SO en

**Ingénierie des technologies de l'information avec
orientation en Logiciels et systèmes complexes**

Septembre 2017

Professeur HES responsable TB
Florent GLÜCK

Table des matières

1	Mode d'emploi	3
1.1	Pré requis	3
1.2	Compilation	3
1.2.1	<i>Plugin</i> block	4
1.2.2	<i>Plugin</i> S3	4
1.3	Création d'un espace de stockage	4
1.4	Montage de l'espace de stockage	4
2	Conventions de codage	5
3	Code Source	6
3.1	<i>Config</i>	6
3.1.1	mtfsCreate.cpp	6
3.1.2	mtfsMount.cpp	16
3.2	<i>Plugin System</i>	17
3.2.1	Plugin.h	17
3.2.2	PluginManager	19
3.2.3	BlockDevice	24
3.2.4	S3	43
3.3	<i>Wrapper</i>	56
3.3.1	FuseBase	56
3.3.2	FuseCallback	65
3.3.3	MtfsFuse	75
3.4	<i>Core</i>	79
3.4.1	structs.h	79
3.4.2	Acces.h	87
3.4.3	Mtfs	89
3.4.4	Pool	122
3.4.5	PoolManager	129
3.4.6	Rule	138
3.4.7	TimeRule	143
3.4.8	UserRightRule	146

3.4.9	Volume	149
3.5	<i>Migrator</i>	159
3.5.1	Migrator.h	159
3.5.2	Migrator.cpp	160
3.6	Tests	161
3.7	Migration test	161
3.8	test de vitesse	162
3.9	Plugin tests	164
3.9.1	block device	164
3.9.2	S3	170
4	CMakeLists	177
4.1	CMakeLists général	177
4.2	CMakeLists mtfs	177
4.3	CMakeLists plugin	179
4.3.1	CMakeLists block	179
4.3.2	CMakeLists S3	179
4.4	CMakeLists Tests	179
4.4.1	CMakeLists Performance tests	180
4.4.2	CMakeLists PluginTests	180

Chapitre 1

Mode d'emploi

1.1 Pré requis

Le système n'est pas livré sous forme d'exécutable pour l'instant il faut donc le compiler à partir des sources. Pour ce fait il faut disposer d'un compilateur C++ capable de compiler au minimum la version C++14 comme g++, et il faut aussi disposer de CMake[1] version 3.5 au minimum.

Les libraires suivantes doivent être installées :

- fuse version 3.0.* <https://github.com/libfuse/libfuse> ;
- boost filesystem http://www.boost.org/doc/libs/1_64_0/more/getting_started/unix-variants.html ;
- aws-sdk-cpp-s3 <https://github.com/aws/aws-sdk-cpp> ;

Outre ces librairies il est fortement recommandé de créer un utilisateur pour MTFS. Il est indispensable de créer un home pour MTFS, par défaut le home `/home/mtfs` est utilisé.

1.2 Compilation

La compilation est effectuée à l'aide de CMake, qui permet de compiler chaque exécutable y compris les tests, et de les installer. Les sources du projet sont disponibles à l'adresse <https://github.com/Dawen18/mtfs>. Pour la suite SRC_PATH correspond au dossier contenant les fichiers récupérés sur GitHub. Il est recommandé de créer un dossier pour les *build* de MTFS. Pour la suite de ces explications, BUILD_PATH correspond à ce dossier de *build*, et HOME_DIR au dossier home de MTFS (par défaut `/home/mtfs`). les commandes suivantes sont à exécuter pour installer MTFS :

- `cd <BUILD_PATH>` ;
- `cmake -DMTFS_HOME_DIR=<HOME_DIR> <SRC_PATH>` ;
- Pour compiler l'intégralité utilisez `make` ou alors spécifiez l'exécutable à compiler avec `make <exécutable>`
- `sudo cp mtFS/mtfsMount mtfs/mtfsCreate /usr/local/bin/. : copie des deux exé-`

cutables dans le dossier des binaires.

Il reste maintenant à installer les *plugins*. Avant tout il faut créer le dossier Libs dans le home de MTFS. `mkdir HOME_DIR/Libs`

1.2.1 *Plugin* block

La commande suivante peut être passée si la totalité du projet a été compilé au point précédent.

```
make block
```

Puis copier la librairie dans le dossier home de MTFS

```
cp Plugin/BlockDevide/libblock.so /home/mtfs/Libs
```

1.2.2 *Plugin* S3

La commande suivante peut être passée si la totalité du projet a été compilé en [section 1.2](#).

```
make s3
```

Puis copier la librairie dans le dossier home de MTFS

```
cp Plugin/S3/libs3.so /home/mtfs/Libs
```

1.3 Création d'un espace de stockage

Les commandes suivantes sont celles utilisées pour la configuration utilisée lors des tests. En cas de doute sur une option, consultez l'aide de l'exécutable avec `mtfsCreate -h`.

```
— mtfsCreate -n test
— mtfsCreate -i 0 -d 0 -b 0 -s 512 -r 0 test
— mtfsCreate -add test
— mtfsCreate -p 1 -add -c '{"highLimit": 2, "plName": "block",
  "params": {"devicePath": "/dev/sdb7", "fsType": "ext4"}}' test
— mtfsCreate -p 1 -add -c '{"lowLimit": 1, "plName": "s3",
  "params": {"bucket": "mtfs", "region": "eu-central-1"}}' test
— mtfsCreate -install test
```

1.4 Montage de l'espace de stockage

Pour monter un espace de stockage voici la commande :

```
mtfsMount [options] <mountpoint> <configname>
```

Par exemple pour monter la configuration de test paramétrée ci-dessus la commande suivante est utilisée :

```
mtfsMount -d -o allow_other /tmp/mtfs test
```

Chapitre 2

Conventions de codage

Le code écrit au cours de ce projet suit des conventions afin de conserver un aspect homogène et faciliter ainsi sa lisibilité :

- Les noms des objets sont écrit en *UpperCamelCase* ;
- Les fonctions, types et variables sont écrits en *lowerCamelCase* ;
- Les structs sont post fixée avec "_st" et les typedef avec "_t"
- Les constantes et macro-définition sont écrites en majuscule et leurs mots sont séparés par des *underscore* (exemple : MTFS_HOME_DIR) ;
- Le code est documenté avec Doxygen :
 - Les fonctions publiques sont documentées uniquement dans le *header*. Ces commentaires ne sont pas dupliqués dans le fichier source.
 - A l'inverse les fonctions privées sont documentées dans le fichier source et non dans le *header*.
-

Chapitre 3

Code Source

3.1 *Config*

3.1.1 mtfsCreate.cpp

```
1  /**
2   * \file mtfsCreate.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Public License V3
7   *
8   * This file is part of MTFS.
9
10   MTFS is free software: you can redistribute it and/or modify
11   it under the terms of the GNU General Public License as published by
12   the Free Software Foundation, either version 3 of the License, or
13   (at your option) any later version.
14
15   Foobar is distributed in the hope that it will be useful,
16   but WITHOUT ANY WARRANTY; without even the implied warranty of
17   MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18   GNU General Public License for more details.
19
20   You should have received a copy of the GNU General Public License
21   along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24  // #include <mtfs/structs.h>
25  #include <option/optionparser.h>
26  #include <iostream>
27  #include <utils/Fs.h>
28  #include <rapidjson/document.h>
29  #include <mtfs/Mtfs.h>
30  #include <rapidjson/stringbuffer.h>
31  #include <rapidjson/prettywriter.h>
32  #include <fstream>
33  #include <rapidjson/istreamwrapper.h>
34  #include <mtfs/TimeRule.h>
35  #include <mtfs/PoolManager.h>
36  #include <algorithm>
37  #include <zconf.h>
38  #include <boost/filesystem.hpp>
39  #include <pluginSystem/PluginManager.h>
```

```

40
41 using namespace std;
42 using namespace mtfs;
43 using namespace rapidjson;
44
45 struct Arg : public option::Arg {
46     static void printError(const char *msg1, const option::Option &opt, const
47         char *msg2) {
48         fprintf(stderr, "ERROR: %s", msg1);
49         fwrite(opt.name, (size_t) opt.namelen, 1, stderr);
50         fprintf(stderr, "%s", msg2);
51     }
52
53     static option::ArgStatus Numeric(const option::Option &option, bool msg) {
54         char *endptr = nullptr;
55         if (option.arg != nullptr && (strtol(option.arg, &endptr, 10) != 0)) {};
56         if (endptr != option.arg && *endptr == 0)
57             return option::ARG_OK;
58
59         if (msg) printError("Option '", option, "' requires a numeric argument\n");
60         return option::ARG_ILLEGAL;
61     }
62
63     static option::ArgStatus NonEmpty(const option::Option &option, bool msg) {
64         if (option.arg != nullptr && option.arg[0] != 0)
65             return option::ARG_OK;
66
67         if (msg) printError("Option '", option, "' requires a non-empty argument\n");
68         return option::ARG_ILLEGAL;
69     }
70
71     static option::ArgStatus Migration(const option::Option &option, bool msg) {
72         return option::ARG_OK;
73     }
74 };
75
76 enum optionIndex {
77     UNKNOWN,
78     NEW,
79     ICACHE,
80     DCACHE,
81     BCACHE,
82     BSIZE,
83     REDUNDANCY,
84     MIGRATION,
85     POOL,
86     VOLUME,
87     ADD,
88     DEL,
89     CONFIG,
90     INSTALL,
91     HELP
92 };
93
94 const option::Descriptor usage[] =
95 {
96     {UNKNOWN, 0, "", "", Arg::None, "USAGE:"
97         mtfsCreate [options] <system_name>"},

```



```

95     {NEW,          0, "n",      "new",          Arg::NonEmpty,      " -n, \t—
new      \tCreate a new storage."},
96     {ICACHE,      0, "i",      "inode-cache", Arg::Numeric,      " -i [SIZE],
\t—inode-cache=[SIZE] \tSet inode cache size."},
97     {DCACHE,      0, "d",      "dir-cache",    Arg::Numeric,      " -d [SIZE],
\t—dir-cache=[SIZE] \tSet directory entry cache size."},
98     {BCACHE,      0, "b",      "block-cache",  Arg::Numeric,      " -b [SIZE],
\t—block-cache=[SIZE] \tSet block cache size."},
99     {BSIZE,        0, "s",      "block-size", Arg::Numeric,      " -s [SIZE],
\t—block-size=[SIZE] \tSet block size. Default 4096 octets"},
100    {REDUNDANCY,    0, "r",      "redundancy", Arg::Numeric,      " -r [NB], \
\t—redundancy=[NB] \tSet number of redundancy for each block or inode."},
101    {MIGRATION,     0, "m",      "migration",  Arg::Migration,    " -m [TYPE],
\t—migration=[TYPE] \tSet migration type in system or pool if -p is used"
},
102    {POOL,          0, "p",      "pool",          Arg::Numeric,      " -p [ID], \
\t—pool=[ID] \tSpecify pool."},
103    {VOLUME,        0, "v",      "volume",        Arg::Numeric,      " -v [ID], \
\t—volume=[ID] \tSpecify volume in pool."},
104    {ADD,           0, "",        "add",           Arg::None,         " \t—add
\tAdd a pool or volume."},
105    {DEL,           0, "",        "del",           Arg::None,         " \t—del
\tDelete a pool or volume."},
106    {CONFIG,        0, "c",      "config",        Arg::NonEmpty,     " -c [CONFIG
], \t—config=[CONFIG] \tSet config for plugin or migration (json string or
filename which contains de json config."},
107    {INSTALL,       0, "",        "install",       Arg::None,         " -i, \t—
install \tInstall the config."},
108    {HELP,          0, "h",      "help",          option::Arg::None, " -h, \t—
help \tPrint this help and exit."},
109    {UNKNOWN,       0, "",        "",              option::Arg::None, "\nExamples
:\n"

110                                     " Create system:\n"
111                                     " mtfscCreate -n mtfscRoot\n"
112                                     " Add pool:\n"
113                                     " mtfscCreate -p mtfscRoot\n"
114                                     " Add volume in pool:\n"
115                                     " mtfscCreate -v mtfscRoot 1\n"},
116    {0,              0, nullptr, nullptr,        nullptr,           nullptr}
117 };
118
119 bool configExist(const string &name);
120
121 void loadConfig(superblock_t &sb, const string &name);
122
123 bool writeConfig(superblock_t &superblock, const string &confName);
124
125 uint32_t addVolume(pool_t &pool, volume_t &volume);
126
127 uint32_t addPool(superblock_t &sb, pool_t &pool);
128
129 uint32_t findMissing(std::vector<uint32_t> &x, uint32_t number);
130
131 void installConfig(superblock_t &superblock, const string &name);
132
133
134 int main(int argc, char **argv) {
135     std::cout << MTFSC_CONFIG_DIR << std::endl;
136     argc -= (argc > 0);

```

```

137 argv += (argc > 0);
138 option::Stats stats(usage, argc, argv);
139 option::Option options[stats.options_max], buffer[stats.buffer_max];
140 option::Parser parse(usage, argc, argv, options, buffer);
141
142 if (parse.error())
143     return 1;
144
145 if ((nullptr != options[HELP]) || 0 == argc) {
146     option::printUsage(cerr, usage);
147     return 0;
148 }
149
150 for (option::Option *opt = options[UNKNOWN]; opt; opt = opt->next())
151     cerr << "Unknown option: " << opt->name << "\n";
152 if (options[UNKNOWN] != nullptr)
153     return -1;
154
155 #ifdef DEBUG
156     for (int i = 0; i < parse.nonOptionsCount(); ++i)
157         std::cout << "Non-option #" << i << ": " << parse.nonOption(i) << "\n";
158 #endif
159
160 if (0 != chdir(MIFS_HOME_DIR)) {
161     return errno;
162 }
163
164 superblock_t superblock;
165 memset(&superblock, 0, sizeof(superblock_t));
166 superblock.pools.clear();
167
168 string confName;
169
170 // if option new.
171 if (options[NEW] != nullptr) {
172 #ifdef DEBUG
173     if (configExist(options[NEW].arg)) {
174         cerr << "Config '" << options[NEW].arg << "' already exist" << endl;
175         return -1;
176     }
177 #endif
178
179     confName = options[NEW].arg;
180
181     superblock.iCacheSz = superblock.bCacheSz = superblock.dCacheSz =
182     superblock.blockSz = 4096;
183     superblock.redundancy = 1;
184     superblock.migration = Rule::TIME_MIGRATION;
185 } else {
186     confName = parse.nonOption(0);
187     if (!configExist(confName)) {
188         cerr << "Config not exist" << endl;
189         return -1;
190     }
191
192 #ifdef DEBUG
193     cout << "Load config " << confName << endl;
194 #endif

```

```

195     loadConfig(superblock, confName);
196 }
197
198
199 if (nullptr != options[ICACHE]) {
200     superblock.iCacheSz = static_cast<size_t>(stoi(options[ICACHE].arg));
201 }
202
203 if (nullptr != options[DCACHE]) {
204     superblock.dCacheSz = static_cast<size_t>(stoi(options[DCACHE].arg));
205 }
206
207 if (nullptr != options[BCACHE]) {
208     superblock.bCacheSz = static_cast<size_t>(stoi(options[BCACHE].arg));
209 }
210
211 if (nullptr != options[BFSIZE]) {
212     superblock.blockSz = static_cast<size_t>(stoi(options[BFSIZE].arg));
213 }
214
215 if (nullptr != options[REDUNDANCY]) {
216     superblock.redundancy = static_cast<size_t>(stoi(options[REDUNDANCY].arg));
217 }
218
219 if (nullptr != options[MIGRATION] && nullptr == options[POOL]) {
220     string migration = options[MIGRATION].arg;
221
222     if ("time" == migration) {
223         superblock.migration = Rule::TIME_MIGRATION;
224     } else if ("user" == migration) {
225         superblock.migration = Rule::RIGHT_MIGRATION;
226     } else {
227         cerr << "unknow migration";
228         return -1;
229     }
230 } else if (nullptr != options[MIGRATION] && nullptr != options[POOL]) {
231 }
232
233 if (nullptr != options[ADD] && nullptr == options[POOL]) {
234     // ADD one pool
235     #ifdef DEBUG
236         cout << "add a pool" << endl;
237     #endif
238     pool_t pool;
239     pool.migration = Rule::TIME_MIGRATION;
240     pool.volumes.clear();
241
242     if (nullptr != options[CONFIG]) {
243         string arg = options[CONFIG].arg;
244         Document tmpDoc;
245         #ifdef DEBUG
246             cout << "pool arg: " << arg << endl;
247         #endif
248
249         if (".json" == arg.substr(arg.length() - 5)) {
250             // TODO Parse json file
251         } else
252             tmpDoc.Parse(arg.c_str());
253

```

```

254     if (0 > Rule::rulesAreValid(superblock.migration, tmpDoc)) {
255         cerr << "Invalid config!" << endl;
256         return -1;
257     }
258
259     if (tmpDoc.HasMember(Rule::MIGRATION))
260         pool.migration = tmpDoc[Rule::MIGRATION].GetInt();
261
262     if (tmpDoc.HasMember(TimeRule::TIME_LOW_LIMIT)) {
263         tmpDoc[TimeRule::TIME_LOW_LIMIT].SetUint(tmpDoc[TimeRule::
TIME_LOW_LIMIT].GetUint() * 60);
264     }
265
266     if (tmpDoc.HasMember(TimeRule::TIME_HIGH_LIMIT))
267         tmpDoc[TimeRule::TIME_HIGH_LIMIT].SetUint(tmpDoc[TimeRule::
TIME_HIGH_LIMIT].GetUint() * 60);
268
269     pool.rule = Rule::buildRule(superblock.migration, tmpDoc);
270 } else {
271     // cerr << "config needed!" << endl;
272     // return -1;
273     Document d;
274     d.SetObject();
275     Value v;
276     v.SetInt(0);
277     d.AddMember(StringRef("timeHighLimit"), v, d.GetAllocator());
278     pool.rule = Rule::buildRule(superblock.migration, d);
279 }
280
281 int newPoolId = addPool(superblock, pool);
282 if (0 < newPoolId)
283     cout << "new pool:" << newPoolId << endl;
284
285 } else if (nullptr != options[ADD] && nullptr != options[POOL]) {
286     // ADD one volume.
287     if (superblock.pools.end() == superblock.pools.find(stoul(options[POOL].arg
))) {
288         cerr << "pool '" << options[POOL].arg << "' not found" << endl;
289         return -1;
290     }
291
292     uint32_t poolId = (uint32_t) stoul(options[POOL].arg);
293     volume_t volume;
294     pool_t *pool = &superblock.pools[poolId];
295
296     if (nullptr != options[CONFIG]) {
297         string arg = options[CONFIG].arg;
298         Document tmpDoc;
299 #ifdef DEBUG
300         cout << "pool arg: " << arg << endl;
301 #endif
302
303         if (".json" == arg.substr(arg.length() - 5)) {
304             // TODO Parse json file
305         } else
306             tmpDoc.Parse(arg.c_str());
307
308         if (0 > Rule::rulesAreValid(pool->migration, tmpDoc)) {
309             cerr << "Invalid config!" << endl;

```

```

310     return -1;
311 }
312
313 if (!tmpDoc.HasMember(pluginSystem::Plugin::TYPE)) {
314     cerr << R"(config need a plugin name eg: {"plName":"block"})" << endl;
315     return -1;
316 }
317 volume.pluginName = tmpDoc[pluginSystem::Plugin::TYPE].GetString();
318
319 if (tmpDoc.HasMember(pluginSystem::Plugin::PARAMS)) {
320     for (auto &&param: tmpDoc[pluginSystem::Plugin::PARAMS].GetObject()) {
321         volume.params.insert(make_pair(param.name.GetString(), param.value.
322 GetString()));
323     }
324 }
325
326 if (tmpDoc.HasMember(TimeRule::TIME_LOW_LIMIT)) {
327     tmpDoc[TimeRule::TIME_LOW_LIMIT].SetUint(tmpDoc[TimeRule::
328 TIME_LOW_LIMIT].GetUint() * 60);
329 }
330
331 if (tmpDoc.HasMember(TimeRule::TIME_HIGH_LIMIT))
332     tmpDoc[TimeRule::TIME_HIGH_LIMIT].SetUint(tmpDoc[TimeRule::
333 TIME_HIGH_LIMIT].GetUint() * 60);
334
335 volume.rule = Rule::buildRule(pool->migration, tmpDoc);
336 } else {
337     cerr << "config needed!" << endl;
338     return -1;
339 }
340
341 int newVolumeId = addVolume(*pool, volume);
342 if (0 < newVolumeId)
343     cout << "new volume:" << newVolumeId << endl;
344 }
345
346 // TODO catch -p, -v anc -c options
347
348 if (nullptr != options[INSTALL]) {
349     installConfig(superblock, confName);
350     return EXIT_SUCCESS;
351 }
352
353 writeConfig(superblock, confName);
354
355 return 0;
356 }
357
358 //
359 // ////////////////////////////////////// Definitions
360 //
361
362 bool configExist(const string &name) {

```

```

361     return Fs::fileExists(MTFS_CONFIG_DIR, name + ".json");
362 }
363
364 void loadConfig(superblock_t &sb, const string &name) {
365     string filename = string(MTFS_CONFIG_DIR) + name + ".json";
366     ifstream file(filename);
367     if (!file.is_open()) {
368         return;
369     }
370
371     IStreamWrapper wrapper(file);
372     Document configFile;
373
374     configFile.ParseStream(wrapper);
375
376     Mtfs::jsonToStruct(configFile, sb);
377 }
378
379 bool writeConfig(superblock_t &superblock, const string &confName) {
380     Document d;
381     Document::AllocatorType &allocator = d.GetAllocator();
382     d.SetObject();
383
384     Value v;
385
386     Mtfs::structToJson(superblock, d);
387
388     StringBuffer sb;
389     PrettyWriter<StringBuffer> pw(sb);
390     d.Accept(pw);
391
392     string filename = string(MTFS_CONFIG_DIR) + confName + ".json";
393     ofstream configFile(filename);
394     configFile << sb.GetString() << endl;
395     configFile.close();
396
397     return true;
398 }
399
400 /**
401  * @brief Add volume in pool
402  *
403  * @param pool
404  * @param volume
405  *
406  * @return The volume Id
407  */
408 uint32_t addVolume(pool_t &pool, volume_t &volume) {
409     uint32_t volumeId = 0;
410     vector<uint32_t> ids;
411     ids.clear();
412
413     for (auto &item : pool.volumes) {
414         ids.push_back(item.first);
415     }
416     if (ids.empty())
417         volumeId = 1;
418     else
419         volumeId = findMissing(ids, 0);

```

```

420
421     if (pool.volumes.end() != pool.volumes.find(volumeId))
422         return 0;
423
424     pool.volumes.insert(make_pair(volumeId, volume));
425
426     return volumeId;
427 }
428
429 /**
430  * @brief add a pool in system
431  *
432  * @param sb    The superblock
433  * @param pool  The pool to add
434  *
435  * @return the pool id
436  */
437 uint32_t addPool(superblock_t &sb, pool_t &pool) {
438     uint32_t poolId;
439     vector<uint32_t> ids;
440     ids.clear();
441
442     for (auto &item : sb.pools) {
443         ids.push_back(item.first);
444     }
445     if (ids.empty())
446         poolId = 1;
447     else
448         poolId = findMissing(ids, 0);
449
450     if (sb.pools.end() != sb.pools.find(poolId))
451         return 0;
452
453     sb.pools.insert(make_pair(poolId, pool));
454
455     return poolId;
456 }
457
458 /**
459  * Function find on stackOverflow
460  *
461  * http://stackoverflow.com/questions/28176191/find-first-missing-element-in-a-vector
462  *
463  * @param x vector with content
464  * @param number Number from which to search
465  * @return The missing number.
466  */
467 uint32_t findMissing(std::vector<uint32_t> &x, uint32_t number) {
468     std::sort(x.begin(), x.end());
469     auto pos = std::upper_bound(x.begin(), x.end(), number);
470
471     if (*pos - number > 1)
472         return number + 1;
473
474     std::vector<int> diffs;
475     std::adjacent_difference(pos, x.end(), std::back_inserter(diffs));
476     auto pos2 = std::find_if(diffs.begin() + 1, diffs.end(), [](int x) { return x
477         > 1; });

```

```

477     return *(pos + (pos2 - diffs.begin() - 1)) + 1;
478
479 }
480
481 void installConfig(superblock_t &superblock, const string &name) {
482     // create ident for rootInode.
483     superblock.rootInodes.clear();
484     const int rootRedundancy = max(3, (int) superblock.redundancy);
485     int i = 0;
486     for (auto &&pool: superblock.pools) {
487         for (auto &&volume: pool.second.volumes) {
488             if (rootRedundancy > i)
489                 superblock.rootInodes.emplace_back(0, volume.first, pool.first);
490             else
491                 break;
492             i++;
493         }
494     }
495
496     writeConfig(superblock, name);
497
498     string filename = "superblock.json";
499     boost::filesystem::create_directory(MTFS_INSTALL_DIR + name);
500     string src = string(MTFS_CONFIG_DIR) + name + ".json";
501     string dst = string(MTFS_INSTALL_DIR) + name + "/" + filename;
502     boost::filesystem::copy_file(src, dst, boost::filesystem::copy_option::
        overwrite_if_exists);
503     string rootFilename = string(MTFS_INSTALL_DIR) + name + "/root.json";
504
505     inode_t rootInode;
506     Mtfs::createRootInode(rootInode);
507
508     Document rootJSON(kObjectType);
509     rootInode.toJson(rootJSON);
510
511     // Attach all plugins for write superblock and rootInode
512     // Only 3 firsts for rootInode
513     pluginSystem::PluginManager *manager = pluginSystem::PluginManager::
        getInstance();
514     i = 0;
515     for (auto &&pool: superblock.pools) {
516         for (auto &&volume: pool.second.volumes) {
517             volume.second.params["home"] = MTFS_PLUGIN_HOME;
518             volume.second.params["blockSize"] = to_string(superblock.blockSz);
519
520             pluginSystem::Plugin *plugin;
521
522             plugin = manager->getPlugin(volume.second.pluginName);
523             plugin->attach(volume.second.params);
524
525             if (rootRedundancy > i)
526                 plugin->put(0, &rootInode, INODE, false);
527             plugin->putSuperblock(superblock);
528
529             i++;
530
531             plugin->detach();
532         }
533     }

```



```

534   StringBuffer sb;
535   PrettyWriter<StringBuffer> pw(sb);
536   rootJSON.Accept(pw);
537
538
539   ofstream configFile(rootFilename);
540   configFile << sb.GetString() << endl;
541   configFile.close();
542 }

```

Listing 3.1 – "mtfsCreate.cpp"

3.1.2 mtfsMount.cpp

```

1  /**
2   * \file mtfsMount.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10   MTFS is free software: you can redistribute it and/or modify
11   it under the terms of the GNU General Public License as published by
12   the Free Software Foundation, either version 3 of the License, or
13   (at your option) any later version.
14
15   Foobar is distributed in the hope that it will be useful,
16   but WITHOUT ANY WARRANTY; without even the implied warranty of
17   MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18   GNU General Public License for more details.
19
20   You should have received a copy of the GNU General Public License
21   along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #define FUSE_USE_VERSION 30
25
26 #include <iostream>
27 #include <unistd.h>
28 #include <rapidjson/document.h>
29 #include <rapidjson/istreamwrapper.h>
30 #include <fstream>
31 #include <pluginSystem/PluginManager.h>
32 #include <mtfs/PoolManager.h>
33 #include <mtfs/Mtfs.h>
34 #include <wrapper/MtfsFuse.h>
35 #include <utils/Fs.h>
36 #include <boost/filesystem.hpp>
37 #include <utils/Logger.h>
38
39
40 using namespace std;
41 using namespace rapidjson;
42
43 int main(int argc, char **argv) {
44     if (!Fs::dirExists(MTFS_HOME_DIR)) {

```

```

45     cerr << "Sorry no configured system found." << endl;
46     cerr << "Please configure one or recover with -r device_in_system "
47         "where device_in_system is a device wich was in the configuration." <<
    endl;
48     return -1;
49 }
50
51
52 string sysName = argv[argc - 1];
53 string filename = sysName + ".json";
54 string filepath = string(MTFS_HOME_DIR) + "/" + mtfs::Mtfs::CONFIG_DIR + "/"
    + filename;
55 argc--;
56
57 if (!boost::filesystem::exists(filepath)) {
58     cerr << "File not found" << endl;
59     return -1;
60 }
61
62 chdir(MTFS_HOME_DIR);
63 ifstream file(string(mtfs::Mtfs::CONFIG_DIR) + "/" + filename);
64 if (!file.is_open()) {
65     cerr << "error openning file " << strerror(errno) << endl;
66     return -1;
67 }
68
69 IStreamWrapper wrapper(file);
70 Document d;
71 d.ParseStream(wrapper);
72
73 // validate config file.
74 if (!mtfs::Mtfs::validate(d)) {
75     cerr << "Invalid or corrupted JSON!" << endl;
76     return -1;
77 }
78
79 // build mtfs
80 mtfs::Mtfs::start(d, MTFS_HOME_DIR, sysName);
81
82 // build mtsfFUSE.
83 wrapper::MtfsFuse *mtfsFuse;
84 mtfsFuse = new wrapper::MtfsFuse();
85
86 int ret = mtfsFuse->run(argc, argv);
87
88 mtfs::Mtfs::stop();
89
90 return ret;
91 }

```

Listing 3.2 – "mtfsMount.cpp"

3.2 *Plugin System*

3.2.1 Plugin.h

```

1 /**

```

```

2  * \file Plugin.h
3  * \brief
4  * \author David Wittwer
5  * \version 0.0.1
6  * \copyright GNU Publis License V3
7  *
8  * This file is part of MTFS.
9
10 MTFS is free software: you can redistribute it and/or modify
11 it under the terms of the GNU General Public License as published by
12 the Free Software Foundation, either version 3 of the License, or
13 (at your option) any later version.
14
15 FooBar is distributed in the hope that it will be useful,
16 but WITHOUT ANY WARRANTY; without even the implied warranty of
17 MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18 GNU General Public License for more details.
19
20 You should have received a copy of the GNU General Public License
21 along with FooBar. If not, see <http://www.gnu.org/licenses/>.
22 */
23 #ifndef PLUGINSYSTEM_PLUGIN_H
24 #define PLUGINSYSTEM_PLUGIN_H
25
26 #include <string>
27 #include <vector>
28 #include <list>
29 #include <iostream>
30 #include <cassert>
31
32 #include <mtfs/structs.h>
33 #include <map>
34
35 namespace pluginSystem {
36     typedef struct {
37         std::string name;
38         std::vector<std::string> params;
39     } pluginInfo_t;
40
41     class Plugin {
42     public:
43         static constexpr const char *TYPE = "plName";
44         static constexpr const char *PARAMS = "params";
45
46         enum statusCode {
47             SUCCESS,
48         };
49
50         /**
51          * Get the name of plugin
52          *
53          * @return name of plugin
54          */
55         virtual std::string getName()=0;
56
57         virtual bool attach(std::map<std::string, std::string> params)=0;
58
59         virtual bool detach()=0;
60

```

```

61     virtual int add(uint64_t *id, const mtfs::blockType &type)=0;
62
63     virtual int del(const uint64_t &id, const mtfs::blockType &type)=0;
64
65     virtual int get(const uint64_t &id, void *data, const mtfs::blockType &type
66     , bool metas)=0;
67
68     virtual int put(const uint64_t &id, const void *data, const mtfs::blockType
69     &type, bool metas)=0;
70
71     virtual bool getSuperblock(mtfs::superblock_t &superblock)=0;
72
73     virtual bool putSuperblock(const mtfs::superblock_t &superblock)=0;
74
75 };
76 } // namespace Plugin
77 #endif

```

Listing 3.3 – "pluginSystem/Plugin.h

3.2.2 PluginManager

PluginManager.h

```

1  /**
2   * \file PluginManager.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #ifndef TRAVAIL_BACHELOR_PLUGINMANAGER_H
25 #define TRAVAIL_BACHELOR_PLUGINMANAGER_H
26
27 #include <map>
28 #include <pluginSystem/Plugin.h>
29
30 namespace pluginSystem {
31     typedef struct {
32         pluginSystem::Plugin *(*createObj)();
33

```

```

34     void (*destroyObj)(Plugin *plugin);
35
36     pluginInfo_t (*getInfo)();
37 } plugin_t;
38
39 class PluginManager {
40 public:
41     static const int SUCCESS = 0;
42     static const int PLUGIN_FOUND = 0;
43     static const int PLUGIN_NOT_FOUND = -1;
44
45     /**
46      * @brief Get manager instance
47      *
48      * @return Instance of PluginManager
49      */
50     static PluginManager *getInstance();
51
52     /**
53      * @brief Get plugin
54      *
55      * @param pluginName Name of plugin
56      * @return nullptr if no plugin found
57      */
58     Plugin *getPlugin(std::string pluginName);
59
60     /**
61      * @brief destroy plugin object
62      *
63      * @param pluginName Name of plugin.
64      * @param plugin Tthe object to destroy.
65      */
66     void freePlugin(std::string pluginName, Plugin *plugin);
67
68     /**
69      * @brief Get plugin info.
70      *
71      * @param pluginName Name of plugin.
72      * @param info Struct who contains the info @see pluginInfo_t.
73      * @return SUCCESS or PLUGIN_NOT_FOUND
74      */
75     int getInfo(std::string pluginName, pluginInfo_t &info);
76
77     /**
78      * @brief Get last error code.
79      *
80      * @return The error code.
81      */
82     int getError();
83
84 private:
85     static constexpr const char *PLUGIN_DIR = "Libs";
86
87     std::map<std::string, plugin_t> pluginMap;
88     static PluginManager *instance;
89     int lastError;
90
91     PluginManager();
92

```

```

93     bool pluginExist(std::string name);
94
95     Plugin *loadPlugin(std::string name);
96
97 };
98 }
99
100
101
102 #endif //TRAVAIL_BACHELOR_PLUGINMANAGER_H

```

Listing 3.4 – "pluginSystem/PluginManager.h"

PluginManager.cpp

```

1  /**
2   * \file PluginManager.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10   MTFS is free software: you can redistribute it and/or modify
11   it under the terms of the GNU General Public License as published by
12   the Free Software Foundation, either version 3 of the License, or
13   (at your option) any later version.
14
15   Foobar is distributed in the hope that it will be useful,
16   but WITHOUT ANY WARRANTY; without even the implied warranty of
17   MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18   GNU General Public License for more details.
19
20   You should have received a copy of the GNU General Public License
21   along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22   */
23
24 #include <dlfcn.h>
25 #include <boost/filesystem.hpp>
26 #include <boost/range/iterator_range.hpp>
27 #include <pluginSystem/PluginManager.h>
28 #include <pluginSystem/Plugin.h>
29
30 #include "pluginSystem/PluginManager.h"
31 #include "../..../Plugin/S3/S3.h"
32
33 using namespace std;
34 using namespace boost::filesystem;
35
36 namespace pluginSystem {
37     PluginManager *PluginManager::instance = 0;
38
39     /**
40      *
41      *
42      *
43      *
44      *
45      *
46      *
47      *
48      *
49      *
50      *
51      *
52      *
53      *
54      *
55      *
56      *
57      *
58      *
59      *
60      *
61      *
62      *
63      *
64      *
65      *
66      *
67      *
68      *
69      *
70      *
71      *
72      *
73      *
74      *
75      *
76      *
77      *
78      *
79      *
80      *
81      *
82      *
83      *
84      *
85      *
86      *
87      *
88      *
89      *
90      *
91      *
92      *
93      *
94      *
95      *
96      *
97      *
98      *
99      *
100     */

```

```

42  */
43  PluginManager *PluginManager::getInstance() {
44      if (!instance)
45          instance = new PluginManager();
46
47      return instance;
48  }
49
50  Plugin *PluginManager::getPlugin(string pluginName) {
51      map<string, plugin_t>::iterator it;
52
53      it = pluginMap.find(pluginName);
54      if (it != pluginMap.end()) {
55          this->lastError = PLUGIN_FOUND;
56          return it->second.createObj();
57      }
58
59      if (!pluginExist(pluginName)) {
60          this->lastError = PLUGIN_NOT_FOUND;
61          return nullptr;
62      }
63
64      return loadPlugin(pluginName);
65  }
66
67  void PluginManager::freePlugin(std::string pluginName, Plugin *plugin) {
68      if (this->pluginMap.find(pluginName) == this->pluginMap.end()) {
69          this->lastError = PLUGIN_NOT_FOUND;
70          return;
71      }
72
73      this->pluginMap[pluginName].destroyObj(plugin);
74      this->lastError = SUCCESS;
75  }
76
77  int PluginManager::getError() {
78      return this->lastError;
79  }
80
81  int PluginManager::getInfo(std::string pluginName, pluginInfo_t &info) {
82      if (this->pluginMap.find(pluginName) == this->pluginMap.end()) {
83          this->lastError = PLUGIN_NOT_FOUND;
84          return PLUGIN_NOT_FOUND;
85      }
86
87      this->lastError = SUCCESS;
88      info = this->pluginMap[pluginName].getInfo();
89
90      return SUCCESS;
91  }
92
93  /*
94  /*
95  /*
96  /*
97  /*
98  /*
99  /*
100  /*
101  /*
102  /*
103  /*
104  /*
105  /*
106  /*
107  /*
108  /*
109  /*
110  /*
111  /*
112  /*
113  /*
114  /*
115  /*
116  /*
117  /*
118  /*
119  /*
120  /*
121  /*
122  /*
123  /*
124  /*
125  /*
126  /*
127  /*
128  /*
129  /*
130  /*
131  /*
132  /*
133  /*
134  /*
135  /*
136  /*
137  /*
138  /*
139  /*
140  /*
141  /*
142  /*
143  /*
144  /*
145  /*
146  /*
147  /*
148  /*
149  /*
150  /*
151  /*
152  /*
153  /*
154  /*
155  /*
156  /*
157  /*
158  /*
159  /*
160  /*
161  /*
162  /*
163  /*
164  /*
165  /*
166  /*
167  /*
168  /*
169  /*
170  /*
171  /*
172  /*
173  /*
174  /*
175  /*
176  /*
177  /*
178  /*
179  /*
180  /*
181  /*
182  /*
183  /*
184  /*
185  /*
186  /*
187  /*
188  /*
189  /*
190  /*
191  /*
192  /*
193  /*
194  /*
195  /*
196  /*
197  /*
198  /*
199  /*
200  /*
201  /*
202  /*
203  /*
204  /*
205  /*
206  /*
207  /*
208  /*
209  /*
210  /*
211  /*
212  /*
213  /*
214  /*
215  /*
216  /*
217  /*
218  /*
219  /*
220  /*
221  /*
222  /*
223  /*
224  /*
225  /*
226  /*
227  /*
228  /*
229  /*
230  /*
231  /*
232  /*
233  /*
234  /*
235  /*
236  /*
237  /*
238  /*
239  /*
240  /*
241  /*
242  /*
243  /*
244  /*
245  /*
246  /*
247  /*
248  /*
249  /*
250  /*
251  /*
252  /*
253  /*
254  /*
255  /*
256  /*
257  /*
258  /*
259  /*
260  /*
261  /*
262  /*
263  /*
264  /*
265  /*
266  /*
267  /*
268  /*
269  /*
270  /*
271  /*
272  /*
273  /*
274  /*
275  /*
276  /*
277  /*
278  /*
279  /*
280  /*
281  /*
282  /*
283  /*
284  /*
285  /*
286  /*
287  /*
288  /*
289  /*
290  /*
291  /*
292  /*
293  /*
294  /*
295  /*
296  /*
297  /*
298  /*
299  /*
300  /*
301  /*
302  /*
303  /*
304  /*
305  /*
306  /*
307  /*
308  /*
309  /*
310  /*
311  /*
312  /*
313  /*
314  /*
315  /*
316  /*
317  /*
318  /*
319  /*
320  /*
321  /*
322  /*
323  /*
324  /*
325  /*
326  /*
327  /*
328  /*
329  /*
330  /*
331  /*
332  /*
333  /*
334  /*
335  /*
336  /*
337  /*
338  /*
339  /*
340  /*
341  /*
342  /*
343  /*
344  /*
345  /*
346  /*
347  /*
348  /*
349  /*
350  /*
351  /*
352  /*
353  /*
354  /*
355  /*
356  /*
357  /*
358  /*
359  /*
360  /*
361  /*
362  /*
363  /*
364  /*
365  /*
366  /*
367  /*
368  /*
369  /*
370  /*
371  /*
372  /*
373  /*
374  /*
375  /*
376  /*
377  /*
378  /*
379  /*
380  /*
381  /*
382  /*
383  /*
384  /*
385  /*
386  /*
387  /*
388  /*
389  /*
390  /*
391  /*
392  /*
393  /*
394  /*
395  /*
396  /*
397  /*
398  /*
399  /*
400  /*
401  /*
402  /*
403  /*
404  /*
405  /*
406  /*
407  /*
408  /*
409  /*
410  /*
411  /*
412  /*
413  /*
414  /*
415  /*
416  /*
417  /*
418  /*
419  /*
420  /*
421  /*
422  /*
423  /*
424  /*
425  /*
426  /*
427  /*
428  /*
429  /*
430  /*
431  /*
432  /*
433  /*
434  /*
435  /*
436  /*
437  /*
438  /*
439  /*
440  /*
441  /*
442  /*
443  /*
444  /*
445  /*
446  /*
447  /*
448  /*
449  /*
450  /*
451  /*
452  /*
453  /*
454  /*
455  /*
456  /*
457  /*
458  /*
459  /*
460  /*
461  /*
462  /*
463  /*
464  /*
465  /*
466  /*
467  /*
468  /*
469  /*
470  /*
471  /*
472  /*
473  /*
474  /*
475  /*
476  /*
477  /*
478  /*
479  /*
480  /*
481  /*
482  /*
483  /*
484  /*
485  /*
486  /*
487  /*
488  /*
489  /*
490  /*
491  /*
492  /*
493  /*
494  /*
495  /*
496  /*
497  /*
498  /*
499  /*
500  /*
501  /*
502  /*
503  /*
504  /*
505  /*
506  /*
507  /*
508  /*
509  /*
510  /*
511  /*
512  /*
513  /*
514  /*
515  /*
516  /*
517  /*
518  /*
519  /*
520  /*
521  /*
522  /*
523  /*
524  /*
525  /*
526  /*
527  /*
528  /*
529  /*
530  /*
531  /*
532  /*
533  /*
534  /*
535  /*
536  /*
537  /*
538  /*
539  /*
540  /*
541  /*
542  /*
543  /*
544  /*
545  /*
546  /*
547  /*
548  /*
549  /*
550  /*
551  /*
552  /*
553  /*
554  /*
555  /*
556  /*
557  /*
558  /*
559  /*
560  /*
561  /*
562  /*
563  /*
564  /*
565  /*
566  /*
567  /*
568  /*
569  /*
570  /*
571  /*
572  /*
573  /*
574  /*
575  /*
576  /*
577  /*
578  /*
579  /*
580  /*
581  /*
582  /*
583  /*
584  /*
585  /*
586  /*
587  /*
588  /*
589  /*
590  /*
591  /*
592  /*
593  /*
594  /*
595  /*
596  /*
597  /*
598  /*
599  /*
600  /*
601  /*
602  /*
603  /*
604  /*
605  /*
606  /*
607  /*
608  /*
609  /*
610  /*
611  /*
612  /*
613  /*
614  /*
615  /*
616  /*
617  /*
618  /*
619  /*
620  /*
621  /*
622  /*
623  /*
624  /*
625  /*
626  /*
627  /*
628  /*
629  /*
630  /*
631  /*
632  /*
633  /*
634  /*
635  /*
636  /*
637  /*
638  /*
639  /*
640  /*
641  /*
642  /*
643  /*
644  /*
645  /*
646  /*
647  /*
648  /*
649  /*
650  /*
651  /*
652  /*
653  /*
654  /*
655  /*
656  /*
657  /*
658  /*
659  /*
660  /*
661  /*
662  /*
663  /*
664  /*
665  /*
666  /*
667  /*
668  /*
669  /*
670  /*
671  /*
672  /*
673  /*
674  /*
675  /*
676  /*
677  /*
678  /*
679  /*
680  /*
681  /*
682  /*
683  /*
684  /*
685  /*
686  /*
687  /*
688  /*
689  /*
690  /*
691  /*
692  /*
693  /*
694  /*
695  /*
696  /*
697  /*
698  /*
699  /*
700  /*
701  /*
702  /*
703  /*
704  /*
705  /*
706  /*
707  /*
708  /*
709  /*
710  /*
711  /*
712  /*
713  /*
714  /*
715  /*
716  /*
717  /*
718  /*
719  /*
720  /*
721  /*
722  /*
723  /*
724  /*
725  /*
726  /*
727  /*
728  /*
729  /*
730  /*
731  /*
732  /*
733  /*
734  /*
735  /*
736  /*
737  /*
738  /*
739  /*
740  /*
741  /*
742  /*
743  /*
744  /*
745  /*
746  /*
747  /*
748  /*
749  /*
750  /*
751  /*
752  /*
753  /*
754  /*
755  /*
756  /*
757  /*
758  /*
759  /*
760  /*
761  /*
762  /*
763  /*
764  /*
765  /*
766  /*
767  /*
768  /*
769  /*
770  /*
771  /*
772  /*
773  /*
774  /*
775  /*
776  /*
777  /*
778  /*
779  /*
780  /*
781  /*
782  /*
783  /*
784  /*
785  /*
786  /*
787  /*
788  /*
789  /*
790  /*
791  /*
792  /*
793  /*
794  /*
795  /*
796  /*
797  /*
798  /*
799  /*
800  /*
801  /*
802  /*
803  /*
804  /*
805  /*
806  /*
807  /*
808  /*
809  /*
810  /*
811  /*
812  /*
813  /*
814  /*
815  /*
816  /*
817  /*
818  /*
819  /*
820  /*
821  /*
822  /*
823  /*
824  /*
825  /*
826  /*
827  /*
828  /*
829  /*
830  /*
831  /*
832  /*
833  /*
834  /*
835  /*
836  /*
837  /*
838  /*
839  /*
840  /*
841  /*
842  /*
843  /*
844  /*
845  /*
846  /*
847  /*
848  /*
849  /*
850  /*
851  /*
852  /*
853  /*
854  /*
855  /*
856  /*
857  /*
858  /*
859  /*
860  /*
861  /*
862  /*
863  /*
864  /*
865  /*
866  /*
867  /*
868  /*
869  /*
870  /*
871  /*
872  /*
873  /*
874  /*
875  /*
876  /*
877  /*
878  /*
879  /*
880  /*
881  /*
882  /*
883  /*
884  /*
885  /*
886  /*
887  /*
888  /*
889  /*
890  /*
891  /*
892  /*
893  /*
894  /*
895  /*
896  /*
897  /*
898  /*
899  /*
900  /*
901  /*
902  /*
903  /*
904  /*
905  /*
906  /*
907  /*
908  /*
909  /*
910  /*
911  /*
912  /*
913  /*
914  /*
915  /*
916  /*
917  /*
918  /*
919  /*
920  /*
921  /*
922  /*
923  /*
924  /*
925  /*
926  /*
927  /*
928  /*
929  /*
930  /*
931  /*
932  /*
933  /*
934  /*
935  /*
936  /*
937  /*
938  /*
939  /*
940  /*
941  /*
942  /*
943  /*
944  /*
945  /*
946  /*
947  /*
948  /*
949  /*
950  /*
951  /*
952  /*
953  /*
954  /*
955  /*
956  /*
957  /*
958  /*
959  /*
960  /*
961  /*
962  /*
963  /*
964  /*
965  /*
966  /*
967  /*
968  /*
969  /*
970  /*
971  /*
972  /*
973  /*
974  /*
975  /*
976  /*
977  /*
978  /*
979  /*
980  /*
981  /*
982  /*
983  /*
984  /*
985  /*
986  /*
987  /*
988  /*
989  /*
990  /*
991  /*
992  /*
993  /*
994  /*
995  /*
996  /*
997  /*
998  /*
999  /*
1000  /*

```

```

96 PluginManager::PluginManager() {
97 }
98
99
100 bool PluginManager::pluginExist(std::string name) {
101     string dir("./");
102     dir += PLUGIN_DIR;
103
104     if (is_directory(dir)) {
105         for (auto &entry : boost::make_iterator_range(directory_iterator(dir))) {
106             if (entry.path().filename() == (string("lib") + name + ".so")) {
107                 return true;
108             }
109         }
110     }
111     return false;
112 }
113
114 Plugin *PluginManager::loadPlugin(string name) {
115     plugin_t plugin{};
116     string path = string("./") + PLUGIN_DIR + "/lib" + name + ".so";
117
118     // Open plugin
119     void *library = dlopen(path.c_str(), RTLD_LAZY);
120     if (nullptr == library) {
121         cerr << "Cannot load plugin '" << name << "': " << dlerror() << endl;
122         return nullptr;
123     }
124     dlerror();
125
126     // Load createObj symbol
127     plugin.createObj = (pluginSystem::Plugin *(*)) dlsym(library, "createObj");
128     char *dlsym_error = dlerror();
129     if (nullptr != dlsym_error) {
130         cerr << "Cannot load symbol create: " << dlsym_error << endl;
131         return nullptr;
132     }
133     dlerror();
134
135     plugin.destroyObj = (void (*)(Plugin *)) dlsym(library, "destroyObj");
136     dlsym_error = dlerror();
137     if (dlsym_error) {
138         cerr << "Cannot load symbol destroy: " << dlsym_error << endl;
139         return nullptr;
140     }
141     dlerror();
142
143     plugin.getInfo = (pluginInfo_t (*)(*)) dlsym(library, "getInfo");
144     dlsym_error = dlerror();
145     if (dlsym_error) {
146         cerr << "Cannot load symbol info: " << dlsym_error << endl;
147         return nullptr;
148     }
149     dlerror();
150
151     dlclose(library);
152
153     this->pluginMap.insert(make_pair(plugin.getInfo().name, plugin));

```



```

154
155     return plugin.createObj();
156 }
157 }

```

Listing 3.5 – "pluginSystem/PluginManager.cpp"

3.2.3 BlockDevice

BlockDevice.h

```

1  /**
2   * \file BlockDevice.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24  #ifndef PLUGINSYSTEM_BLOCK_DEVICE_H
25  #define PLUGINSYSTEM_BLOCK_DEVICE_H
26
27  #include <string>
28  #include <vector>
29  #include <list>
30  #include <iostream>
31  #include <cassert>
32
33  #include <pluginSystem/Plugin.h>
34  #include <mtfs/structs.h>
35  #include <map>
36  #include <queue>
37  #include <mutex>
38
39  namespace pluginSystem {
40      class BlockDevice : public Plugin {
41          static constexpr const char *INODES_DIR = "inodes";
42          static constexpr const char *BLOCKS_DIR = "blocks";
43          static constexpr const char *DIR_BLOCKS_DIR = "dirBlocks";
44          static constexpr const char *METAS_DIR = "metas";
45          static constexpr const char *INODE_METAS_DIR = "metas/inodes";
46          static constexpr const char *DIR_BLOCK_METAS_DIR = "metas/dirBlocks";
47          static constexpr const char *BLOCK_METAS_DIR = "metas/blocks";

```

```

48     static const int SUCCESS = 0;
49
50
51 public:
52     static constexpr const char *NAME = "block";
53
54     BlockDevice();
55
56     virtual ~BlockDevice();
57
58     bool attach(std::map<std::string, std::string> params) override;
59
60     bool detach() override;
61
62     int add(uint64_t *id, const mfts::blockType &type) override;
63
64     int del(const uint64_t &id, const mfts::blockType &type) override;
65
66     int get(const uint64_t &id, void *data, const mfts::blockType &type, bool
metas) override;
67
68     int put(const uint64_t &id, const void *data, const mfts::blockType &type,
bool metas) override;
69
70     bool getSuperblock(mfts::superblock_t &superblock) override;
71
72     bool putSuperblock(const mfts::superblock_t &superblock) override;
73
74     std::string getName() override;
75
76 private:
77     int blockSize;
78     std::string mountpoint;
79     std::string devicePath;
80     std::string fsType;
81
82     std::mutex inodeMutex;
83     std::vector<uint64_t> freeInodes;
84     uint64_t nextFreeInode;
85
86     std::mutex dirBlockMutex;
87     std::vector<uint64_t> freeDirBlocks;
88     uint64_t nextFreeDirBlock;
89
90     std::mutex blockMutex;
91     std::vector<uint64_t> freeBlocks;
92     uint64_t nextFreeBlock;
93
94
95     int addInode(uint64_t *inodeId);
96
97     int addDirBlock(uint64_t *id);
98
99     int addBlock(uint64_t *blockId);
100
101     int delInode(const uint64_t &inodeId);
102
103     int getInode(const uint64_t &inodeId, mfts::inode_st &inode);
104

```

```

105     int putInode(const uint64_t &inodeId, const mtfs::inode_st &inode);
106
107     int getInodeMetas(const uint64_t &inodeId, mtfs::blockInfo_t &metas);
108
109     int putInodeMetas(const uint64_t &inodeId, const mtfs::blockInfo_t &metas);
110
111     int delDirBlock(const uint64_t &id);
112
113     int getDirBlock(const uint64_t &id, mtfs::dirBlock_t &block);
114
115     int putDirBlock(const uint64_t &id, const mtfs::dirBlock_t &block);
116
117     int getDirBlockMetas(const uint64_t &id, mtfs::blockInfo_t &metas);
118
119     int putDirBlockMetas(const uint64_t &id, const mtfs::blockInfo_t &metas);
120
121     int delBlock(const uint64_t &blockId);
122
123     int getBlock(const uint64_t &blockId, std::uint8_t *buffer);
124
125     int putBlock(const uint64_t &blockId, const uint8_t *buffer);
126
127     int getBlockMetas(const uint64_t &blockId, mtfs::blockInfo_t &metas);
128
129     int putBlockMetas(const uint64_t &blockId, const mtfs::blockInfo_t &metas);
130
131
132     void initDirHierarchie();
133
134     void initInodes();
135
136     void initDirBlocks();
137
138     void initBlocks();
139
140     void writeMetas();
141
142     void logError(std::string message);
143
144     bool dirExists(std::string path);
145
146     int createFile(std::string path);
147
148     int deleteFile(std::string path);
149
150     int getMetas(const std::string &filename, mtfs::blockInfo_t &infos);
151
152     int putMetas(const std::string &filename, const mtfs::blockInfo_t &infos);
153
154 };
155
156 } // namespace Plugin
157 #endif

```

Listing 3.6 – "pluginSystem/BlockDevice.h"

BlockDevice.cpp

```

1  /**
2   * \file BlockDevice.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #include <iostream>
25 #include <sys/mount.h>
26 #include <map>
27 #include <sys/stat.h>
28 #include <fcntl.h>
29 #include <rapidjson/document.h>
30 #include <rapidjson/stringbuffer.h>
31 #include <rapidjson/prettywriter.h>
32 #include <rapidjson/istreamwrapper.h>
33 #include <fstream>
34 #include <memory>
35 #include <utils/Logger.h>
36
37 #include "BlockDevice.h"
38
39 #define DEBUG
40
41 using namespace std;
42 using namespace rapidjson;
43 using namespace mtfs;
44
45 namespace pluginSystem {
46     BlockDevice::BlockDevice() : nextFreeBlock(0), nextFreeDirBlock(0),
47         nextFreeInode(1) {
48         this->freeBlocks.clear();
49         this->freeDirBlocks.clear();
50         this->freeInodes.clear();
51     }
52
53     BlockDevice::~BlockDevice() {
54         this->freeBlocks.clear();
55         this->freeDirBlocks.clear();
56         this->freeInodes.clear();
57     }
58
59     string BlockDevice::getName() {

```

```

59     return NAME;
60 }
61
62 bool BlockDevice::attach(std::map<string, string> params) {
63     if (params.find("devicePath") == params.end() || params.find("home") ==
64         params.end() ||
65         params.find("fsType") == params.end())
66         return false;
67
68     string home;
69     home = params["home"];
70     if (home.back() != '/')
71         home += '/';
72     string parentDir = home + "BlockDevices/";
73     string bs = params["blockSize"];
74     this->blockSize = stoi(params.at("blockSize"));
75     this->fsType = params["fsType"];
76     this->devicePath = params["devicePath"];
77     this->mountpoint = parentDir + this->devicePath.substr(this->devicePath.
78         find('/', 1) + 1) + '/';
79
80 #ifdef DEBUG
81     Logger::getInstance()->log("BLOCK", "create dir in home", Logger::L_INFO);
82 #endif
83
84     if (!dirExists(parentDir)) {
85         if (mkdir(parentDir.c_str(), 0700) != 0) {
86             logError("mkdir error " + parentDir + " " + strerror(errno));
87             return false;
88         }
89     }
90
91     if (!dirExists(this->mountpoint)) {
92         if (mkdir(this->mountpoint.c_str(), 0700) != 0) {
93             logError("mkdir error " + this->mountpoint + " " + strerror(errno));
94             return false;
95         }
96     }
97
98 #ifndef DEBUG
99     // Mount device
100     mount(this->devicePath.c_str(), this->mountpoint.c_str(), this->fsType.
101         c_str(), 0, nullptr);
102 #endif
103
104     initDirHierarchie();
105
106     initInodes();
107     initDirBlocks();
108     initBlocks();
109
110     return true;
111 }
112
113 bool BlockDevice::detach() {
114     writeMetas();
115
116 #ifndef DEBUG
117     if (umount2(this->mountpoint.c_str(), MNT_DETACH) != 0) {

```

```

115     cerr << "ERROR Failed to umount: " << this->mountpoint << endl;
116     cerr << "reason: " << strerror(errno) << " [" << errno << "]" << endl;
117     return false;
118 }
119 #endif
120
121     return true;
122 }
123
124 int BlockDevice::add(uint64_t *id, const blockType &type) {
125     int ret;
126     switch (type) {
127     case INODE:
128         ret = this->addInode(id);
129         break;
130     case DIR_BLOCK:
131         ret = this->addDirBlock(id);
132         break;
133     case DATA_BLOCK:
134         ret = this->addBlock(id);
135         break;
136     default:
137         ret = ENOSYS;
138         break;
139     }
140     return ret;
141 }
142
143 int BlockDevice::del(const uint64_t &id, const mfts::blockType &type) {
144     int ret;
145
146     switch (type) {
147     case INODE:
148         ret = this->delInode(id);
149         break;
150     case DIR_BLOCK:
151         ret = this->delDirBlock(id);
152         break;
153     case DATA_BLOCK:
154         ret = this->delBlock(id);
155         break;
156     default:
157         ret = ENOSYS;
158         break;
159     }
160
161     return ret;
162 }
163
164 int BlockDevice::get(const uint64_t &id, void *data, const mfts::blockType &
165     type, const bool metas) {
166     int ret;
167
168     switch (type) {
169     case INODE:
170         if (metas)
171             ret = this->getInodeMetas(id, *(blockInfo_t *) data);
172         else
173             ret = this->getInode(id, *(inode_t *) data);

```

```

173     break;
174 case DIR_BLOCK:
175     if (metas)
176         ret = this->getDirBlockMetas(id, *(blockInfo_t *) data);
177     else
178         ret = this->getDirBlock(id, *(dirBlock_t *) data);
179     break;
180 case DATA_BLOCK:
181     if (metas)
182         ret = this->getBlockMetas(id, *(blockInfo_t *) data);
183     else
184         ret = this->getBlock(id, (uint8_t *) data);
185     break;
186 default:
187     ret = ENOSYS;
188     break;
189 }
190
191 return ret;
192 }
193
194 int BlockDevice::put(const uint64_t &id, const void *data, const mtfs::
195     blockType &type, const bool metas) {
196     int ret;
197
198     switch (type) {
199     case INODE:
200         if (metas)
201             ret = this->putInodeMetas(id, *(blockInfo_t *) data);
202         else
203             ret = this->putInode(id, *(inode_t *) data);
204         break;
205     case DIR_BLOCK:
206         if (metas)
207             ret = this->putDirBlockMetas(id, *(blockInfo_t *) data);
208         else
209             ret = this->putDirBlock(id, *(dirBlock_t *) data);
210         break;
211     case DATA_BLOCK:
212         if (metas)
213             ret = this->putBlockMetas(id, *(blockInfo_t *) data);
214         else
215             ret = this->putBlock(id, (uint8_t *) data);
216         break;
217     default:
218         ret = ENOSYS;
219         break;
220     }
221
222     return ret;
223 }
224
225 int BlockDevice::addInode(uint64_t *inodeId) {
226     unique_lock<mutex> lk(this->inodeMutex);
227     if (this->freeInodes.empty()) {
228         *inodeId = this->nextFreeInode;
229         this->nextFreeInode++;
230     } else {
231         *inodeId = this->freeInodes[0];

```

```

231     this->freeInodes.erase(freeInodes.begin());
232 }
233 lk.unlock();
234
235 string filename = this->mountpoint;
236 if ( '/' != filename.back() )
237     filename += '/';
238
239 filename = filename + INODES_DIR + "/" + to_string(*inodeId);
240
241 return createFile(filename);
242 }
243
244 int BlockDevice::delInode(const uint64_t &inodeId) {
245     unique_lock<mutex> lk(this->inodeMutex);
246     if (this->nextFreeInode - 1 == inodeId)
247         this->nextFreeInode--;
248     else {
249         this->freeInodes.push_back(inodeId);
250     }
251     lk.unlock();
252
253     string filename = this->mountpoint;
254     if ( '/' != filename.back() )
255         filename += '/';
256
257     filename = filename + INODES_DIR + "/" + to_string(inodeId);
258
259     return deleteFile(filename);
260 }
261
262 int BlockDevice::getInode(const uint64_t &inodeId, mtf::inode_st &inode) {
263     string filename = this->mountpoint;
264     if ( '/' != filename.back() )
265         filename += '/';
266
267     filename = filename + INODES_DIR + "/" + to_string(inodeId);
268     ifstream file(filename);
269     if (!file.is_open())
270         return errno != 0 ? errno : ENOENT;
271
272     IStreamWrapper wrapper(file);
273
274     Document d;
275     d.ParseStream(wrapper);
276
277     assert(d.IsObject());
278     assert(d.HasMember(IN_MODE));
279     assert(d.HasMember(IN_UID));
280     assert(d.HasMember(IN_GID));
281     assert(d.HasMember(IN_SIZE));
282     assert(d.HasMember(IN_LINKS));
283     assert(d.HasMember(IN_ACCESS));
284     assert(d.HasMember(IN_BLOCKS));
285
286     inode.accesRight = (uint16_t) d[IN_MODE].GetUint();
287     inode.uid = (uint16_t) d[IN_UID].GetUint();
288     inode.gid = (uint16_t) d[IN_GID].GetUint();
289     inode.size = d[IN_SIZE].GetUint64();

```



```

290     inode.linkCount = (uint8_t) d[IN_LINKS].GetUint();
291     inode.atime = d[IN_ACCESS].GetUint64();
292
293     const Value &dataArray = d[IN_BLOCKS];
294     assert(dataArray.IsArray());
295     inode.dataBlocks.clear();
296     for (auto &a : dataArray.GetArray()) {
297         vector<ident_t> redundancy;
298
299         assert(a.IsArray());
300         for (auto &v : a.GetArray()) {
301             ident_t ident;
302
303             assert(v.IsObject());
304             assert(v.HasMember(ID_POOL));
305             assert(v.HasMember(ID_VOLUME));
306             assert(v.HasMember(ID_ID));
307
308             ident.poolId = (uint16_t) v[ID_POOL].GetUint();
309             ident.volumeId = (uint16_t) v[ID_VOLUME].GetUint();
310             ident.id = v[ID_ID].GetUint64();
311
312             redundancy.push_back(ident);
313         }
314
315         inode.dataBlocks.push_back(redundancy);
316     }
317
318     return this->SUCCESS;
319 }
320
321 int BlockDevice::putInode(const uint64_t &inodeId, const inode_st &inode) {
322     Document d;
323     d.SetObject();
324     Document::AllocatorType &alloc = d.GetAllocator();
325
326     inode.toJson(d);
327
328     StringBuffer sb;
329     PrettyWriter<StringBuffer> pw(sb);
330     d.Accept(pw);
331
332     string filename = this->mountpoint;
333     if ('/' != filename.back())
334         filename += '/';
335
336     filename = filename + BlockDevice::INODES_DIR + "/" + to_string(inodeId);
337     ofstream inodeFile(filename);
338     inodeFile << sb.GetString() << endl;
339     inodeFile.close();
340
341     return this->SUCCESS;
342 }
343
344 int BlockDevice::getInodeMetas(const uint64_t &inodeId, mtf::blockInfo_t &
345     metas) {
346     string filename = this->mountpoint;
347     if ('/' != filename.back())
348         filename += '/';

```

```

348     filename = filename + INODE_METAS_DIR + "/" + to_string(inodeId) + ".json";
349
350
351     return getMetas(filename, metas);
352 }
353
354 int BlockDevice::putInodeMetas(const uint64_t &inodeId, const mfts::
    blockInfo_t &metas) {
355     string filename = this->mountpoint;
356     if ('/' != filename.back())
357         filename += '/';
358
359     filename = filename + INODE_METAS_DIR + "/" + to_string(inodeId) + ".json";
360
361     return putMetas(filename, metas);
362 }
363
364 int BlockDevice::addDirBlock(uint64_t *id) {
365     unique_lock<mutex> lk(this->dirBlockMutex);
366     if (this->freeDirBlocks.empty()) {
367         *id = this->nextFreeDirBlock;
368         this->nextFreeDirBlock++;
369     } else {
370         *id = this->freeDirBlocks.front();
371         iter_swap(this->freeDirBlocks.begin(), this->freeDirBlocks.end());
372         this->freeDirBlocks.pop_back();
373     }
374     lk.unlock();
375
376     string filename = this->mountpoint;
377     if ('/' != filename.back())
378         filename += '/';
379
380     filename = filename + DIR_BLOCKS_DIR + "/" + to_string(*id);
381
382     return createFile(filename);
383 }
384
385 int BlockDevice::delDirBlock(const uint64_t &id) {
386     unique_lock<mutex> lk(this->dirBlockMutex);
387     if (id == this->nextFreeDirBlock - 1)
388         this->nextFreeDirBlock--;
389     else
390         this->freeDirBlocks.push_back(id);
391     lk.unlock();
392
393     string filename = this->mountpoint;
394     if ('/' != filename.back())
395         filename += '/';
396
397     filename = filename + DIR_BLOCKS_DIR + "/" + to_string(id);
398
399     return deleteFile(filename);
400 }
401
402 int BlockDevice::getDirBlock(const uint64_t &id, dirBlock_t &block) {
403     string filename = this->mountpoint;
404     if ('/' != filename.back())
405         filename += '/';

```

```

406     filename = filename + DIR_BLOCKS_DIR + "/" + to_string(id);
407
408     ifstream file(filename);
409     if (!file.is_open())
410         return errno != 0 ? errno : ENOENT;
411
412
413     IStreamWrapper wrapper(file);
414
415     Document d;
416     d.ParseStream(wrapper);
417
418     assert(d.IsObject());
419
420     for (auto &&item: d.GetObject()) {
421         assert(item.value.IsArray());
422
423         vector<ident_t> ids;
424
425         for (auto &&ident: item.value.GetArray()) {
426             ident_t i;
427             i.fromJson(ident);
428             ids.push_back(i);
429         }
430         block.entries.insert(make_pair(item.name.GetString(), ids));
431     }
432     return this->SUCCESS;
433 }
434
435 int BlockDevice::putDirBlock(const uint64_t &id, const dirBlock_t &block) {
436     Document d;
437     d.SetObject();
438
439     Document::AllocatorType &allocator = d.GetAllocator();
440
441     for (auto &&item: block.entries) {
442         Value r(kArrayType);
443         for (auto &&ident: item.second) {
444             Value v(kObjectType);
445             ident.toJson(v, allocator);
446             r.PushBack(v, allocator);
447         }
448         d.AddMember(StringRef(item.first.c_str()), r, allocator);
449     }
450
451     StringBuffer sb;
452     PrettyWriter<StringBuffer> pw(sb);
453     d.Accept(pw);
454
455     string filename = this->mountpoint;
456     if ( '/' != filename.back() )
457         filename += '/';
458
459     filename = filename + BlockDevice::DIR_BLOCKS_DIR + "/" + to_string(id);
460     ofstream file(filename);
461     if (!file.is_open())
462         return errno != 0 ? errno : ENOENT;
463     file << sb.GetString() << endl;
464 }

```

```

465     file.close();
466
467     return this->SUCCESS;
468 }
469
470 int BlockDevice::getDirBlockMetas(const uint64_t &id, mfts::blockInfo_t &
471     metas) {
472     string filename = this->mountpoint;
473     if ( '/' != filename.back() )
474         filename += '/';
475
476     filename = filename + DIR_BLOCK_METAS_DIR + "/" + to_string(id) + ".json";
477     return getMetas(filename, metas);
478 }
479
480 int BlockDevice::putDirBlockMetas(const uint64_t &id, const mfts::blockInfo_t
481     &metas) {
482     string filename = this->mountpoint;
483     if ( '/' != filename.back() )
484         filename += '/';
485
486     filename = filename + DIR_BLOCK_METAS_DIR + "/" + to_string(id) + ".json";
487     return putMetas(filename, metas);
488 }
489
490 int BlockDevice::addBlock(uint64_t *blockId) {
491     unique_lock<mutex> lk(this->blockMutex);
492     if (this->freeBlocks.empty()) {
493         *blockId = this->nextFreeBlock;
494         this->nextFreeBlock++;
495     } else {
496         *blockId = this->freeBlocks[this->freeBlocks.size() - 1];
497         this->freeBlocks.pop_back();
498     }
499     lk.unlock();
500
501     string filename = this->mountpoint;
502     if ( '/' != filename.back() )
503         filename += '/';
504
505     filename = filename + BLOCKS_DIR + "/" + to_string(*blockId);
506
507     return createFile(filename);
508 }
509
510 int BlockDevice::delBlock(const uint64_t &blockId) {
511     unique_lock<mutex> lk(this->blockMutex);
512     if (this->nextFreeBlock - 1 == blockId)
513         nextFreeBlock--;
514     else
515         this->freeBlocks.push_back(blockId);
516     lk.unlock();
517
518     string filename = this->mountpoint;
519     if ( '/' != filename.back() )
520         filename += '/';
521
522     filename = filename + BLOCKS_DIR + "/" + to_string(blockId);

```

```

522     return deleteFile(filename);
523 }
524
525 int BlockDevice::getBlock(const uint64_t &blockId, std::uint8_t *buffer) {
526     string filename = this->mountpoint;
527     if ('/' != filename.back())
528         filename += '/';
529
530     filename = filename + BLOCKS_DIR + "/" + to_string(blockId);
531
532     ifstream file(filename);
533     if (!file.is_open())
534         return errno != 0 ? errno : ENOENT;
535
536     file.read((char *) buffer, this->blockSize);
537     file.close();
538     return this->SUCCESS;
539 }
540
541 int BlockDevice::putBlock(const uint64_t &blockId, const uint8_t *buffer) {
542     string filename = this->mountpoint;
543     if ('/' != filename.back())
544         filename += '/';
545
546     filename = filename + BLOCKS_DIR + "/" + to_string(blockId);
547     ofstream file(filename);
548     if (!file.is_open())
549         return errno != 0 ? errno : ENOENT;
550
551     file.write((const char *) buffer, this->blockSize);
552     file.close();
553     return this->SUCCESS;
554 }
555
556 int BlockDevice::getBlockMetas(const uint64_t &blockId, mtf::blockInfo_t &
557     metas) {
558     string filename = this->mountpoint;
559     if ('/' != filename.back())
560         filename += '/';
561
562     filename = filename + BLOCK_METAS_DIR + "/" + to_string(blockId) + ".json";
563
564     return getMetas(filename, metas);
565 }
566
567 int BlockDevice::putBlockMetas(const uint64_t &blockId, const blockInfo_t &
568     metas) {
569     string filename = this->mountpoint;
570     if ('/' != filename.back())
571         filename += '/';
572
573     filename = filename + BLOCK_METAS_DIR + "/" + to_string(blockId) + ".json";
574
575     return putMetas(filename, metas);
576 }
577
578 bool BlockDevice::getSuperblock(mtf::superblock_t &superblock) {
579     return false;
580 }

```

```

579
580 bool BlockDevice::putSuperblock(const superblock_t &superblock) {
581     string filename = this->mountpoint;
582     if ( '/' != filename.back() )
583         filename += '/';
584
585     filename = filename + METAS_DIR + "/superblock";
586     ofstream sbFile(filename, ios::binary);
587     sbFile.write((const char *) &superblock, sizeof(superblock));
588     sbFile.close();
589
590     return true;
591 }
592
593 ////////////////Private method////////////////////
594
595 void BlockDevice::initDirHierarchie() {
596     if (!dirExists(this->mountpoint + BlockDevice::INODES_DIR))
597         mkdir((this->mountpoint + BlockDevice::INODES_DIR).c_str(), 0700);
598     if (!dirExists(this->mountpoint + BlockDevice::DIR_BLOCKS_DIR))
599         mkdir((this->mountpoint + BlockDevice::DIR_BLOCKS_DIR).c_str(), 0700);
600     if (!dirExists(this->mountpoint + BlockDevice::BLOCKS_DIR))
601         mkdir((this->mountpoint + BlockDevice::BLOCKS_DIR).c_str(), 0700);
602     if (!dirExists(this->mountpoint + BlockDevice::METAS_DIR))
603         mkdir((this->mountpoint + BlockDevice::METAS_DIR).c_str(), 0700);
604     if (!dirExists(this->mountpoint + BlockDevice::INODE_METAS_DIR))
605         mkdir((this->mountpoint + BlockDevice::INODE_METAS_DIR).c_str(), 0700);
606     if (!dirExists(this->mountpoint + BlockDevice::DIR_BLOCK_METAS_DIR))
607         mkdir((this->mountpoint + BlockDevice::DIR_BLOCK_METAS_DIR).c_str(),
608             0700);
609     if (!dirExists(this->mountpoint + BlockDevice::BLOCK_METAS_DIR))
610         mkdir((this->mountpoint + BlockDevice::BLOCK_METAS_DIR).c_str(), 0700);
611 }
612
613 void BlockDevice::initInodes() {
614     string filename = this->mountpoint;
615     if ( '/' != filename.back() )
616         filename += '/';
617
618     filename = filename + METAS_DIR + "/inodes.json";
619
620     ifstream inodeFile(filename);
621     if (!inodeFile.is_open())
622         return;
623
624     IStreamWrapper isw(inodeFile);
625
626     Document d;
627     d.ParseStream(isw);
628
629     assert(d.IsObject());
630     assert(d.HasMember("nextFreeInode"));
631     assert(d.HasMember("freeInodes"));
632     this->nextFreeInode = d["nextFreeInode"].GetUint64();
633     const Value &inodeArray = d["freeInodes"];
634
635     this->freeInodes.clear();
636     assert(inodeArray.IsArray());
637     for (SizeType i = 0; i < inodeArray.Size(); i++) {

```

```

637     this->freeInodes.push_back(inodeArray[i].GetUint64());
638 }
639 }
640
641 void BlockDevice::initDirBlocks() {
642     string filename = this->mountpoint;
643     if ('/' != filename.back())
644         filename += '/';
645
646     filename = filename + METAS_DIR + "/dirBlocks.json";
647
648     ifstream inodeFile(filename);
649     if (!inodeFile.is_open())
650         return;
651
652     IStreamWrapper isw(inodeFile);
653
654     Document d;
655     d.ParseStream(isw);
656
657     assert(d.IsObject());
658     assert(d.HasMember("nextFreeDirBlock"));
659     assert(d.HasMember("freeDirBlocks"));
660     this->nextFreeDirBlock = d["nextFreeDirBlock"].GetUint64();
661     const Value &array = d["freeDirBlocks"];
662
663     this->freeDirBlocks.clear();
664     assert(array.IsArray());
665     for (SizeType i = 0; i < array.Size(); i++) {
666         this->freeDirBlocks.push_back(array[i].GetUint64());
667     }
668 }
669
670 void BlockDevice::initBlocks() {
671     string filename = this->mountpoint;
672     if ('/' != filename.back())
673         filename += '/';
674
675     filename = filename + METAS_DIR + "/blocks.json";
676
677     ifstream inodeFile(filename);
678     if (!inodeFile.is_open())
679         return;
680
681     IStreamWrapper isw(inodeFile);
682
683     Document d;
684     d.ParseStream(isw);
685
686     assert(d.IsObject());
687     assert(d.HasMember("nextFreeBlock"));
688     assert(d.HasMember("freeBlocks"));
689     this->nextFreeBlock = d["nextFreeBlock"].GetUint64();
690     const Value &inodeArray = d["freeBlocks"];
691
692     this->freeBlocks.clear();
693     assert(inodeArray.IsArray());
694     for (SizeType i = 0; i < inodeArray.Size(); i++) {
695         this->freeBlocks.push_back(inodeArray[i].GetUint64());

```

```

696     }
697 }
698
699 void BlockDevice::writeMetas() {
700     {
701         // writeInodeMeta
702         Document d;
703         d.SetObject();
704
705         Document::AllocatorType &allocator = d.GetAllocator();
706
707         Value value(kObjectType);
708         value.SetUint64(this->nextFreeInode);
709         d.AddMember(StringRef("nextFreeInode"), value, allocator);
710
711         Value inodeList(kArrayType);
712         Value inode(kObjectType);
713
714         for (unsigned long &freeInode : this->freeInodes) {
715             inode.SetUint64(freeInode);
716             inodeList.PushBack(Value(freeInode), allocator);
717         }
718
719         d.AddMember(StringRef("freeInodes"), inodeList, allocator);
720
721         StringBuffer strBuff;
722         PrettyWriter<StringBuffer> writer(strBuff);
723         d.Accept(writer);
724
725         string filename = this->mountpoint;
726         if ('/' != filename.back())
727             filename += '/';
728
729         filename = filename + METAS_DIR + "/inodes.json";
730
731         ofstream inodeFile;
732         inodeFile.open(filename);
733         inodeFile << strBuff.GetString() << endl;
734         inodeFile.close();
735     }
736
737     {
738         // writeDirBlocksMeta
739         Document d;
740         d.SetObject();
741
742         Document::AllocatorType &allocator = d.GetAllocator();
743
744         Value freeInode(kObjectType);
745         freeInode.SetUint64(this->nextFreeDirBlock);
746         d.AddMember(StringRef("nextFreeDirBlock"), freeInode, allocator);
747
748         Value inodeList(kArrayType);
749         Value inode(kObjectType);
750
751         for (unsigned long &freeDirBlock : this->freeDirBlocks) {
752             inode.SetUint64(freeDirBlock);
753             inodeList.PushBack(Value(freeDirBlock), allocator);
754         }

```



```

755     d.AddMember(StringRef("freeDirBlocks"), inodeList, allocator);
756
757     StringBuffer strBuff;
758     PrettyWriter<StringBuffer> writer(strBuff);
759     d.Accept(writer);
760
761     string filename = this->mountpoint;
762     if ('/' != filename.back())
763         filename += '/';
764
765     filename = filename + METAS_DIR + "/dirBlocks.json";
766
767     ofstream inodeFile;
768     inodeFile.open(filename);
769     inodeFile << strBuff.GetString() << endl;
770     inodeFile.close();
771 }
772
773 {
774     // Write block metas
775     Document db;
776     db.SetObject();
777
778     Document::AllocatorType &blAllocator = db.GetAllocator();
779
780     Value freeBlock(kObjectType);
781     freeBlock.SetUint64(this->nextFreeBlock);
782     db.AddMember(StringRef("nextFreeBlock"), freeBlock, blAllocator);
783
784     Value blockList(kArrayType);
785
786     for (auto b : this->freeBlocks) {
787         blockList.PushBack(Value(b), blAllocator);
788     }
789
790     db.AddMember(StringRef("freeBlocks"), blockList, blAllocator);
791
792     StringBuffer bStrBuff;
793     PrettyWriter<StringBuffer> bWriter(bStrBuff);
794     db.Accept(bWriter);
795
796     string filename = this->mountpoint;
797     if ('/' != filename.back())
798         filename += '/';
799
800     filename = filename + METAS_DIR + "/blocks.json";
801
802     ofstream blockFile;
803     blockFile.open(filename);
804     blockFile << bStrBuff.GetString() << endl;
805     blockFile.close();
806 }
807 }
808
809 void BlockDevice::logError(string message) {
810     cerr << message << endl;
811 }
812
813

```

```

814 bool BlockDevice::dirExists(string path) {
815     struct stat info{};
816
817     if (stat(path.c_str(), &info) != 0)
818         return false;
819     return (info.st_mode & S_IFDIR) != 0;
820 }
821
822 int BlockDevice::createFile(string path) {
823     if (creat(path.c_str(), 0600) < 0) {
824         string message = "create " + path + " error " + string(strerror(errno)) +
825             " [" + to_string(errno) + "]";
826         logError(message);
827         return errno;
828     }
829     return this->SUCCESS;
830 }
831
832 int BlockDevice::deleteFile(std::string path) {
833     if (remove(path.c_str()) != 0) {
834         string message = "ERROR: deleting file " + path;
835         logError(message);
836         return errno;
837     }
838     return SUCCESS;
839 }
840
841 int BlockDevice::getMetas(const std::string &filename, mfts::blockInfo_t &
842     infos) {
843     ifstream file(filename);
844     if (!file.is_open())
845         return 0 != errno ? errno : EAGAIN;
846
847     IStreamWrapper wrapper(file);
848
849     Document d;
850     d.ParseStream(wrapper);
851
852     assert(d.HasMember(BI_REFF));
853     assert(d[BI_REFF].IsArray());
854     assert(d.HasMember(BI_ACCESS));
855
856     for (auto &id : d[BI_REFF].GetArray()) {
857         assert(id.IsObject());
858         assert(id.HasMember(ID_POOL));
859         assert(id.HasMember(ID_VOLUME));
860         assert(id.HasMember(ID_ID));
861
862         uint32_t poolId = id[ID_POOL].GetUint();
863         uint32_t volumeId = id[ID_VOLUME].GetUint();
864         uint64_t i = id[ID_ID].GetUint64();
865
866         infos.referenceId.emplace_back(i, volumeId, poolId);
867     }
868
869     infos.lastAccess = d[BI_ACCESS].GetUint64();
870
871     return 0;
872 }

```

```

871
872 int BlockDevice::putMetas(const std::string &filename, const mtfs::
    blockInfo_t &infos) {
873     Document d;
874     d.SetObject();
875     Document::AllocatorType &allocator = d.GetAllocator();
876
877     Value v;
878
879     Value a(kArrayType);
880     for (auto &&id : infos.referenceId) {
881         v.SetObject();
882
883         v.AddMember(StringRef(ID_POOL), Value(id.poolId), allocator);
884         v.AddMember(StringRef(ID_VOLUME), Value(id.volumeId), allocator);
885         v.AddMember(StringRef(ID_ID), Value(id.id), allocator);
886
887         a.PushBack(v, allocator);
888     }
889     d.AddMember(StringRef(BI_REFF), a, allocator);
890
891     v.SetUint64(infos.lastAccess);
892     d.AddMember(StringRef(BI_ACCESS), v, allocator);
893
894     StringBuffer bStrBuff;
895     PrettyWriter<StringBuffer> bWriter(bStrBuff);
896     d.Accept(bWriter);
897
898     ofstream blockFile;
899     blockFile.open(filename);
900     if (!blockFile.is_open())
901         return -1;
902
903     blockFile << bStrBuff.GetString() << endl;
904     blockFile.close();
905
906     return 0;
907 }
908
909 } // namespace Plugin
910
911 extern "C" pluginSystem::Plugin *createObj() {
912     return new pluginSystem::BlockDevice();
913 }
914
915 extern "C" void destroyObj(pluginSystem::Plugin *plugin) {
916     delete plugin;
917 }
918
919 extern "C" pluginSystem::pluginInfo_t getInfo() {
920     vector<string> params;
921     params.emplace_back("devicePath");
922     params.emplace_back("fsType");
923
924     pluginSystem::pluginInfo_t info = {
925         .name = pluginSystem::BlockDevice::NAME,
926         .params = params,
927     };
928 }

```

```

929
930     return info;
931 }

```

Listing 3.7 – "pluginSystem/BlockDevice.cpp"

3.2.4 S3

S3.h

```

1  /**
2   * \file S3.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #ifndef MTFS_S3_H
25 #define MTFS_S3_H
26
27 #include <pluginSystem/Plugin.h>
28 #include <aws/core/Aws.h>
29 #include <aws/s3/S3Client.h>
30
31 namespace pluginSystem {
32     class S3 : public Plugin {
33     public:
34
35         static constexpr const char *NAME = "s3";
36
37         S3();
38
39         std::string getName() override;
40
41         bool attach(std::map<std::string, std::string> params) override;
42
43         bool detach() override;
44
45         int add(uint64_t *id, const mtfs::blockType &type) override;
46
47         int del(const uint64_t &id, const mtfs::blockType &type) override;
48

```

```

49     int get(const uint64_t &id, void *data, const mtfs::blockType &type, bool
metas) override;
50
51     int put(const uint64_t &id, const void *data, const mtfs::blockType &type,
bool metas) override;
52
53     bool getSuperblock(mtfs::superblock_t &superblock) override;
54
55     bool putSuperblock(const mtfs::superblock_t &superblock) override;
56
57 private:
58     static constexpr const char *INODE_PREFIX = "inodes/";
59     static constexpr const char *DIRECTORY_PREFIX = "directories/";
60     static constexpr const char *DATA_PREFIX = "datas/";
61
62     Aws::S3::S3Client *s3Client;
63     Aws::SDKOptions *options;
64     std::string home;
65     size_t blocksize;
66     Aws::String bucket;
67
68     std::mutex inodeMutex;
69     std::vector<uint64_t> freeInodes;
70     uint64_t nextFreeInode;
71
72     std::mutex dirBlockMutex;
73     std::vector<uint64_t> freeDirectory;
74     uint64_t nextFreeDirectory;
75
76     std::mutex blockMutex;
77     std::vector<uint64_t> freeDatas;
78     uint64_t nextFreeData;
79
80     int writeTmpFile(const std::string &filename, const void *data, const mtfs
::blockType &type);
81
82     int writeInode(const std::string &filename, const mtfs::inode_t *inode);
83
84     int writeDirBlock(const std::string &filename, const mtfs::dirBlock_t *
dirBlock);
85
86     int writeDataBlock(const std::string &filename, const uint8_t *datas);
87
88     int writeSuperblock(const std::string &filename, const mtfs::superblock_t *
superblock);
89
90     int dlObj(const std::string &filename, const std::string &key);
91
92     int upObj(const std::string &filename, const std::string &key);
93
94     bool dirExists(std::string path);
95
96     void initIds();
97
98     void writeMetas();
99
100    void delObj(const std::string &key);
101 };
102 }

```

```

103
104 #endif //MTFS_S3_H

```

Listing 3.8 – "pluginSystem/S3.h"

S3.cpp

```

1  /**
2   * \file S3.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #include <aws/s3/model/PutObjectRequest.h>
25 #include <aws/s3/model/GetObjectRequest.h>
26 #include <aws/s3/model/DeleteObjectRequest.h>
27 #include <iostream>
28 #include <fstream>
29 #include <rapidjson/stringbuffer.h>
30 #include <rapidjson/prettywriter.h>
31 #include <sys/stat.h>
32 #include <utils/Logger.h>
33 #include <zconf.h>
34 #include <rapidjson/istreamwrapper.h>
35 #include <mtfs/structs.h>
36
37 #include "S3.h"
38
39 using namespace std;
40 using namespace Aws::S3;
41 using namespace rapidjson;
42
43 namespace pluginSystem {
44
45     S3::S3() : nextFreeInode(1), nextFreeDirectory(0), nextFreeData(0) {
46         this->freeDatas.clear();
47         this->freeDirectory.clear();
48         this->freeInodes.clear();
49     }
50
51     std::string S3::getName() {
52         return S3::NAME;

```

```

53 }
54
55 bool S3::attach(std::map<std::string, std::string> params) {
56     if (params.end() == params.find("home") || params.end() == params.find("
    blockSize") ||
57         params.end() == params.find("region") || params.end() == params.find("
    bucket"))
58         return false;
59
60     this->blocksize = static_cast<size_t>(stoi(params["blockSize"]));
61     this->bucket = params["bucket"].c_str();
62     this->home = params["home"] + this->bucket.c_str() + "/";
63
64     this->options = new Aws::SDKOptions();
65     Aws::InitAPI(*this->options);
66
67     Aws::Client::ClientConfiguration configuration;
68     configuration.region = params["region"].c_str();
69     this->s3Client = new S3Client(configuration);
70
71     if (!dirExists(this->home)) {
72         if (mkdir(this->home.c_str(), 0700) != 0) {
73             Logger::getInstance()->log("S3", "mkdir error " + this->home + " " +
    strerror(errno), Logger::L_ERROR);
74             return false;
75         }
76     }
77
78     this->initIds();
79
80     return true;
81 }
82
83 bool S3::detach() {
84     this->writeMetas();
85
86     Aws::ShutdownAPI(*this->options);
87
88     return true;
89 }
90
91 int S3::add(uint64_t *id, const mtfs::blockType &type) {
92
93     mutex *mu = nullptr;
94     vector<uint64_t> *ve = nullptr;
95     uint64_t *nextId;
96
97     switch (type) {
98     case mtfs::INODE:
99         mu = &this->inodeMutex;
100         ve = &this->freeInodes;
101         nextId = &this->nextFreeInode;
102         break;
103     case mtfs::DIR_BLOCK:
104         mu = &this->dirBlockMutex;
105         ve = &this->freeDirectory;
106         nextId = &this->nextFreeDirectory;
107         break;
108     case mtfs::DATA_BLOCK:

```

```

109     mu = &this->blockMutex;
110     ve = &this->freeDatas;
111     nextId = &this->nextFreeData;
112     break;
113 default:
114     return ENOSYS;
115 }
116
117 unique_lock<mutex> lk(*mu);
118 if (ve->empty()) {
119     *id = *nextId;
120     (*nextId)++;
121 } else {
122     *id = (*ve)[0];
123     ve->erase(ve->begin());
124 }
125 lk.unlock();
126
127 return 0;
128 }
129
130 int S3::del(const uint64_t &id, const mtf::blockType &type) {
131     mutex *mu = nullptr;
132     vector<uint64_t> *ve = nullptr;
133     uint64_t *nextId;
134
135     string key;
136
137     switch (type) {
138     case mtf::INODE:
139         mu = &this->inodeMutex;
140         ve = &this->freeInodes;
141         nextId = &this->nextFreeInode;
142         key = INODE_PREFIX;
143         break;
144     case mtf::DIR_BLOCK:
145         mu = &this->dirBlockMutex;
146         ve = &this->freeDirectory;
147         nextId = &this->nextFreeDirectory;
148         key = DIRECTORY_PREFIX;
149         break;
150     case mtf::DATA_BLOCK:
151         mu = &this->blockMutex;
152         ve = &this->freeDatas;
153         nextId = &this->nextFreeData;
154         key = DATA_PREFIX;
155         break;
156     default:
157         return ENOSYS;
158     }
159
160     unique_lock<mutex> lk(*mu);
161     if (id == *nextId - 1) {
162         (*nextId)--;
163     } else {
164         ve->push_back(id);
165     }
166     lk.unlock();
167

```



```

168     key += to_string(id);
169     this->delObj(key);
170
171     return 0;
172 }
173
174 int S3::get(const uint64_t &id, void *data, const mtfs::blockType &type, bool
    metas) {
175     string filename = this->home, key;
176     switch (type) {
177         case mtfs::INODE:
178             filename += "in";
179             key += INODE_PREFIX;
180             break;
181         case mtfs::DIR_BLOCK:
182             filename += "di";
183             key += DIRECTORY_PREFIX;
184             break;
185         case mtfs::DATA_BLOCK:
186             filename += "da";
187             key += DATA_PREFIX;
188             break;
189         default:
190             return ENOSYS;
191     }
192
193     if (metas) {
194         filename += "me";
195         key = "metas/" + key;
196     }
197
198     key += to_string(id);
199
200     if (0 != this->dObj(filename, key))
201         return 1;
202
203     if (mtfs::DATA_BLOCK == type) {
204         if (metas) {
205             ifstream file(filename);
206             if (!file.is_open())
207                 return errno != 0 ? errno : ENOENT;
208
209             IStreamWrapper wrapper(file);
210
211             Document d;
212             d.ParseStream(wrapper);
213
214             ((mtfs::blockInfo_t *) data)->fromJson(d);
215         } else {
216             ifstream file(filename);
217             if (!file.is_open())
218                 return errno != 0 ? errno : ENOENT;
219
220             file.read((char *) data, this->blocksize);
221             file.close();
222         }
223     } else {
224         ifstream file(filename);
225         if (!file.is_open())

```

```

226         return errno != 0 ? errno : ENOENT;
227
228     IStreamWrapper wrapper(file);
229
230     Document d;
231     d.ParseStream(wrapper);
232
233     switch (type) {
234     case mtfs::INODE:
235         if (metas)
236             ((mtfs::blockInfo_t *) data)->fromJson(d);
237         else
238             ((mtfs::inode_t *) data)->fromJson(d);
239         break;
240     case mtfs::DIR_BLOCK:
241         if (metas)
242             ((mtfs::blockInfo_t *) data)->fromJson(d);
243         else
244             ((mtfs::dirBlock_t *) data)->fromJson(d);
245         break;
246     case mtfs::SUPERBLOCK:
247         break;
248     default:
249         return ENOSYS;
250     }
251
252     file.close();
253 }
254
255 unlink(filename.c_str());
256 return 0;
257 }
258
259 int S3::put(const uint64_t &id, const void *data, const mtfs::blockType &type
260 , bool metas) {
261     string filename = this->home, key;
262     switch (type) {
263     case mtfs::INODE:
264         filename += "in";
265         key += INODE_PREFIX;
266         break;
267     case mtfs::DIR_BLOCK:
268         filename += "di";
269         key += DIRECTORY_PREFIX;
270         break;
271     case mtfs::DATA_BLOCK:
272         filename += "da";
273         key += DATA_PREFIX;
274         break;
275     default:
276         return ENOSYS;
277     }
278
279     if (metas) {
280         filename += "me";
281         key = "metas/" + key;
282     }
283
284     key += to_string(id);

```

```

284
285 if (mtfs::DATA_BLOCK == type) {
286     if (metas) {
287         Document d;
288
289         ((mtfs::blockInfo_t *) data)->toJson(d);
290
291         StringBuffer sb;
292         PrettyWriter<StringBuffer> pw(sb);
293         d.Accept(pw);
294
295         ofstream file(filename);
296         if (!file.is_open())
297             return errno != 0 ? errno : ENOENT;
298         file << sb.GetString() << endl;
299         file.close();
300     } else {
301         ofstream file(filename);
302         if (!file.is_open())
303             return errno != 0 ? errno : ENOENT;
304
305         file.write((char *) data, this->blocksize);
306         file.close();
307     }
308 } else {
309     Document d;
310
311     switch (type) {
312     case mtfs::INODE:
313         if (metas)
314             ((mtfs::blockInfo_t *) data)->toJson(d);
315         else
316             ((mtfs::inode_t *) data)->toJson(d);
317         break;
318     case mtfs::DIR_BLOCK:
319         if (metas)
320             ((mtfs::blockInfo_t *) data)->toJson(d);
321         else
322             ((mtfs::dirBlock_t *) data)->toJson(d);
323         break;
324     case mtfs::SUPERBLOCK:
325         break;
326     default:
327         return ENOSYS;
328     }
329
330     StringBuffer sb;
331     PrettyWriter<StringBuffer> pw(sb);
332     d.Accept(pw);
333
334     ofstream file(filename);
335     if (!file.is_open())
336         return errno != 0 ? errno : ENOENT;
337     file << sb.GetString() << endl;
338     file.close();
339 }
340
341 if (0 != this->upObj(filename, key))
342     return 1;

```

```

343     unlink(filename.c_str());
344     return 0;
345 }
346
347 bool S3::getSuperblock(mtf::superblock_t &superblock) {
348     return false;
349 }
350
351 bool S3::putSuperblock(const mtf::superblock_t &superblock) {
352     return false;
353 }
354
355 int S3::writeTmpFile(const string &filename, const void *data, const mtf::
356     blockType &type) {
357     switch (type) {
358         case mtf::INODE:
359             return this->writeInode(filename, (mtf::inode_t *) data);
360             break;
361         case mtf::DIR_BLOCK:
362             return this->writeDirBlock(filename, (mtf::dirBlock_t *) data);
363             break;
364         case mtf::DATA_BLOCK:
365             return this->writeDataBlock(filename, (uint8_t *) data);
366             break;
367         case mtf::SUPERBLOCK:
368             return this->writeSuperblock(filename, (mtf::superblock_t *) data);
369             break;
370         default:
371             return ENOSYS;
372     }
373 }
374
375 int S3::writeInode(const std::string &filename, const mtf::inode_t *inode) {
376     Document d;
377     d.SetObject();
378     Document::AllocatorType &alloc = d.GetAllocator();
379
380     inode->toJson(d);
381
382     StringBuffer sb;
383     PrettyWriter<StringBuffer> pw(sb);
384     d.Accept(pw);
385
386     ofstream inodeFile(filename);
387     inodeFile << sb.GetString() << endl;
388     inodeFile.close();
389
390     return this->SUCCESS;
391 }
392
393 int S3::writeDirBlock(const std::string &filename, const mtf::dirBlock_t *
394     dirBlock) {
395     Document d;
396     d.SetObject();
397
398     Document::AllocatorType &allocator = d.GetAllocator();
399
400     for (auto &item: dirBlock->entries) {

```

```

400     Value r(kArrayType);
401     for (auto &&ident: item.second) {
402         Value v(kObjectType);
403         ident.toJson(v, allocator);
404         r.PushBack(v, allocator);
405     }
406     d.AddMember(StringRef(item.first.c_str()), r, allocator);
407 }
408
409 StringBuffer sb;
410 PrettyWriter<StringBuffer> pw(sb);
411 d.Accept(pw);
412
413 ofstream file(filename);
414 if (!file.is_open())
415     return errno != 0 ? errno : ENOENT;
416 file << sb.GetString() << endl;
417 file.close();
418
419 return this->SUCCESS;
420 }
421
422 int S3::writeDataBlock(const std::string &filename, const uint8_t *datas) {
423     ofstream file(filename);
424     if (!file.is_open())
425         return errno != 0 ? errno : ENOENT;
426
427     file.write((const char *) datas, this->blocksize);
428     file.close();
429     return this->SUCCESS;
430 }
431
432 int S3::writeSuperblock(const std::string &filename, const mtfs::superblock_t
    *superblock) {
433     ofstream sbFile(filename, ios::binary);
434     sbFile.write((const char *) superblock, sizeof(*superblock));
435     sbFile.close();
436 }
437
438 int S3::dlObj(const std::string &filename, const std::string &key) {
439     Model::GetObjectRequest objectRequest;
440     objectRequest.WithBucket(this->bucket).WithKey(key.c_str());
441
442     auto getObjectOutcome = this->s3Client->GetObject(objectRequest);
443
444     if (getObjectOutcome.IsSuccess()) {
445         Aws::OFStream localFile;
446         localFile.open(filename, ios_base::out | ios_base::binary);
447         localFile << getObjectOutcome.GetResult().GetBody().rdbuf();
448         return 0;
449     }
450
451     return 1;
452 }
453
454 int S3::upObj(const std::string &filename, const std::string &key) {
455     Model::PutObjectRequest objectRequest;
456     objectRequest.WithBucket(this->bucket).WithKey(key.c_str());
457

```

```

458     auto inputData = Aws::MakeShared<Aws::FStream>("PutObjectInputStream",
459         filename.c_str(), std::ios_base::in);
460
461     objectRequest.SetBody(inputData);
462
463     auto putObjectOutcome = this->s3Client->PutObject(objectRequest);
464
465     if (putObjectOutcome.IsSuccess()) {
466 //         std::cout << "Done!" << std::endl;
467         return 0;
468     }
469     string message = string("PutObject error: ") + putObjectOutcome.GetError().
470     GetExceptionName().c_str() + " " +
471         putObjectOutcome.GetError().GetMessage().c_str();
472     Logger::getInstance()->log("S3", message, Logger::L_ERROR);
473
474     return 1;
475 }
476
477 void S3::delObj(const string &key) {
478     Aws::S3::Model::DeleteObjectRequest objectRequest;
479     objectRequest.WithBucket(this->bucket).WithKey(key.c_str());
480
481     auto deleteObjectOutcome = this->s3Client->DeleteObject(objectRequest);
482
483     if (deleteObjectOutcome.IsSuccess()) {
484         Logger::getInstance()->log("S3", "Delete Success", Logger::L_INFO);
485     } else {
486         string message =
487             string("PutObject error: ") + deleteObjectOutcome.GetError().
488             GetExceptionName().c_str() + " " +
489             deleteObjectOutcome.GetError().GetMessage().c_str();
490         Logger::getInstance()->log("S3", message, Logger::L_ERROR);
491     }
492 }
493
494 bool S3::dirExists(string path) {
495     struct stat info{};
496
497     if (stat(path.c_str(), &info) != 0)
498         return false;
499     return (info.st_mode & S_IFDIR) != 0;
500 }
501
502 void S3::initIds() {
503     string keyName = "ids.json";
504     string filename = this->home + keyName;
505
506     if (0 != dlObj(filename, keyName)) {
507         return;
508     }
509
510     ifstream inodeFile(filename);
511     if (!inodeFile.is_open())
512         return;
513
514     IStreamWrapper isw(inodeFile);
515
516     Document d;

```

```

515     d.ParseStream(isw);
516
517     assert(d.IsObject());
518
519     // Init inode ids
520     assert(d.HasMember("inode"));
521     Value o = d["inode"].GetObject();
522     assert(o.HasMember("next"));
523     this->nextFreeInode = o["next"].GetUInt64();
524     assert(o.HasMember("frees"));
525     for (auto &&id : o["frees"].GetArray()) {
526         this->freeInodes.push_back(id.GetUInt64());
527     }
528
529     // Init Dir block ids
530     assert(d.HasMember("directory"));
531     o = d["directory"].GetObject();
532     assert(o.HasMember("next"));
533     this->nextFreeDirectory = o["next"].GetUInt64();
534     assert(o.HasMember("frees"));
535     for (auto &&id : o["frees"].GetArray()) {
536         this->freeDirectory.push_back(id.GetUInt64());
537     }
538
539     // Init data Block Ids
540     assert(d.HasMember("data"));
541     o = d["data"].GetObject();
542     assert(o.HasMember("next"));
543     this->nextFreeData = o["next"].GetUInt64();
544     assert(o.HasMember("frees"));
545     for (auto &&id : o["frees"].GetArray()) {
546         this->freeDatas.push_back(id.GetUInt64());
547     }
548
549
550     unlink(filename.c_str());
551
552     Logger::getInstance()->log("S3", "INIT SUCCESS", Logger::L_INFO);
553 }
554
555 void S3::writeMetas() {
556     Document d;
557     d.SetObject();
558
559     Document::AllocatorType &allocator = d.GetAllocator();
560
561     Value o(kObjectType);
562     Value a(kArrayType);
563     Value v;
564     v.SetUInt64(this->nextFreeInode);
565     o.AddMember(StringRef("next"), v, allocator);
566     for (auto &&id : this->freeInodes) {
567         v.SetUInt64(id);
568         a.PushBack(v, allocator);
569     }
570     o.AddMember(StringRef("frees"), a, allocator);
571     d.AddMember(StringRef("inode"), o, allocator);
572
573     o.SetObject();

```

```

574     a.SetArray();
575     v.SetUInt64(this->nextFreeDirectory);
576     o.AddMember(StringRef("next"), v, allocator);
577     for (auto &&id : this->freeDirectory) {
578         v.SetUInt64(id);
579         a.PushBack(v, allocator);
580     }
581     o.AddMember(StringRef("frees"), a, allocator);
582     d.AddMember(StringRef("directory"), o, allocator);
583
584     o.SetObject();
585     a.SetArray();
586     v.SetUInt64(this->nextFreeData);
587     o.AddMember(StringRef("next"), v, allocator);
588     for (auto &&id : this->freeDatas) {
589         v.SetUInt64(id);
590         a.PushBack(v, allocator);
591     }
592     o.AddMember(StringRef("frees"), a, allocator);
593     d.AddMember(StringRef("data"), o, allocator);
594
595     StringBuffer strBuff;
596     PrettyWriter<StringBuffer> writer(strBuff);
597     d.Accept(writer);
598
599     string key = "ids.json";
600     string filename = this->home + key;
601     ofstream file;
602     file.open(filename);
603     file << strBuff.GetString() << endl;
604     file.close();
605
606     this->upObj(filename, key);
607
608     unlink(filename.c_str());
609 }
610 }
611
612 extern "C" pluginSystem::Plugin *createObj() {
613     return new pluginSystem::S3();
614 }
615
616 extern "C" void destroyObj(pluginSystem::Plugin *plugin) {
617     delete plugin;
618 }
619
620 extern "C" pluginSystem::pluginInfo_t getInfo() {
621     vector<string> params;
622     params.emplace_back("bucket");
623     params.emplace_back("region");
624
625     pluginSystem::pluginInfo_t info = {
626         .name = pluginSystem::S3::NAME,
627         .params = params,
628     };
629
630     return info;
631 }

```


Listing 3.9 – "pluginSystem/S3.cpp"

3.3 *Wrapper*

3.3.1 FuseBase

FuseBase.h

```

1  /**
2   * \file FuseBase.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #ifndef MTFSFUSE_FUSE_BASE_H
25 #define MTFSFUSE_FUSE_BASE_H
26
27 #include <fuse3/fuse_lowlevel.h>
28 #include <string>
29 #include <vector>
30 #include <list>
31 #include <iostream>
32 #include <assert.h>
33
34 #include "wrapper/FuseCallback.h"
35
36
37 namespace wrapper {
38     class FuseCallback;
39
40     class FuseBase {
41     public:
42         FuseBase();
43
44         virtual ~FuseBase() = 0;
45
46         int run(int argc, char **argv);
47
48 
```

```

49 protected:
50     friend class FuseCallback;
51
52     virtual bool runPrepare(int argc, char **argv)=0;
53
54     virtual void init(void *userdata, fuse_conn_info *conn);
55
56     virtual void destroy(void *userdata);
57
58     virtual void lookup(fuse_req_t req, fuse_ino_t parent, const char *name);
59
60     virtual void forget(fuse_req_t req, fuse_ino_t ino, uint64_t nlookup);
61
62     virtual void getAttr(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi);
63
64     virtual void setattr(fuse_req_t req, fuse_ino_t ino, struct stat *attr, int
65         to_set, fuse_file_info *fi);
66
67     virtual void readlink(fuse_req_t req, fuse_ino_t ino);
68
69     virtual void mknod(fuse_req_t req, fuse_ino_t parent, const char *name,
70         mode_t mode, dev_t rdev);
71
72     virtual void mkdir(fuse_req_t req, fuse_ino_t parent, const char *name,
73         mode_t mode);
74
75     virtual void unlink(fuse_req_t req, fuse_ino_t parent, const char *name);
76
77     virtual void rmdir(fuse_req_t req, fuse_ino_t parent, const char *name);
78
79     virtual void symlink(fuse_req_t req, const char *link, fuse_ino_t parent,
80         const char *name);
81
82     virtual void
83         rename(fuse_req_t req, fuse_ino_t parent, const char *name, fuse_ino_t
84             newparent, const char *newname,
85             unsigned int flags);
86
87     virtual void link(fuse_req_t req, fuse_ino_t ino, fuse_ino_t newparent,
88         const char *newname);
89
90     virtual void open(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi);
91
92     virtual void read(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
93         fuse_file_info *fi);
94
95     virtual void write(fuse_req_t req, fuse_ino_t ino, const char *buf, size_t
96         size, off_t off, fuse_file_info *fi);
97
98     virtual void flush(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi);
99
100    virtual void release(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi);
101
102    virtual void fsync(fuse_req_t req, fuse_ino_t ino, int datasync,
103        fuse_file_info *fi);
104
105    virtual void opendir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi);

```

```

99     virtual void readdir(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off
, fuse_file_info *fi);
100
101     virtual void releasedir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi)
;
102
103     virtual void fsyncdir(fuse_req_t req, fuse_ino_t ino, int datasync,
fuse_file_info *fi);
104
105     virtual void statfs(fuse_req_t req, fuse_ino_t ino);
106
107     virtual void
108     setattr(fuse_req_t req, fuse_ino_t ino, const char *name, const char *
value, size_t size, int flags);
109
110     virtual void getxattr(fuse_req_t req, fuse_ino_t ino, const char *name,
size_t size);
111
112     virtual void listxattr(fuse_req_t req, fuse_ino_t ino, size_t size);
113
114     virtual void removexattr(fuse_req_t req, fuse_ino_t ino, const char *name);
115
116     virtual void access(fuse_req_t req, fuse_ino_t ino, int mask);
117
118     virtual void create(fuse_req_t req, fuse_ino_t parent, const char *name,
mode_t mode, fuse_file_info *fi);
119
120     virtual void getlk(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi,
struct flock *lock);
121
122     virtual void setlk(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi,
struct flock *lock, int sleep);
123
124     virtual void bmap(fuse_req_t req, fuse_ino_t ino, size_t blocksize,
uint64_t idx);
125
126     virtual void ioctl(fuse_req_t req, fuse_ino_t ino, int cmd, void *arg,
fuse_file_info *fi, unsigned int flags,
127         const void *in_buf, size_t in_bufsz, size_t out_bufsz);
128
129     virtual void poll(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi,
fuse_pollhandle *ph);
130
131     virtual void write_buf(fuse_req_t req, fuse_ino_t ino, fuse_bufvec *bufv,
off_t off, fuse_file_info *fi);
132
133     virtual void retrieve_reply(fuse_req_t req, void *cookie, fuse_ino_t ino,
off_t offset, fuse_bufvec *bufv);
134
135     virtual void forget_multi(fuse_req_t req, size_t count, fuse_forget_data *
forgets);
136
137     virtual void flock(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi, int
op);
138
139     virtual void
140     fallocate(fuse_req_t req, fuse_ino_t ino, int mode, off_t offset, off_t
length, fuse_file_info *fi);
141

```

```

142     virtual void readdirplus(fuse_req_t req, fuse_ino_t ino, size_t size, off_t
        off, fuse_file_info *fi);
143
144     private:
145         FuseCallback *callbacks;
146
147     };
148
149 } // namespace wrapper
150 #endif

```

Listing 3.10 – "wrapper/FuseBase.h"

FuseBase.cpp

```

1  /**
2   * \file FuseBase.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10   MTFS is free software: you can redistribute it and/or modify
11   it under the terms of the GNU General Public License as published by
12   the Free Software Foundation, either version 3 of the License, or
13   (at your option) any later version.
14
15   Foobar is distributed in the hope that it will be useful,
16   but WITHOUT ANY WARRANTY; without even the implied warranty of
17   MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18   GNU General Public License for more details.
19
20   You should have received a copy of the GNU General Public License
21   along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #include <utils/Logger.h>
25 #include <fstream>
26 #include "wrapper/FuseBase.h"
27
28 namespace wrapper {
29
30     FuseBase::FuseBase() {
31         callbacks = FuseCallback::getInstance();
32     }
33
34     FuseBase::~FuseBase() = default;
35
36     int FuseBase::run(int argc, char **argv) {
37
38         struct fuse_args args = FUSE_ARGS_INIT(argc, argv);
39         struct fuse_session *se;
40         struct fuse_cmdline_opts opts{};
41         int ret = -1;
42
43         if (!runPrepare(argc, argv))
44             return -1;

```

```

45 //      std::cout << "run" << std::endl;
46
47 callbacks->setBase(this);
48
49 if (fuse_parse_cmdline(&args, &opts) != 0)
50     return 1;
51 if (0 != opts.show_help) {
52     std::cerr << "usage: " << argv[0] << " [options] <mountpoint> <configName>
>" << std::endl << std::endl;
53     fuse_cmdline_help();
54     fuse_lowlevel_help();
55     ret = 0;
56     goto err_out1;
57 } else if (0 != opts.show_version) {
58     std::cerr << "FUSE library version " << fuse_pkgversion() << std::endl;
59     fuse_lowlevel_version();
60     ret = 0;
61     goto err_out1;
62 }
63
64 if (0 != opts.debug) {
65     std::ofstream log("/tmp/mtfs.log");
66     Logger::init(std::cerr, Logger::L_DEBUG);
67 } else {
68     std::ofstream log;
69     log.open("/var/log/mtfs.log");
70     Logger::init(log, Logger::L_ERROR);
71 }
72
73 se = fuse_session_new(&args, &FuseCallback::ops, sizeof(FuseCallback::ops),
74     nullptr);
75
76 if (se == nullptr)
77     goto err_out1;
78
79 if (fuse_set_signal_handlers(se) != 0)
80     goto err_out2;
81
82 if (fuse_session_mount(se, opts.mountpoint) != 0)
83     goto err_out3;
84
85 fuse_daemonize(opts.foreground);
86
87 /* Block until ctrl+c or fusermount -u */
88 if (0 != opts.singlethread)
89     ret = fuse_session_loop(se);
90 else
91     ret = fuse_session_loop_mt(se, opts.clone_fd);
92
93 fuse_session_unmount(se);
94 err_out3:
95 fuse_remove_signal_handlers(se);
96 err_out2:
97 fuse_session_destroy(se);
98 err_out1:
99 free(opts.mountpoint);
100 fuse_opt_free_args(&args);
101
102 return 0 != ret ? 1 : 0;

```

```

102 }
103
104 void FuseBase::init(void *userdata, fuse_conn_info *conn) {
105     (void) userdata, conn;
106 }
107
108 void FuseBase::destroy(void *userdata) {
109     (void) userdata;
110 }
111
112 void FuseBase::lookup(fuse_req_t req, fuse_ino_t parent, const char *name) {
113     (void) parent, name;
114     fuse_reply_err(req, ENOSYS);
115 }
116
117 void FuseBase::forget(fuse_req_t req, fuse_ino_t ino, uint64_t nlookup) {
118     (void) ino, nlookup;
119     fuse_reply_err(req, ENOSYS);
120 }
121
122 void FuseBase::getattr(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
123     (void) ino, fi;
124     fuse_reply_err(req, ENOSYS);
125 }
126
127 void FuseBase::setattr(fuse_req_t req, fuse_ino_t ino, struct stat *attr, int
    to_set, fuse_file_info *fi) {
128     (void) ino, attr, to_set, fi;
129     fuse_reply_err(req, ENOSYS);
130 }
131
132 void FuseBase::readlink(fuse_req_t req, fuse_ino_t ino) {
133     (void) ino;
134     fuse_reply_err(req, ENOSYS);
135 }
136
137 void FuseBase::mknod(fuse_req_t req, fuse_ino_t parent, const char *name,
    mode_t mode, dev_t rdev) {
138     (void) parent, name, mode, rdev;
139     fuse_reply_err(req, ENOSYS);
140 }
141
142 void FuseBase::mkdir(fuse_req_t req, fuse_ino_t parent, const char *name,
    mode_t mode) {
143     (void) parent, name, mode;
144     fuse_reply_err(req, ENOSYS);
145 }
146
147 void FuseBase::unlink(fuse_req_t req, fuse_ino_t parent, const char *name) {
148     (void) parent, name;
149     fuse_reply_err(req, ENOSYS);
150 }
151
152 void FuseBase::rmdir(fuse_req_t req, fuse_ino_t parent, const char *name) {
153     (void) parent, name;
154     fuse_reply_err(req, ENOSYS);
155 }
156

```

```

157 void FuseBase::symlink(fuse_req_t req, const char *link, fuse_ino_t parent,
158   const char *name) {
159     (void) link, parent, name;
160     fuse_reply_err(req, ENOSYS);
161 }
162
163 void
164 FuseBase::rename(fuse_req_t req, fuse_ino_t parent, const char *name,
165   fuse_ino_t newparent, const char *newname,
166   unsigned int flags) {
167     (void) parent, name, newparent, newname, flags;
168     fuse_reply_err(req, ENOSYS);
169 }
170
171 void FuseBase::link(fuse_req_t req, fuse_ino_t ino, fuse_ino_t newparent,
172   const char *newname) {
173     (void) ino, newparent, newname;
174     fuse_reply_err(req, ENOSYS);
175 }
176
177 void FuseBase::open(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
178     (void) ino, fi;
179     fuse_reply_err(req, ENOSYS);
180 }
181
182 void FuseBase::read(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
183   fuse_file_info *fi) {
184     (void) ino, size, off, fi;
185     fuse_reply_err(req, ENOSYS);
186 }
187
188 void FuseBase::write(fuse_req_t req, fuse_ino_t ino, const char *buf, size_t
189   size, off_t off, fuse_file_info *fi) {
190     (void) ino, buf, size, off, fi;
191     fuse_reply_err(req, ENOSYS);
192 }
193
194 void FuseBase::flush(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
195     (void) ino, fi;
196     fuse_reply_err(req, ENOSYS);
197 }
198
199 void FuseBase::release(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
200     (void) ino, fi;
201     fuse_reply_err(req, ENOSYS);
202 }
203
204 void FuseBase::fsync(fuse_req_t req, fuse_ino_t ino, int datasync,
205   fuse_file_info *fi) {
206     (void) ino, datasync, fi;
207     fuse_reply_err(req, ENOSYS);
208 }
209
210 void FuseBase::opendir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
211     (void) ino, fi;
212     fuse_reply_err(req, ENOSYS);
213 }

```

```

209 void FuseBase::readdir(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off
    , fuse_file_info *fi) {
210     (void) ino, size, off, fi;
211     fuse_reply_err(req, ENOSYS);
212 }
213
214 void FuseBase::releasedir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi)
    {
215     (void) ino, fi;
216     fuse_reply_err(req, ENOSYS);
217 }
218
219 void FuseBase::fsyncdir(fuse_req_t req, fuse_ino_t ino, int datasync,
    fuse_file_info *fi) {
220     (void) ino, datasync, fi;
221     fuse_reply_err(req, ENOSYS);
222 }
223
224 void FuseBase::statfs(fuse_req_t req, fuse_ino_t ino) {
225     (void) ino;
226     fuse_reply_err(req, ENOSYS);
227 }
228
229 void
230 FuseBase::setxattr(fuse_req_t req, fuse_ino_t ino, const char *name, const
    char *value, size_t size, int flags) {
231     (void) ino, name, value, size, flags;
232     fuse_reply_err(req, ENOSYS);
233 }
234
235 void FuseBase::getxattr(fuse_req_t req, fuse_ino_t ino, const char *name,
    size_t size) {
236     (void) ino, name, size;
237     fuse_reply_err(req, ENOSYS);
238 }
239
240 void FuseBase::listxattr(fuse_req_t req, fuse_ino_t ino, size_t size) {
241     (void) ino, size;
242     fuse_reply_err(req, ENOSYS);
243 }
244
245 void FuseBase::removexattr(fuse_req_t req, fuse_ino_t ino, const char *name)
    {
246     (void) ino, name;
247     fuse_reply_err(req, ENOSYS);
248 }
249
250 void FuseBase::access(fuse_req_t req, fuse_ino_t ino, int mask) {
251     (void) ino, mask;
252     fuse_reply_err(req, ENOSYS);
253 }
254
255 void FuseBase::create(fuse_req_t req, fuse_ino_t parent, const char *name,
    mode_t mode, fuse_file_info *fi) {
256     (void) parent, name, mode, fi;
257     fuse_reply_err(req, ENOSYS);
258 }
259

```



```

260 void FuseBase::getlk(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi,
    struct flock *lock) {
261     (void) ino, fi, lock;
262     fuse_reply_err(req, ENOSYS);
263 }
264
265 void FuseBase::setlk(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi,
    struct flock *lock, int sleep) {
266     (void) ino, fi, lock, sleep;
267     fuse_reply_err(req, ENOSYS);
268 }
269
270 void FuseBase::bmap(fuse_req_t req, fuse_ino_t ino, size_t blocksize,
    uint64_t idx) {
271     (void) ino, blocksize, idx;
272     fuse_reply_err(req, ENOSYS);
273 }
274
275 void FuseBase::ioctl(fuse_req_t req, fuse_ino_t ino, int cmd, void *arg,
    fuse_file_info *fi, unsigned int flags,
276     const void *in_buf, size_t in_bufsz, size_t out_bufsz) {
277     (void) ino, cmd, arg, fi, flags, in_buf, in_bufsz, out_bufsz;
278     fuse_reply_err(req, ENOSYS);
279 }
280
281 void FuseBase::poll(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi,
    fuse_pollhandle *ph) {
282     (void) ino, fi, ph;
283     fuse_reply_err(req, ENOSYS);
284 }
285
286 void FuseBase::write_buf(fuse_req_t req, fuse_ino_t ino, fuse_bufvec *bufv,
    off_t off, fuse_file_info *fi) {
287     (void) ino, bufv, off, fi;
288     fuse_reply_err(req, ENOSYS);
289 }
290
291 void FuseBase::retrive_reply(fuse_req_t req, void *cookie, fuse_ino_t ino,
    off_t offset, fuse_bufvec *bufv) {
292     (void) cookie, ino, offset, bufv;
293     fuse_reply_err(req, ENOSYS);
294 }
295
296 void FuseBase::forget_multi(fuse_req_t req, size_t count, fuse_forget_data *
    forgets) {
297     (void) count, forgets;
298     fuse_reply_err(req, ENOSYS);
299 }
300
301 void FuseBase::flock(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi, int
    op) {
302     (void) ino, fi, op;
303     fuse_reply_err(req, ENOSYS);
304 }
305
306 void FuseBase::fallocate(fuse_req_t req, fuse_ino_t ino, int mode, off_t
    offset, off_t length, fuse_file_info *fi) {
307     (void) ino, mode, offset, length, fi;
308     fuse_reply_err(req, ENOSYS);

```

```

309 }
310
311 void FuseBase::readdirplus(fuse_req_t req, fuse_ino_t ino, size_t size, off_t
    off, fuse_file_info *fi) {
312     (void) ino, size, off, fi;
313     fuse_reply_err(req, ENOSYS);
314 }
315
316
317 } // namespace wrapper

```

Listing 3.11 – "wrapper/FuseBase.cpp"

3.3.2 FuseCallback

FuseCallback.h

```

1  /**
2   * \file FuseCallback.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10   MTFS is free software: you can redistribute it and/or modify
11   it under the terms of the GNU General Public License as published by
12   the Free Software Foundation, either version 3 of the License, or
13   (at your option) any later version.
14
15   Foobar is distributed in the hope that it will be useful,
16   but WITHOUT ANY WARRANTY; without even the implied warranty of
17   MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18   GNU General Public License for more details.
19
20   You should have received a copy of the GNU General Public License
21   along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #ifndef MTFSFUSE_FUSE_CALLBACK_H
25 #define MTFSFUSE_FUSE_CALLBACK_H
26
27 #include <fuse3/fuse_lowlevel.h>
28 #include <string>
29 #include <vector>
30 #include <list>
31 #include <iostream>
32
33 #include "wrapper/FuseBase.h"
34
35
36 namespace wrapper {
37     class FuseBase;
38
39     class FuseCallback {
40     public:
41         static FuseCallback *getInstance();

```

```

42     void setBase(FuseBase *base);
43
44     static struct fuse_lowlevel_ops ops;
45
46
47 private:
48     FuseCallback();
49
50     static void init(void *userdata, struct fuse_conn_info *conn);
51
52     static void destroy(void *userdata);
53
54     static void lookup(fuse_req_t req, fuse_ino_t parent, const char *name);
55
56     static void forget(fuse_req_t req, fuse_ino_t ino, std::uint64_t nlookup);
57
58     static void getAttr(fuse_req_t req, fuse_ino_t ino, struct fuse_file_info *
59 fi);
60
61     static void setattr(fuse_req_t req, fuse_ino_t ino, struct stat *attr, int
62 to_set, struct fuse_file_info *fi);
63
64     static void readlink(fuse_req_t req, fuse_ino_t ino);
65
66     static void mknod(fuse_req_t req, fuse_ino_t parent, const char *name,
67 mode_t mode, dev_t rdev);
68
69     static void mkdir(fuse_req_t req, fuse_ino_t parent, const char *name,
70 mode_t mode);
71
72     static void unlink(fuse_req_t req, fuse_ino_t parent, const char *name);
73
74     static void rmdir(fuse_req_t req, fuse_ino_t parent, const char *name);
75
76     static void symlink(fuse_req_t req, const char *link, fuse_ino_t parent,
77 const char *name);
78
79     static void
80 rename(fuse_req_t req, fuse_ino_t parent, const char *name, fuse_ino_t
81 newparent, const char *newname,
82         unsigned int flags);
83
84     static void link(fuse_req_t req, fuse_ino_t ino, fuse_ino_t newparent,
85 const char *newname);
86
87     static void open(fuse_req_t req, fuse_ino_t ino, struct fuse_file_info *fi)
88 ;
89
90     static void read(fuse_req_t req, fuse_ino_t ino, std::size_t size, off_t
91 off, struct fuse_file_info *fi);
92
93     static void
94 write(fuse_req_t req, fuse_ino_t ino, const char *buf, std::size_t size,
95 off_t off, struct fuse_file_info *fi);
96
97     static void flush(fuse_req_t req, fuse_ino_t ino, struct fuse_file_info *fi
98 );
99

```

```

90     static void release(fuse_req_t req, fuse_ino_t ino, struct fuse_file_info *
91     fi);
92
93     static void fsync(fuse_req_t req, fuse_ino_t ino, int datasync, struct
94     fuse_file_info *fi);
95
96     static void opendir(fuse_req_t req, fuse_ino_t ino, struct fuse_file_info *
97     fi);
98
99     static void readdir(fuse_req_t req, fuse_ino_t ino, struct
100     fuse_file_info *fi);
101
102     static void fsyncdir(fuse_req_t req, fuse_ino_t ino, int datasync, struct
103     fuse_file_info *fi);
104
105     static void statfs(fuse_req_t req, fuse_ino_t ino);
106
107     static void
108     setattr(fuse_req_t req, fuse_ino_t ino, const char *name, const char *
109     value, std::size_t size, int flags);
110
111     static void getxattr(fuse_req_t req, fuse_ino_t ino, const char *name, std
112     ::size_t size);
113
114     static void listxattr(fuse_req_t req, fuse_ino_t ino, std::size_t size);
115
116     static void removexattr(fuse_req_t req, fuse_ino_t ino, const char *name);
117
118     static void access(fuse_req_t req, fuse_ino_t ino, int mask);
119
120     static void create(fuse_req_t req, fuse_ino_t parent, const char *name,
121     mode_t mode, struct fuse_file_info *fi);
122
123     static void getlk(fuse_req_t req, fuse_ino_t ino, struct fuse_file_info *fi
124     , struct flock *lock);
125
126     static void setlk(fuse_req_t req, fuse_ino_t ino, struct fuse_file_info *fi
127     , struct flock *lock, int sleep);
128
129     static void bmap(fuse_req_t req, fuse_ino_t ino, std::size_t blocksize, std
130     ::uint64_t idx);
131
132     static void
133     ioctl(fuse_req_t req, fuse_ino_t ino, int cmd, void *arg, struct
134     fuse_file_info *fi, unsigned int flags,
135         const void *in_buf, std::size_t in_bufsz, std::size_t out_bufsz);
136
137     static void poll(fuse_req_t req, fuse_ino_t ino, struct fuse_file_info *fi,
138     struct fuse_pollhandle *ph);
139
140     static void
141     write_buf(fuse_req_t req, fuse_ino_t ino, struct fuse_bufvec *bufv, off_t
142     off, struct fuse_file_info *fi);
143
144     static void retrieve_reply(fuse_req_t req, void *cookie, fuse_ino_t ino,
145     off_t offset, struct fuse_bufvec *bufv);

```

```

133     static void forget_multi(fuse_req_t req, std::size_t count, struct
134     fuse_forget_data *forgets);
135
136     static void flock(fuse_req_t req, fuse_ino_t ino, struct fuse_file_info *fi
137     , int op);
138
139     static void
140     fallocate(fuse_req_t req, fuse_ino_t ino, int mode, off_t offset, off_t
141     length, struct fuse_file_info *fi);
142
143     static void readdirplus(fuse_req_t req, fuse_ino_t ino, std::size_t size,
144     off_t off, struct fuse_file_info *fi);
145
146     static FuseCallback *self;
147
148     static FuseBase *base;
149
150 };
151 } // namespace wrapper
152 #endif

```

Listing 3.12 – "wrapper/FuseCallback.h"

FuseCallback.cpp

```

1 #include <string>
2 #include <vector>
3 #include <list>
4 /**
5  * \file FuseCallback.cpp
6  * \brief
7  * \author David Wittwer
8  * \version 0.0.1
9  * \copyright GNU Publis License V3
10  *
11  * This file is part of MTEFS.
12
13  MTEFS is free software: you can redistribute it and/or modify
14  it under the terms of the GNU General Public License as published by
15  the Free Software Foundation, either version 3 of the License, or
16  (at your option) any later version.
17
18  Fooobar is distributed in the hope that it will be useful,
19  but WITHOUT ANY WARRANTY; without even the implied warranty of
20  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
21  GNU General Public License for more details.
22
23  You should have received a copy of the GNU General Public License
24  along with Fooobar. If not, see <http://www.gnu.org/licenses/>.
25  */
26
27 #include <iostream>
28 #include <assert.h>
29
30 #include <wrapper/FuseCallback.h>
31

```

```

32 namespace wrapper {
33     FuseCallback *FuseCallback::self = 0;
34     FuseBase *FuseCallback::base = 0;
35
36     struct fuse_lowlevel_ops FuseCallback::ops =
37     {
38         .init = FuseCallback::init ,
39         .destroy=FuseCallback::destroy ,
40         .lookup=FuseCallback::lookup ,
41         .forget=FuseCallback::forget ,
42         .getattr=FuseCallback::getattr ,
43         .setattr=FuseCallback::setattr ,
44         .readlink=FuseCallback::readlink ,
45         .mknod=FuseCallback::mknod ,
46         .mkdir=FuseCallback::mkdir ,
47         .unlink=FuseCallback::unlink ,
48         .rmdir=FuseCallback::rmdir ,
49         .symlink=FuseCallback::symlink ,
50         .rename=FuseCallback::rename ,
51         .link=FuseCallback::link ,
52         .open=FuseCallback::open ,
53         .read=FuseCallback::read ,
54         .write=FuseCallback::write ,
55         .flush=FuseCallback::flush ,
56         .release=FuseCallback::release ,
57         .fsync=FuseCallback::fsync ,
58         .opendir=FuseCallback::opendir ,
59         .readdir=FuseCallback::readdir ,
60         .releasedir=FuseCallback::releasedir ,
61         .fsyncdir=FuseCallback::fsyncdir ,
62         .statfs=FuseCallback::statfs ,
63         .setxattr=FuseCallback::setxattr ,
64         .getxattr=FuseCallback::getxattr ,
65         .listxattr=FuseCallback::listxattr ,
66         .removexattr=FuseCallback::removexattr ,
67         .access=FuseCallback::access ,
68         .create=FuseCallback::create ,
69         .getlk=FuseCallback::getlk ,
70         .setlk=FuseCallback::setlk ,
71         .bmap=FuseCallback::bmap ,
72         .ioctl=FuseCallback::ioctl ,
73         .poll=FuseCallback::poll ,
74         .write_buf=FuseCallback::write_buf ,
75         .retrieve_reply=FuseCallback::retrieve_reply ,
76         .forget_multi=FuseCallback::forget_multi ,
77         .flock=FuseCallback::flock ,
78         .fallocate=FuseCallback::fallocate ,
79         .readdirplus=FuseCallback::readdirplus ,
80     };
81
82     FuseCallback::FuseCallback() {
83         this->base = 0;
84     }
85
86     FuseCallback *FuseCallback::getInstance() {
87         if (!self)
88             self = new FuseCallback();
89
90         return self;

```

```

91 }
92
93 void FuseCallback::setBase(FuseBase *base) {
94     this->base = base;
95 }
96
97 void FuseCallback::init(void *userdata, struct fuse_conn_info *conn) {
98     if (base)
99         base->init(userdata, conn);
100 }
101
102 void FuseCallback::destroy(void *userdata) {
103     if (base)
104         base->destroy(userdata);
105 }
106
107 void FuseCallback::lookup(fuse_req_t req, fuse_ino_t parent, const char *name) {
108     if (base)
109         base->lookup(req, parent, name);
110     else
111         fuse_reply_err(req, ENOSYS);
112 }
113
114 void FuseCallback::forget(fuse_req_t req, fuse_ino_t ino, uint64_t nlookup) {
115     if (base)
116         base->forget(req, ino, nlookup);
117     else
118         fuse_reply_err(req, ENOSYS);
119 }
120
121 void FuseCallback::getattr(fuse_req_t req, fuse_ino_t ino, struct
122     fuse_file_info *fi) {
123     if (base)
124         base->getattr(req, ino, fi);
125     else
126         fuse_reply_err(req, ENOSYS);
127 }
128
129 void FuseCallback::setattr(fuse_req_t req, fuse_ino_t ino, struct stat *attr, int
130     to_set, struct fuse_file_info *fi) {
131     if (base)
132         base->setattr(req, ino, attr, to_set, fi);
133     else
134         fuse_reply_err(req, ENOSYS);
135 }
136
137 void FuseCallback::readlink(fuse_req_t req, fuse_ino_t ino) {
138     if (base)
139         base->readlink(req, ino);
140     else
141         fuse_reply_err(req, ENOSYS);
142 }
143
144 void FuseCallback::mknod(fuse_req_t req, fuse_ino_t parent, const char *name,
145     mode_t mode, dev_t rdev) {
146     if (base)
147         base->mknod(req, parent, name, mode, rdev);
148 }

```

```

146     else
147         fuse_reply_err(req, ENOSYS);
148 }
149
150 void FuseCallback::mkdir(fuse_req_t req, fuse_ino_t parent, const char *name,
151                          mode_t mode) {
152     if (base)
153         base->mkdir(req, parent, name, mode);
154     else
155         fuse_reply_err(req, ENOSYS);
156 }
157
158 void FuseCallback::unlink(fuse_req_t req, fuse_ino_t parent, const char *name)
159 {
160     if (base)
161         base->unlink(req, parent, name);
162     else
163         fuse_reply_err(req, ENOSYS);
164 }
165
166 void FuseCallback::rmdir(fuse_req_t req, fuse_ino_t parent, const char *name)
167 {
168     if (base)
169         base->rmdir(req, parent, name);
170     else
171         fuse_reply_err(req, ENOSYS);
172 }
173
174 void FuseCallback::symlink(fuse_req_t req, const char *link, fuse_ino_t
175                          parent, const char *name) {
176     if (base)
177         base->symlink(req, link, parent, name);
178     else
179         fuse_reply_err(req, ENOSYS);
180 }
181
182 void
183 FuseCallback::rename(fuse_req_t req, fuse_ino_t parent, const char *name,
184                    fuse_ino_t newparent, const char *newname,
185                    unsigned int flags) {
186     if (base)
187         base->rename(req, parent, name, newparent, newname, flags);
188     else
189         fuse_reply_err(req, ENOSYS);
190 }
191
192 void FuseCallback::link(fuse_req_t req, fuse_ino_t ino, fuse_ino_t newparent,
193                        const char *newname) {
194     if (base)
195         base->link(req, ino, newparent, newname);
196     else
197         fuse_reply_err(req, ENOSYS);
198 }
199
200 void FuseCallback::open(fuse_req_t req, fuse_ino_t ino, struct fuse_file_info
201                        *fi) {
202     if (base)
203         base->open(req, ino, fi);
204     else
205         fuse_reply_err(req, ENOSYS);
206 }

```



```

198     fuse_reply_err(req, ENOSYS);
199 }
200
201 void FuseCallback::read(fuse_req_t req, fuse_ino_t ino, size_t size, off_t
    off, struct fuse_file_info *fi) {
202     if (base)
203         base->read(req, ino, size, off, fi);
204     else
205         fuse_reply_err(req, ENOSYS);
206 }
207
208 void FuseCallback::write(fuse_req_t req, fuse_ino_t ino, const char *buf,
    size_t size, off_t off,
209                         struct fuse_file_info *fi) {
210     if (base)
211         base->write(req, ino, buf, size, off, fi);
212     else
213         fuse_reply_err(req, ENOSYS);
214 }
215
216 void FuseCallback::flush(fuse_req_t req, fuse_ino_t ino, struct
    fuse_file_info *fi) {
217     if (base)
218         base->flush(req, ino, fi);
219     else
220         fuse_reply_err(req, ENOSYS);
221 }
222
223 void FuseCallback::release(fuse_req_t req, fuse_ino_t ino, struct
    fuse_file_info *fi) {
224     if (base)
225         base->release(req, ino, fi);
226     else
227         fuse_reply_err(req, ENOSYS);
228 }
229
230 void FuseCallback::fsync(fuse_req_t req, fuse_ino_t ino, int datasync, struct
    fuse_file_info *fi) {
231     if (base)
232         base->fsync(req, ino, datasync, fi);
233     else
234         fuse_reply_err(req, ENOSYS);
235 }
236
237 void FuseCallback::opendir(fuse_req_t req, fuse_ino_t ino, struct
    fuse_file_info *fi) {
238     if (base)
239         base->opendir(req, ino, fi);
240     else
241         fuse_reply_err(req, ENOSYS);
242 }
243
244 void FuseCallback::readdir(fuse_req_t req, fuse_ino_t ino, size_t size, off_t
    off, struct fuse_file_info *fi) {
245     if (base)
246         base->readdir(req, ino, size, off, fi);
247     else
248         fuse_reply_err(req, ENOSYS);
249 }

```

```

250
251 void FuseCallback::releasedir(fuse_req_t req, fuse_ino_t ino, struct
    fuse_file_info *fi) {
252     if (base)
253         base->releasedir(req, ino, fi);
254     else
255         fuse_reply_err(req, ENOSYS);
256 }
257
258 void FuseCallback::fsyncdir(fuse_req_t req, fuse_ino_t ino, int datasync,
    struct fuse_file_info *fi) {
259     if (base)
260         base->fsyncdir(req, ino, datasync, fi);
261     else
262         fuse_reply_err(req, ENOSYS);
263 }
264
265 void FuseCallback::statfs(fuse_req_t req, fuse_ino_t ino) {
266     if (base)
267         base->statfs(req, ino);
268     else
269         fuse_reply_err(req, ENOSYS);
270 }
271
272 void FuseCallback::setxattr(fuse_req_t req, fuse_ino_t ino, const char *name,
    const char *value, size_t size,
273                          int flags) {
274     if (base)
275         base->setxattr(req, ino, name, value, size, flags);
276     else
277         fuse_reply_err(req, ENOSYS);
278 }
279
280 void FuseCallback::getxattr(fuse_req_t req, fuse_ino_t ino, const char *name,
    size_t size) {
281     if (base)
282         base->getxattr(req, ino, name, size);
283     else
284         fuse_reply_err(req, ENOSYS);
285 }
286
287 void FuseCallback::listxattr(fuse_req_t req, fuse_ino_t ino, size_t size) {
288     if (base)
289         base->listxattr(req, ino, size);
290     else
291         fuse_reply_err(req, ENOSYS);
292 }
293
294 void FuseCallback::removexattr(fuse_req_t req, fuse_ino_t ino, const char *
    name) {
295     if (base)
296         base->removexattr(req, ino, name);
297     else
298         fuse_reply_err(req, ENOSYS);
299 }
300
301 void FuseCallback::access(fuse_req_t req, fuse_ino_t ino, int mask) {
302     if (base)
303         base->access(req, ino, mask);

```

```

304     else
305         fuse_reply_err(req, ENOSYS);
306 }
307
308 void
309 FuseCallback::create(fuse_req_t req, fuse_ino_t parent, const char *name,
310                     mode_t mode, struct fuse_file_info *fi) {
311     if (base)
312         base->create(req, parent, name, mode, fi);
313     else
314         fuse_reply_err(req, ENOSYS);
315 }
316
317 void FuseCallback::getlk(fuse_req_t req, fuse_ino_t ino, struct
318                         fuse_file_info *fi, struct flock *lock) {
319     if (base)
320         base->getlk(req, ino, fi, lock);
321     else
322         fuse_reply_err(req, ENOSYS);
323 }
324
325 void FuseCallback::setlk(fuse_req_t req, fuse_ino_t ino, struct
326                         fuse_file_info *fi, struct flock *lock, int sleep) {
327     if (base)
328         base->setlk(req, ino, fi, lock, sleep);
329     else
330         fuse_reply_err(req, ENOSYS);
331 }
332
333 void FuseCallback::bmap(fuse_req_t req, fuse_ino_t ino, size_t blocksize,
334                        uint64_t idx) {
335     if (base)
336         base->bmap(req, ino, blocksize, idx);
337     else
338         fuse_reply_err(req, ENOSYS);
339 }
340
341 void FuseCallback::ioctl(fuse_req_t req, fuse_ino_t ino, int cmd, void *arg,
342                          struct fuse_file_info *fi,
343                          unsigned int flags,
344                          const void *in_buf, size_t in_bufsz, size_t out_bufsz) {
345     if (base)
346         base->ioctl(req, ino, cmd, arg, fi, flags, in_buf, in_bufsz, out_bufsz);
347     else
348         fuse_reply_err(req, ENOSYS);
349 }
350
351 void FuseCallback::poll(fuse_req_t req, fuse_ino_t ino, struct fuse_file_info
352                        *fi, struct fuse_pollhandle *ph) {
353     if (base)
354         base->poll(req, ino, fi, ph);
355     else
356         fuse_reply_err(req, ENOSYS);
357 }
358
359 void FuseCallback::write_buf(fuse_req_t req, fuse_ino_t ino, struct
360                             fuse_bufvec *bufv, off_t off,
361                             struct fuse_file_info *fi) {
362     if (base)

```

```

356     base->write_buf(req, ino, bufv, off, fi);
357     else
358         fuse_reply_err(req, ENOSYS);
359 }
360
361 void
362 FuseCallback::retrive_reply(fuse_req_t req, void *cookie, fuse_ino_t ino,
363     off_t offset, struct fuse_bufvec *bufv) {
364     if (base)
365         base->retrive_reply(req, cookie, ino, offset, bufv);
366     else
367         fuse_reply_err(req, ENOSYS);
368 }
369
370 void FuseCallback::forget_multi(fuse_req_t req, size_t count, struct
371     fuse_forget_data *forgets) {
372     if (base)
373         base->forget_multi(req, count, forgets);
374     else
375         fuse_reply_err(req, ENOSYS);
376 }
377
378 void FuseCallback::flock(fuse_req_t req, fuse_ino_t ino, struct
379     fuse_file_info *fi, int op) {
380     if (base)
381         base->flock(req, ino, fi, op);
382     else
383         fuse_reply_err(req, ENOSYS);
384 }
385
386 void
387 FuseCallback::fallocate(fuse_req_t req, fuse_ino_t ino, int mode, off_t
388     offset, off_t length,
389     struct fuse_file_info *fi) {
390     if (base)
391         base->fallocate(req, ino, mode, offset, length, fi);
392     else
393         fuse_reply_err(req, ENOSYS);
394 }
395
396 void
397 FuseCallback::readdirplus(fuse_req_t req, fuse_ino_t ino, size_t size, off_t
398     off, struct fuse_file_info *fi) {
399     if (base)
400         base->readdirplus(req, ino, size, off, fi);
401     else
402         fuse_reply_err(req, ENOSYS);
403 }
404
405 } // namespace wrapper

```

Listing 3.13 – "wrapper/FuseCallback.cpp"

3.3.3 MtfsFuse

MtfsFuse.h

```

1  /**
2   * \file MtfsFuse.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24  #ifndef MTFSFUSE_MTFS_FUSE_H
25  #define MTFSFUSE_MTFS_FUSE_H
26
27  #include <string>
28  #include <vector>
29  #include <list>
30  #include <iostream>
31  #include <cassert>
32
33  #include "wrapper/FuseBase.h"
34
35  namespace wrapper {
36  class MtfsFuse : public FuseBase {
37
38  public:
39
40      ~MtfsFuse() override;
41
42  protected:
43      bool runPrepare(int argc, char **argv) override;
44
45
46      void init(void *userdata, fuse_conn_info *conn) override;
47
48      void destroy(void *userdata) override;
49
50      void lookup(fuse_req_t req, fuse_ino_t parent, const char *name) override;
51
52      void getAttr(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) override;
53
54      void setattr(fuse_req_t req, fuse_ino_t ino, struct stat *attr, int to_set,
55                  fuse_file_info *fi) override;
56
57      void mknod(fuse_req_t req, fuse_ino_t parent, const char *name, mode_t mode,
58                dev_t rdev) override;
59

```

```

58 void open(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) override;
59
60 void release(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) override;
61
62 void opendir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) override;
63
64 void access(fuse_req_t req, fuse_ino_t ino, int mask) override;
65
66 void readdirplus(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
67 fuse_file_info *fi) override;
68
69 void releasedir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi)
70 override;
71
72 void readdir(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
73 fuse_file_info *fi) override;
74
75 void mkdir(fuse_req_t req, fuse_ino_t parent, const char *name, mode_t mode
76 ) override;
77
78 void
79 write(fuse_req_t req, fuse_ino_t ino, const char *buf, size_t size, off_t
80 off, fuse_file_info *fi) override;
81
82 void write_buf(fuse_req_t req, fuse_ino_t ino, fuse_bufvec *bufv, off_t off
83 , fuse_file_info *fi) override;
84
85 void read(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
86 fuse_file_info *fi) override;
87
88 void unlink(fuse_req_t req, fuse_ino_t parent, const char *name) override;
89
90 void rmdir(fuse_req_t req, fuse_ino_t parent, const char *name) override;
91 };
92
93 } // namespace wrapper
94 #endif

```

Listing 3.14 – "wrapper/MtfsFuse.h"

MtfsFuse.cpp

```

1 /**
2  * \file MtfsFuse.cpp
3  * \brief
4  * \author David Wittwer
5  * \version 0.0.1
6  * \copyright GNU Public License V3
7  *
8  * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the

```

```

18 GNU General Public License for more details.
19
20 You should have received a copy of the GNU General Public License
21 along with FooBar. If not, see <http://www.gnu.org/licenses/>.
22 */
23
24 #include <string>
25 #include <mtfs/Mtfs.h>
26
27 #include "wrapper/MtfsFuse.h"
28
29
30 namespace wrapper {
31
32     MtfsFuse::~MtfsFuse() {
33
34     }
35
36     bool MtfsFuse::runPrepare(int argc, char **argv) {
37         (void) argc, argv;
38         return true;
39     }
40
41     void MtfsFuse::init(void *userdata, fuse_conn_info *conn) {
42         mtfs::Mtfs::getInstance()->init(userdata, conn);
43     }
44
45     void MtfsFuse::destroy(void *userdata) {
46         mtfs::Mtfs::getInstance()->destroy(userdata);
47     }
48
49     void MtfsFuse::lookup(fuse_req_t req, fuse_ino_t parent, const char *name) {
50         mtfs::Mtfs::getInstance()->lookup(req, parent, name);
51     }
52
53     void MtfsFuse::getattr(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
54         mtfs::Mtfs::getInstance()->getattr(req, ino, fi);
55     }
56
57     void MtfsFuse::setattr(fuse_req_t req, fuse_ino_t ino, struct stat *attr, int
58         to_set, fuse_file_info *fi) {
59         mtfs::Mtfs::getInstance()->setAttr(req, ino, attr, to_set, fi);
60     }
61
62     void MtfsFuse::mknod(fuse_req_t req, fuse_ino_t parent, const char *name,
63         mode_t mode, dev_t rdev) {
64         mtfs::Mtfs::getInstance()->mknod(req, parent, name, mode, rdev);
65     }
66
67     void MtfsFuse::mkdir(fuse_req_t req, fuse_ino_t parent, const char *name,
68         mode_t mode) {
69         mtfs::Mtfs::getInstance()->mkdir(req, parent, name, mode);
70     }
71
72     void MtfsFuse::unlink(fuse_req_t req, fuse_ino_t parent, const char *name) {
73         mtfs::Mtfs::getInstance()->unlink(req, parent, name);
74     }
75

```

```

74 void Mtf Fuse::rmdir(fuse_req_t req, fuse_ino_t parent, const char *name) {
75     mtf s::Mtf s::getInstance()->rmdir(req, parent, name);
76 }
77
78 void Mtf Fuse::open(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
79     mtf s::Mtf s::getInstance()->open(req, ino, fi);
80 }
81
82 void Mtf Fuse::release(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
83     mtf s::Mtf s::getInstance()->release(req, ino, fi);
84 }
85
86 void Mtf Fuse::opendir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
87     mtf s::Mtf s::getInstance()->opendir(req, ino, fi);
88 }
89
90 void Mtf Fuse::access(fuse_req_t req, fuse_ino_t ino, int mask) {
91     mtf s::Mtf s::getInstance()->access(req, ino, mask);
92 }
93
94 void Mtf Fuse::readdir(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
95     fuse_file_info *fi) {
96     mtf s::Mtf s::getInstance()->readdir(req, ino, size, off, fi);
97 }
98
99 void Mtf Fuse::readdirplus(fuse_req_t req, fuse_ino_t ino, size_t size, off_t
100     off, fuse_file_info *fi) {
101     mtf s::Mtf s::getInstance()->readdirplus(req, ino, size, off, fi);
102 }
103
104 void Mtf Fuse::releasedir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi)
105 {
106     mtf s::Mtf s::getInstance()->releasedir(req, ino, fi);
107 }
108
109 void Mtf Fuse::write(fuse_req_t req, fuse_ino_t ino, const char *buf, size_t
110     size, off_t off, fuse_file_info *fi) {
111     mtf s::Mtf s::getInstance()->write(req, ino, buf, size, off, fi);
112 }
113
114 void Mtf Fuse::write_buf(fuse_req_t req, fuse_ino_t ino, fuse_bufvec *bufv,
115     off_t off, fuse_file_info *fi) {
116     mtf s::Mtf s::getInstance()->write_buf(req, ino, bufv, off, fi);
117 }
118
119 void Mtf Fuse::read(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
120     fuse_file_info *fi) {
121     mtf s::Mtf s::getInstance()->read(req, ino, size, off, fi);
122 }
123
124 } // namespace wrapper

```

Listing 3.15 – "wrapper/Mtf Fuse.cpp"

3.4 Core

3.4.1 structs.h


```

1 /**
2  * \file structs.h
3  * \brief
4  * \author David Wittwer
5  * \version 0.0.1
6  * \copyright GNU Publis License V3
7  *
8  * This file is part of MIFS.
9
10  MIFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #ifndef TRAVAIL_BACHELOR_STRUCTS_H
25 #define TRAVAIL_BACHELOR_STRUCTS_H
26
27 #include <map>
28 #include <rapidjson/document.h>
29
30 #define IN_MODE "mode"
31 #define IN_UID "uid"
32 #define IN_GID "gid"
33 #define IN_SIZE "size"
34 #define IN_LINKS "linkCount"
35 #define IN_ACCESS "atime"
36 #define IN_BLOCKS "dataBlocks"
37
38 #define ID_POOL "poolId"
39 #define ID_VOLUME "volumeId"
40 #define ID_ID "id"
41
42 #define RU_MIGRATION "migration"
43
44 #define PO_POOLS "pools"
45
46 #define SB_INODE_CACHE "inodeCacheSize"
47 #define SB_DIR_CACHE "directoryCacheSize"
48 #define SB_BLOCK_CACHE "blockCacheSize"
49 #define SB_BLOCK_SIZE_ST "blockSize"
50 #define SB_REDUNDANCY "redundancy"
51
52 #define BI_REFF "referenceId"
53 #define BI_ACCESS "lastAccess"
54
55 namespace mtfs {
56     class Rule;
57
58     enum blockType {
59         INODE,

```

```

60     DIR_BLOCK,
61     DATA_BLOCK,
62     SUPERBLOCK
63 };
64
65 typedef struct ruleInfo_st {
66     uid_t uid;
67     gid_t gid;
68     uint64_t lastAccess;
69
70     ruleInfo_st() : ruleInfo_st(0, 0, 0) {}
71
72     ruleInfo_st(uid_t uid, gid_t gid, uint64_t atime) : uid(uid), gid(gid),
73     lastAccess(atime) {}
74 } ruleInfo_t;
75
76 typedef struct ident_st {
77     std::uint32_t poolId;
78     std::uint32_t volumeId;
79     std::uint64_t id;
80
81     explicit ident_st(std::uint64_t id = 0, std::uint32_t vid = 0, std::
82     uint32_t pid = 0) : id(id),
83
84                                     volumeId(vid),
85                                     poolId(pid) {}
86
87     ~ident_st() = default;
88
89     ident_st &operator=(const ident_st &id) {
90         this->poolId = id.poolId;
91         this->volumeId = id.volumeId;
92         this->id = id.id;
93
94         return *this;
95     }
96
97     bool operator==(const ident_st &i) const {
98         return (poolId == i.poolId && volumeId == i.volumeId && id == i.id);
99     }
100
101     bool operator!=(const ident_st &i) const {
102         return (poolId != i.poolId || volumeId != i.volumeId || id != i.id);
103     }
104
105     bool operator<(const ident_st &i) const {
106         if (poolId < i.poolId)
107             return true;
108
109         if (volumeId < i.volumeId)
110             return true;
111
112         return id < i.id;
113     }
114
115     void toJson(rapidjson::Value &dest, rapidjson::Document::AllocatorType &
116     alloc) const {
117         rapidjson::Value v;
118
119         v.SetUint(this->poolId);

```

```

116     dest.AddMember(rapidjson::StringRef(ID_POOL), v, alloc);
117
118     v.SetUint(this->volumeId);
119     dest.AddMember(rapidjson::StringRef(ID_VOLUME), v, alloc);
120
121     v.SetUint64(this->id);
122     dest.AddMember(rapidjson::StringRef(ID_ID), v, alloc);
123 }
124
125 void fromJson(rapidjson::Value &src) {
126     assert(src.IsObject());
127     assert(src.HasMember(ID_POOL));
128     assert(src.HasMember(ID_VOLUME));
129     assert(src.HasMember(ID_ID));
130
131     this->poolId = src[ID_POOL].GetUint();
132     this->volumeId = src[ID_VOLUME].GetUint();
133     this->id = src[ID_ID].GetUint64();
134 }
135
136 std::string toString() {
137     return "p:" + std::to_string(this->poolId) + " v:" + std::to_string(this
->volumeId) + " i:" +
138         std::to_string(this->id);
139 }
140
141 } ident_t;
142
143 typedef struct inode_st {
144     mode_t accesRight{0};
145     uid_t uid{0};
146     gid_t gid{0};
147     uint64_t size{0};
148     uint32_t linkCount{1};
149     std::uint64_t atime;
150     std::vector<std::vector<ident_t>> dataBlocks;
151
152     inode_st() : atime((uint64_t) time(nullptr)) {
153         this->dataBlocks.clear();
154     }
155
156     bool operator==(const inode_st &rhs) const {
157         if (accesRight != rhs.accesRight)
158             return false;
159         if (uid != rhs.uid)
160             return false;
161         if (gid != rhs.gid)
162             return false;
163         if (size != rhs.size)
164             return false;
165         if (linkCount != rhs.linkCount)
166             return false;
167         if (atime != rhs.atime)
168             return false;
169
170         if (dataBlocks.size() != rhs.dataBlocks.size()) {
171 //             std::cout << "size differ " << dataBlocks.size() << " " << rhs.
dataBlocks.size() << std::endl;
172             return false;

```

```

173     }
174     for (auto lIter = dataBlocks.begin(), rIter = rhs.dataBlocks.begin();
175         lIter != dataBlocks.end(); lIter++, rIter++) {
176         std::vector<ident_t> red = *lIter, rhsRed = *rIter;
177
178         if (red.size() != rhsRed.size())
179             return false;
180         for (auto lb = red.begin(), rb = rhsRed.begin();
181             lb != red.end(); lb++, rb++) {
182             ident_t li = *lb, ri = *rb;
183
184             if (li != ri)
185                 return false;
186         }
187     }
188
189     return true;
190 }
191
192 void toJson(rapidjson::Document &dest) const {
193     dest.SetObject();
194
195     rapidjson::Value v;
196     rapidjson::Document::AllocatorType &alloc = dest.GetAllocator();
197
198     v.SetInt(this->accessRight);
199     dest.AddMember(rapidjson::StringRef(IN_MODE), v, alloc);
200
201     v.SetInt(this->uid);
202     dest.AddMember(rapidjson::StringRef(IN_UID), v, alloc);
203
204     v.SetInt(this->gid);
205     dest.AddMember(rapidjson::StringRef(IN_GID), v, alloc);
206
207     v.SetUint64(this->size);
208     dest.AddMember(rapidjson::StringRef(IN_SIZE), v, alloc);
209
210     v.SetUint(this->linkCount);
211     dest.AddMember(rapidjson::StringRef(IN_LINKS), v, alloc);
212
213     v.SetUint64(this->atime);
214     dest.AddMember(rapidjson::StringRef(IN_ACCESS), v, alloc);
215
216     rapidjson::Value a(rapidjson::kArrayType);
217     for (auto &&block : dataBlocks) {
218         rapidjson::Value red(rapidjson::kArrayType);
219
220         for (auto &&block : blocks) {
221             rapidjson::Value ident(rapidjson::kObjectType);
222
223             block.toJson(ident, alloc);
224
225             red.PushBack(ident, alloc);
226         }
227
228         a.PushBack(red, alloc);
229     }
230     dest.AddMember(rapidjson::StringRef(IN_BLOCKS), a, alloc);
231 }

```

```

232
233 void fromJson(rapidjson::Document &d) {
234     assert(d.IsObject());
235     assert(d.HasMember(IN_MODE));
236     assert(d.HasMember(IN_UID));
237     assert(d.HasMember(IN_GID));
238     assert(d.HasMember(IN_SIZE));
239     assert(d.HasMember(IN_LINKS));
240     assert(d.HasMember(IN_ACCESS));
241     assert(d.HasMember(IN_BLOCKS));
242
243     this->accesRight = (uint16_t) d[IN_MODE].GetUint();
244     this->uid = (uint16_t) d[IN_UID].GetUint();
245     this->gid = (uint16_t) d[IN_GID].GetUint();
246     this->size = d[IN_SIZE].GetUint64();
247     this->linkCount = (uint8_t) d[IN_LINKS].GetUint();
248     this->atime = d[IN_ACCESS].GetUint64();
249
250     const rapidjson::Value &dataArray = d[IN_BLOCKS];
251     assert(dataArray.IsArray());
252     this->dataBlocks.clear();
253     for (auto &a : dataArray.GetArray()) {
254         std::vector<ident_t> redundancy;
255
256         assert(a.IsArray());
257         for (auto &v : a.GetArray()) {
258             ident_t ident;
259
260             assert(v.IsObject());
261             assert(v.HasMember(ID_POOL));
262             assert(v.HasMember(ID_VOLUME));
263             assert(v.HasMember(ID_ID));
264
265             ident.poolId = (uint16_t) v[ID_POOL].GetUint();
266             ident.volumeId = (uint16_t) v[ID_VOLUME].GetUint();
267             ident.id = v[ID_ID].GetUint64();
268
269             redundancy.push_back(ident);
270         }
271
272         this->dataBlocks.push_back(redundancy);
273     }
274 }
275 } inode_t;
276
277 typedef struct dirBlock_st {
278     std::map<std::string, std::vector<mtfs::ident_t>> entries;
279
280     void fromJson(rapidjson::Document &d) {
281         assert(d.IsObject());
282
283         for (auto &&item: d.GetObject()) {
284             assert(item.value.IsArray());
285
286             std::vector<ident_t> ids;
287
288             for (auto &&ident: item.value.GetArray()) {
289                 ident_t i;
290                 i.fromJson(ident);

```

```

291     ids.push_back(i);
292 }
293 this->entries.emplace(item.name.GetString(), ids);
294 }
295 }
296
297 void toJson(rapidjson::Document &d) {
298     d.SetObject();
299
300     rapidjson::Document::AllocatorType &allocator = d.GetAllocator();
301
302     for (auto &&item: this->entries) {
303         rapidjson::Value r(rapidjson::kArrayType);
304         for (auto &&ident: item.second) {
305             rapidjson::Value v(rapidjson::kObjectType);
306             ident.toJson(v, allocator);
307             r.PushBack(v, allocator);
308         }
309         d.AddMember(rapidjson::StringRef(item.first.c_str()), r, allocator);
310     }
311 }
312
313 } dirBlock_t;
314
315 typedef struct volume_st {
316     std::string pluginName;
317     Rule *rule;
318     std::map<std::string, std::string> params;
319 } volume_t;
320
321 typedef struct pool_st {
322     int migration;
323     Rule *rule;
324     std::map<uint32_t, volume_t> volumes;
325 } pool_t;
326
327 typedef struct superbloc_k_st {
328     size_t iCacheSz;
329     size_t dCacheSz;
330     size_t bCacheSz;
331     size_t blockSz;
332     size_t redundancy;
333     int migration;
334     std::map<uint32_t, pool_t> pools;
335     std::vector<ident_t> rootInodes;
336 } superbloc_k_t;
337
338 typedef struct blockInfo_st {
339     ident_t id;
340     std::vector<ident_t> referenceId;
341     uint64_t lastAccess{0};
342
343     bool operator==(const blockInfo_st &other) const {
344         if (other.id != this->id)
345             return false;
346
347         if (!(other.referenceId == this->referenceId))
348             return false;
349

```

```

350     return (other.lastAccess == this->lastAccess);
351 }
352
353 blockInfo_st() : id(ident_t()) {};
354
355 void toJson(rapidjson::Document &d) {
356     d.SetObject();
357     rapidjson::Document::AllocatorType &allocator = d.GetAllocator();
358
359     rapidjson::Value v;
360
361     rapidjson::Value a(rapidjson::kArrayType);
362     for (auto &&id : this->referenceId) {
363         v.SetObject();
364
365         v.AddMember(rapidjson::StringRef(ID_POOL), rapidjson::Value(id.poolId),
366             allocator);
367         v.AddMember(rapidjson::StringRef(ID_VOLUME), rapidjson::Value(id.
368             volumeId), allocator);
369         v.AddMember(rapidjson::StringRef(ID_ID), rapidjson::Value(id.id),
370             allocator);
371
372         a.PushBack(v, allocator);
373     }
374     d.AddMember(rapidjson::StringRef(BI_REFF), a, allocator);
375
376     v.SetUint64(this->lastAccess);
377     d.AddMember(rapidjson::StringRef(BI_ACCESS), v, allocator);
378 }
379
380 void fromJson(rapidjson::Document &d) {
381     assert(d.IsObject());
382     assert(d.HasMember(BI_REFF));
383     assert(d[BI_REFF].IsArray());
384     assert(d.HasMember(BI_ACCESS));
385
386     for (auto &&id : d[BI_REFF].GetArray()) {
387         assert(id.IsObject());
388         assert(id.HasMember(ID_POOL));
389         assert(id.HasMember(ID_VOLUME));
390         assert(id.HasMember(ID_ID));
391
392         uint32_t poolId = id[ID_POOL].GetUint();
393         uint32_t volumeId = id[ID_VOLUME].GetUint();
394         uint64_t i = id[ID_ID].GetUint64();
395
396         this->referenceId.push_back(ident_st(i, volumeId, poolId));
397     }
398
399     this->lastAccess = d[BI_ACCESS].GetUint64();
400 }
401
402 } blockInfo_t;
403
404 } // namespace mtFS
405 #endif // TRAVAIL_BACHELOR_STRUCTS_H

```

Listing 3.16 – "core/structs.h"

3.4.2 Acces.h

```

1  /**
2   * \file Access.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFs.
9
10  MTFs is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24  #ifndef FILESTORAGE_INODE_ACCES_H
25  #define FILESTORAGE_INODE_ACCES_H
26
27  #include <string>
28  #include <vector>
29  #include <list>
30  #include <iostream>
31  #include <assert.h>
32  #include <mtfs/structs.h>
33
34  namespace mtfs {
35      typedef struct ruleInfo_st ruleInfo_t;
36
37      class Acces {
38      public:
39
40          virtual ~Acces() {};
41
42          /**
43           * Add a block in a volume
44           *
45           * @param [in] info ruleInfo_t To choose a pool or volume in which add the
46           block
47           * @param [out] ids Id for new blocks
48           * @param [in] type Type of block (see enum blockType)
49           * @param [in] nb Numbre of block to add.
50           *
51           * @return 0 if success else std linux error code.
52           */

```



```

52     virtual int add(const ruleInfo_t &info, std::vector<ident_t> &ids,
53     blockType type, size_t nb)=0;
54
55     /**
56     * Dele a block in a volume
57     *
58     * @param [in] id Id of block to delete
59     * @param [in] type Type of block to delete
60     * @return 0 if success else std linux error code.
61     */
62     virtual int del(const ident_t &id, blockType type)=0;
63
64     /**
65     * Get a block in a volume
66     *
67     * @param [in] id Id of block
68     * @param [out] data Pointer on memory already allocate
69     * @param [in] type Type of block
70     *
71     * @return 0 if success else std linux error code.
72     */
73     virtual int get(const ident_t &id, void *data, blockType type)=0;
74
75     /**
76     * Put a block in a volume
77     *
78     * @param [in] id Id of block
79     * @param [in] data Pointer on memory who contains the data.
80     * @param [in] type Type of block
81     *
82     * @return 0 if success else std linux error code.
83     */
84     virtual int put(const ident_t &id, const void *data, blockType type)=0;
85
86     /**
87     * Get metas block info
88     *
89     * @param [in] id Id of block
90     * @param [out] metas Datas
91     * @param [in] type Type of block
92     *
93     * @return 0 if success else std linux error code.
94     */
95     virtual int getMetas(const ident_t &id, blockInfo_t &metas, blockType type)
96     =0;
97
98     /**
99     * Put metas block info
100    *
101    * @param id Id of block
102    * @param metas Datas
103    * @param type Type of block
104    *
105    * @return 0 if success else std linux error code.
106    */
107    virtual int putMetas(const ident_t &id, const blockInfo_t &metas, blockType
    type)=0;
};

```

```

108 } // namespace mtfs
109 #endif
110

```

Listing 3.17 – "core/Acces.h"

3.4.3 Mtfs

Mtfs.h

```

1 /**
2  * \file Mtfs.h
3  * \brief
4  * \author David Wittwer
5  * \version 0.0.1
6  * \copyright GNU Publis License V3
7  *
8  * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #ifndef FILESTORAGE_MTFS_H
25 #define FILESTORAGE_MTFS_H
26
27 #include <thread>
28
29 #include <mtfs/Rule.h>
30 #include <mtfs/Acces.h>
31 #include <rapidjson/document.h>
32 #include <boost/threadpool.hpp>
33 #include <fuse3/fuse_lowlevel.h>
34 #include "structs.h"
35 #include "Migrator.h"
36 #include <mutex>
37 #include <condition_variable>
38 #include <utils/Semaphore.h>
39
40 namespace mtfs {
41     struct internalInode_st;
42     struct dl_st;
43
44     class Mtfs {
45     public:
46         // static constexpr const char *SYSTEMS_DIR = "Systems";
47         static constexpr const char *CONFIG_DIR = "Configs";
48

```

```

49     static constexpr const char *INODE_CACHE = "inodeCacheSize";
50     static constexpr const char *DIR_CACHE = "directoryCacheSize";
51     static constexpr const char *BLOCK_CACHE = "blockCacheSize";
52     static constexpr const char *BLOCK_SIZE_ST = "blockSize";
53     static constexpr const char *REDUNDANCY = "redundancy";
54     static constexpr const char *ROOT_INODES = "rootInodes";
55
56
57 private:
58 //     CONFIG
59     static const size_t SIMULT_DL = 2;
60     static const size_t SIMULT_UP = 2;
61     static const int INIT_DL = 2;
62     static constexpr const double ATTR_TIMEOUT = 1.0;
63
64 //     REQUEST STATUS CODES
65     static const int SUCCESS = 0;
66     static const int PENDING = 9999;
67
68     static MtfS *instance;
69     static boost::threadpool::pool *threadPool;
70     static std::string systemName;
71
72     size_t redundancy;
73     size_t blockSize;
74     int maxEntryPerBlock;
75     internalInode_st *rootIn;
76
77     Acces *inodes;
78     Acces *dirBlocks;
79     Acces *blocks;
80
81     std::thread migratorThr;
82     Migrator::info_st migratorInfo;
83
84 public:
85     static MtfS *getInstance();
86
87     /**
88      * Validate system config
89      *
90      * @param system JSON config of system
91      *
92      * @return true if valid else false
93      */
94     static bool validate(const rapidjson::Value &system);
95
96     /**
97      * Create the root inode
98      *
99      * @param [out] inode The root inode
100     *
101     * @return true if success
102     */
103     static bool createRootInode(inode_t &inode);
104
105     /**
106     * Start MTFs
107     */

```

```

108     * @param system System config
109     * @param homeDir MIFS home dir
110     * @param sysName System name
111     *
112     * @return true if success
113     */
114     static bool start(rapidjson::Document &system, std::string homeDir, std::
string sysName);
115
116     /**
117     * Stop MIFS
118     */
119     static void stop();
120
121     /**
122     * Convert struct superblock to JSON document
123     *
124     * @param [in] sb Superblock struct
125     * @param [out] d JSON document
126     */
127     static void structToJson(const superblock_t &sb, rapidjson::Document &d);
128
129     /**
130     * Convert JSON document to struct superblock
131     *
132     * @param [in] d JSON document
133     * @param [out] sb Struct superblock
134     */
135     static void jsonToStruct(rapidjson::Document &d, superblock_t &sb);
136
137     /*          Fuse handlers          */
138
139     /**
140     * Init handler
141     *
142     * @param userdata
143     * @param conn
144     */
145     void init(void *userdata, fuse_conn_info *conn);
146
147     /**
148     * Destroy handler
149     *
150     * @param userdata
151     */
152     void destroy(void *userdata);
153
154     /**
155     * Lookup handler
156     *
157     * @param req
158     * @param parent
159     * @param name
160     */
161     void lookup(fuse_req_t req, fuse_ino_t parent, const std::string &name);
162
163     /**
164     * Getattr handler
165     *

```

```

166     * @param req
167     * @param ino
168     * @param fi
169     */
170 void getAttr(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi);
171
172 /**
173  * Setattr handler
174  *
175  * @param req
176  * @param ino
177  * @param attr
178  * @param toSet
179  * @param fi
180  */
181 void setattr(fuse_req_t req, fuse_ino_t ino, struct stat *attr, int toSet,
182             fuse_file_info *fi);
183
184 /**
185  * Mknod handler
186  *
187  * @param req
188  * @param parent
189  * @param name
190  * @param mode
191  * @param rdev
192  */
193 void mknod(fuse_req_t req, fuse_ino_t parent, const std::string &name,
194            mode_t mode, dev_t rdev);
195
196 /**
197  * Mkdir handler
198  *
199  * @param req
200  * @param ino
201  * @param name
202  * @param mode
203  */
204 void mkdir(fuse_req_t req, fuse_ino_t ino, const std::string &name, mode_t
205            mode);
206
207 /**
208  * Unlink handler
209  *
210  * @param req
211  * @param parent
212  * @param name
213  */
214 void unlink(fuse_req_t req, fuse_ino_t parent, const std::string &name);
215
216 /**
217  * Rmdir handler
218  *
219  * @param req
220  * @param parent
221  * @param name
222  */
223 void rmdir(fuse_req_t req, fuse_ino_t parent, const std::string &name);

```

```

222  /**
223   * Open handler
224   *
225   * @param req
226   * @param ino
227   * @param fi
228   */
229  void open(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi);
230
231  /**
232   * Read handler
233   *
234   * @param req
235   * @param ino
236   * @param size
237   * @param off
238   * @param fi
239   */
240  void read(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
241            fuse_file_info *fi);
242
243  /**
244   * Write handler
245   *
246   * @param req
247   * @param ino
248   * @param buf
249   * @param size
250   * @param off
251   * @param fi
252   */
253  void write(fuse_req_t req, fuse_ino_t ino, const char *buf, size_t size,
254            off_t off, fuse_file_info *fi);
255
256  /**
257   * Release handler
258   *
259   * @param req
260   * @param ino
261   * @param fi
262   */
263  void release(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi);
264
265  /**
266   * Opendir handler
267   *
268   * @param req
269   * @param ino
270   * @param fi
271   */
272  void opendir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi);
273
274  /**
275   * Readdir handler
276   *
277   * @param req
278   * @param ino
279   * @param size
280   * @param off

```

```

279     * @param fi
280     */
281     void readdir(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
282                 fuse_file_info *fi);
283
284     /**
285     * Releasedir handler
286     *
287     * @param req
288     * @param ino
289     * @param fi
290     */
291     void releasedir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi);
292
293     /**
294     * Access handler
295     *
296     * @param req
297     * @param ino
298     * @param mask
299     */
300     void access(fuse_req_t req, fuse_ino_t ino, int mask);
301
302     /**
303     * Write_buf handler
304     *
305     * @param req
306     * @param ino
307     * @param bufv
308     * @param off
309     * @param fi
310     */
311     void write_buf(fuse_req_t req, fuse_ino_t ino, fuse_bufvec *bufv, off_t off,
312                   fuse_file_info *fi);
313
314     /**
315     * Readdirplus handler
316     *
317     * @param req
318     * @param ino
319     * @param size
320     * @param off
321     * @param fi
322     */
323     void readdirplus(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
324                      fuse_file_info *fi);
325
326 private:
327
328     Mtfs();
329
330     inode_t getRootInode();
331
332     bool build(const superblock_t &superblock);
333
334     void readRootInode();

```

```

335     void writeRootInode();
336
337     int addEntry(internalInode_st *parentInode, std::string name, std::vector<
ident_t> &inodeIds);
338
339     int insertInode(const inode_t &inode, std::vector<ident_t> &idents);
340
341     void dlBlocks(const inode_t &inode, dl_st *dlSt, const blockType type,
const int firstBlockIdx);
342
343     void dlDirBlocks(std::vector<ident_t> &ids, std::queue<dirBlock_t> *q, std
::mutex *queueMutex, Semaphore *sem);
344
345     void dlInodes(dl_st *src, dl_st *dst);
346
347     void initMetas(const internalInode_st &parentInode, const std::vector<
ident_t> &ids, const blockType &type,
348                 boost::threadpool::pool *thPool = nullptr);
349
350     void
351     dlInode(std::vector<ident_t> &ids, std::queue<std::pair<std::string,
inode_t>> *queue, std::mutex *queueMutex,
352            Semaphore *sem,
353            std::string &key);
354
355     ///////////////////////////////////
356     internalInode_st *getIntInode(fuse_ino_t ino);
357
358     void
359     doReaddir(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
fuse_file_info *fi, const bool &plus = false);
360
361     size_t dirBufAdd(fuse_req_t &req, char *buf, size_t &currentSize, std::
string name, internalInode_st &inode,
362                    const bool &plus);
363
364     void buildParam(const internalInode_st &inode, fuse_entry_param &param);
365
366     void buildStat(const internalInode_st &inode, struct stat &st);
367
368     ruleInfo_t getRuleInfo(const inode_t &inode);
369
370     static internalInode_st *newInode(const mode_t &mode, const fuse_ctx *ctx);
371
372     static uint64_t now();
373
374     int doUnlink(internalInode_st *parent, const std::string name);
375
376     int delEntry(std::vector<ident_t> &ids, dirBlock_t &blk, const std::string
&name);
377 };
378
379 // namespace mtfs
380 #endif

```

Listing 3.18 – "core/Mtfs.h"

Mtfs.cpp

```

1  /**
2   * \file Mtfs.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  FooBar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with FooBar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #include <mtfs/Mtfs.h>
25 #include <mtfs/Cache.h>
26 #include <pluginSystem/PluginManager.h>
27 #include <fstream>
28 #include <rapidjson/istreamwrapper.h>
29 #include <rapidjson/stringbuffer.h>
30 #include <rapidjson/prettywriter.h>
31 #include <boost/filesystem.hpp>
32 #include <utility>
33 #include <grp.h>
34 #include <pwd.h>
35 #include <utils/Logger.h>
36
37 #define GROUPS_TO_SEARCH 30
38 //TODO compute this value in function of entry size
39 #define ENTRY_PER_BLOCK 20
40
41 namespace mtfs {
42     using namespace std;
43     using namespace rapidjson;
44     using namespace boost::threadpool;
45
46     struct internalInode_st {
47         inode_t inode;
48         vector<ident_t> idents;
49     };
50
51     struct dl_st {
52         mutex *fifoMu{nullptr};
53         Semaphore *sem{nullptr};
54         union {
55             queue<pair<string, inode_t>> *inodeQueue;
56             queue<dirBlock_t> *dirQueue;
57             queue<uint8_t*> *blkQueue;

```

```

58     } fifo;
59     bool end{false};
60     mutex *endMu{nullptr};
61     pool *dlThreadPool{nullptr};
62
63     dl_st() {
64     }
65 };
66
67 struct fd_st {
68     enum bt {
69         DATA,
70         DIR,
71     };
72
73     bool firstCall{true};
74     bt blockT;
75     dl_st blkDl;
76     dl_st inoDl;
77     boost::thread *blkThr{nullptr};
78     boost::thread *inoThr{nullptr};
79
80     fd_st() : blkDl(dl_st()), inoDl(dl_st()) {}
81
82     ~fd_st() {
83         switch (this->blockT) {
84             case DATA:
85                 delete blkDl.fifo.blkQueue;
86                 break;
87             case DIR:
88                 delete blkDl.fifo.dirQueue;
89                 break;
90         }
91     }
92 };
93
94
95 MtfS *MtfS::instance = nullptr;
96 pool *MtfS::threadPool = nullptr;
97 string MtfS::systemName;
98
99 //
100 ///////////////////////////////////////////////////
101 //                                     STATICS                                     //
102 ///////////////////////////////////////////////////
103
104 MtfS *MtfS::getInstance() {
105     if (instance == nullptr)
106         instance = new MtfS();
107
108     return instance;
109 }
110
111 bool MtfS::validate(const Value &system) {
112     if (!system.IsObject())
113         throw invalid_argument("Not a object!");

```

```

113     if (!system.HasMember(INODE_CACHE))
114         throw invalid_argument("Inode cache missing!");
115
116     if (!system.HasMember(DIR_CACHE))
117         throw invalid_argument("Directory cache missing!");
118
119     if (!system.HasMember(BLOCK_CACHE))
120         throw invalid_argument("Block cache missing!");
121
122     if (!system.HasMember(BLOCK_SIZE_ST))
123         throw invalid_argument("Block size missing!");
124
125     if (!system.HasMember(REDUNDANCY))
126         throw invalid_argument("Redundancy missing!");
127
128     if (!system.HasMember(Pool::POOLS))
129         throw invalid_argument("Pools missing!");
130     if (!system[Pool::POOLS].IsObject())
131         throw invalid_argument("Pool is not a object!");
132
133     int migration = -1;
134     if (system[Pool::POOLS].MemberCount() <= 0)
135         throw invalid_argument("Number of pool invalid!");
136     else if (system[Pool::POOLS].MemberCount() != 1) {
137         if (!system.HasMember(Rule::MIGRATION))
138             throw invalid_argument("Migration missing!");
139
140         migration = system[Rule::MIGRATION].GetInt();
141     }
142
143     for (auto &m: system[Pool::POOLS].GetObject()) {
144         if (Rule::rulesAreValid(migration, m.value) != Rule::VALID_RULES)
145             throw invalid_argument(string("Rules invalid for pool ") + m.name.
146 GetString() + "'");
147
148         if (!Pool::validate(m.value))
149             throw invalid_argument(string("Pool ") + m.name.GetString() + "'
150 invalid!");
151     }
152     return true;
153 }
154
155 bool Mtfs::createRootInode(inode_t &inode) {
156     inode.accessRight = S_IFDIR | 0775;
157     inode.uid = 0;
158     inode.gid = 1001;
159     inode.size = 0;
160     inode.linkCount = 2;
161     inode.atime = (uint64_t) time(nullptr);
162     inode.dataBlocks.clear();
163
164     return true;
165 }
166
167 bool Mtfs::start(Document &system, std::string homeDir, string sysName) {
168     (void) homeDir;
169     sysName = std::move(sysName);

```

```

170     superblock_t superblock;
171     jsonToStruct(system, superblock);
172
173     if (!getInstance()->build(superblock))
174         return false;
175
176     auto nbThread = (unsigned int) (thread::hardware_concurrency() * 1.25);
177     threadPool = new pool(nbThread);
178
179     return true;
180 }
181
182 void Mtf::stop() {
183     delete threadPool;
184
185     delete instance;
186 }
187
188 void Mtf::structToJson(const superblock_t &sb, rapidjson::Document &d) {
189     rapidjson::Document::AllocatorType &allocator = d.GetAllocator();
190
191     d.SetObject();
192
193     rapidjson::Value v;
194
195     v.SetUint((unsigned int) sb.iCacheSz);
196     d.AddMember(rapidjson::StringRef(Mtf::INODE_CACHE), v, allocator);
197
198     v.SetUint((unsigned int) sb.dCacheSz);
199     d.AddMember(rapidjson::StringRef(Mtf::DIR_CACHE), v, allocator);
200
201     v.SetUint((unsigned int) sb.bCacheSz);
202     d.AddMember(rapidjson::StringRef(Mtf::BLOCK_CACHE), v, allocator);
203
204     v.SetUint((unsigned int) sb.blockSz);
205     d.AddMember(rapidjson::StringRef(Mtf::BLOCK_SIZE_ST), v, allocator);
206
207     v.SetUint((unsigned int) sb.redundancy);
208     d.AddMember(rapidjson::StringRef(Mtf::REDUNDANCY), v, allocator);
209
210     v.SetInt(sb.migration);
211     d.AddMember(rapidjson::StringRef(Rule::MIGRATION), v, allocator);
212
213     rapidjson::Value pools(rapidjson::kObjectType);
214     for (auto &item: sb.pools) {
215         rapidjson::Value pool(rapidjson::kObjectType);
216
217         Pool::structToJson(item.second, pool, allocator);
218
219         string id = to_string(item.first);
220         Value index(id.c_str(), (SizeType) id.size(), allocator);
221         pools.AddMember(index, pool, allocator);
222     }
223     d.AddMember(rapidjson::StringRef(Pool::POOLS), pools, allocator);
224
225     v.SetArray();
226     for (auto &id: sb.rootInodes) {
227         Value ident(kObjectType);
228

```

```

229     id.toJson(ident, allocator);
230     v.PushBack(ident, allocator);
231 }
232 d.AddMember(StringRef(ROOT_INODES), v, allocator);
233 }
234
235 void Mtf::jsonToStruct(rapidjson::Document &d, superblock_t &sb) {
236     assert(d.HasMember(INODE_CACHE));
237     sb.iCacheSz = d[INODE_CACHE].GetUint();
238
239     assert(d.HasMember(DIR_CACHE));
240     sb.dCacheSz = d[DIR_CACHE].GetUint();
241
242     assert(d.HasMember(BLOCK_CACHE));
243     sb.bCacheSz = d[BLOCK_CACHE].GetUint();
244
245     assert(d.HasMember(BLOCK_SIZE_ST));
246     sb.blockSz = d[BLOCK_SIZE_ST].GetUint();
247
248     assert(d.HasMember(REDUNDANCY));
249     sb.redundancy = d[REDUNDANCY].GetUint();
250
251     assert(d.HasMember(Rule::MIGRATION));
252     sb.migration = d[Rule::MIGRATION].GetInt();
253
254     assert(d.HasMember(Pool::POOLS));
255     for (auto &&item : d[Pool::POOLS].GetObject()) {
256         string sId = item.name.GetString();
257         auto id = (uint32_t) stoul(sId);
258         pool_t pool;
259         memset(&pool, 0, sizeof(pool_t));
260         pool.volumes.clear();
261
262         pool.rule = Rule::buildRule(sb.migration, item.value);
263
264         Pool::jsonToStruct(item.value, pool);
265
266         sb.pools.insert(make_pair(id, pool));
267     }
268
269     assert(d.HasMember(ROOT_INODES) && d[ROOT_INODES].IsArray());
270     for (auto &&ident : d[ROOT_INODES].GetArray()) {
271         ident_t id = ident_t();
272         id.fromJson(ident);
273         sb.rootInodes.push_back(id);
274     }
275 }
276
277 //
278 ////////////////////////////////////////
279
280 /// MEMBERS ///
281 //
282 ////////////////////////////////////////

```

```

282  ///                                     FUSE fcts                                     ///
283  //
284  //////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////////
285  void Mtfs::init(void *userdata, fuse_conn_info *conn) {
286      (void) userdata, conn;
287
288      getInstance()->readRootInode();
289
290      this->migratorInfo.poolManager = (PoolManager *) this->blocks;
291
292      this->migratorThr = std::thread(Migrator::main, &this->migratorInfo);
293
294  #ifdef DEBUG
295      cerr << "[MTFS]: End init" << endl;
296  #endif
297  }
298
299  void Mtfs::destroy(void *userdata) {
300      (void) userdata;
301
302      unique_lock<mutex> lk(*this->migratorInfo.endMutex);
303      this->migratorInfo.end = true;
304      lk.unlock();
305      this->migratorInfo.condV.notify_all();
306
307      this->migratorThr.join();
308
309      Acces *iptr = this->inodes;
310      Acces *dptr = this->dirBlocks;
311      Acces *bptr = this->blocks;
312
313      if (iptr != dptr && iptr != bptr)
314          delete this->inodes;
315
316      if (dptr != bptr)
317          delete this->dirBlocks;
318      delete this->blocks;
319  }
320
321  void Mtfs::getattr(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
322      (void) fi;
323
324      internalInode_st *inode = this->getIntInode(ino);
325
326      if (nullptr == inode)
327          return (void) fuse_reply_err(req, ENOENT);
328
329      struct stat st{};
330      this->buildStat(*inode, st);
331
332      fuse_reply_attr(req, &st, this->ATTR_TIMEOUT);
333  }
334
335  void Mtfs::setattr(fuse_req_t req, fuse_ino_t ino, struct stat *attr, int
336      toSet, fuse_file_info *fi) {
337      (void) fi;

```

```

338     internalInode_st *inInode = this->getIntInode(ino);
339
340     if (0 != (FUSE_SET_ATTR_MODE & toSet))
341         inInode->inode.accesRight = attr->st_mode;
342
343     if (0 != (FUSE_SET_ATTR_UID & toSet))
344         inInode->inode.uid = attr->st_uid;
345
346     if (0 != (FUSE_SET_ATTR_GID & toSet))
347         inInode->inode.gid = attr->st_gid;
348
349     if (0 != (FUSE_SET_ATTR_SIZE & toSet))
350         inInode->inode.size = (uint64_t) attr->st_size;
351
352     if (0 != (FUSE_SET_ATTR_ATIME & toSet))
353         inInode->inode.atime = (uint64_t) attr->st_atim.tv_sec;
354
355     buildStat(*inInode, *attr);
356     fuse_reply_attr(req, attr, 1.0);
357
358     pool upThPool(this->SIMULT_UP);
359     for (auto &&id: inInode->idents) {
360         upThPool.schedule(bind(&Acces::put, this->inodes, id, &inInode->inode,
361                               INODE));
362     }
363
364 void Mfts::lookup(fuse_req_t req, fuse_ino_t parent, const string &name) {
365     Logger::getInstance()->log("LOOKUP", name, Logger::L_DEBUG);
366
367     internalInode_st *parentInode = this->getIntInode(parent);
368
369 //     dl All blocks
370     dl_st blkDl = dl_st();
371     blkDl.dlThPool = new pool(this->SIMULT_DL);
372     blkDl.sem = new Semaphore();
373     blkDl.fifoMu = new mutex;
374     blkDl.fifo.dirQueue = new queue<dirBlock_t>();
375     blkDl.endMu = new mutex;
376
377     this->dlBlocks(parentInode->inode, &blkDl, blockType::DIR_BLOCK, 0);
378
379     vector<ident_t> inodeIds;
380
381     unique_lock<mutex> endLk(*blkDl.endMu, defer_lock);
382     unique_lock<mutex> fifoLk(*blkDl.fifoMu, defer_lock);
383     lock(endLk, fifoLk);
384     while (!(blkDl.end && blkDl.fifo.dirQueue->empty())) {
385         endLk.unlock();
386         fifoLk.unlock();
387
388         blkDl.sem->wait();
389
390         lock(endLk, fifoLk);
391         if (blkDl.end && blkDl.fifo.dirQueue->empty())
392             break;
393         endLk.unlock();
394
395         dirBlock_t block = blkDl.fifo.dirQueue->front();

```

```

396     blkDl.fifo.dirQueue->pop();
397     fifoLk.unlock();
398
399     //     if entry is find
400     if (block.entries.end() != block.entries.find(name)) {
401         blkDl.dlThPool->clear();
402         inodeIds = block.entries[name];
403
404         lock(endLk, fifoLk);
405         break;
406     }
407
408     lock(endLk, fifoLk);
409 }
410 endLk.unlock();
411 fifoLk.unlock();
412
413
414 //     TODO Doit etre un pointeur car lookup avant open donc inode doit etre en
415 //     memoire.
416 auto *inode = new internalInode_st();
417 inode->idents = inodeIds;
418 int ret;
419
420 if (inodeIds.empty()) {
421     fuse_reply_err(req, ENOENT);
422     return;
423 }
424
425 do {
426     ret = this->inodes->get(inodeIds.back(), &inode->inode, blockType::INODE);
427 } while (0 != ret);
428
429 if (inodeIds.empty()) {
430     break;
431 }
432
433 //     Send reply to fuse
434 if (0 != ret) {
435     fuse_reply_err(req, ret);
436 } else {
437     fuse_entry_param param = fuse_entry_param();
438     this->buildParam(*inode, param);
439     fuse_reply_entry(req, &param);
440 }
441
442 //     free all datas;
443 blkDl.dlThPool->wait();
444
445 delete (blkDl.dlThPool);
446 delete (blkDl.sem);
447 delete (blkDl.fifoMu);
448 delete (blkDl.endMu);
449
450 while (!blkDl.fifo.dirQueue->empty())
451     blkDl.fifo.dirQueue->pop();
452
453 delete (blkDl.fifo.dirQueue);
454 }

```



```

453
454 void Mtfs::mknod(fuse_req_t req, fuse_ino_t parent, const string &name,
455 mode_t mode, dev_t rdev) {
456
457     int ret;
458
459     internalInode_st *parentInode = this->getIntInode(parent);
460
461     // Create and write new inode
462     internalInode_st *inode = this->newInode(mode, fuse_req_ctx(req));
463     vector<ident_t> inodeIds;
464     if (0 != (ret = this->insertInode(inode->inode, inodeIds))) {
465         delete inode;
466         fuse_reply_err(req, ret);
467         return;
468     }
469
470     pool mPool(SIMULT_UP);
471     this->initMetas(*parentInode, inodeIds, blockType::INODE, &mPool);
472
473     // Add entry in dir
474     if (0 != (ret = this->addEntry(parentInode, name, inodeIds))) {
475         delete inode;
476         fuse_reply_err(req, ret);
477         return;
478     }
479
480     // reply to fuse
481     fuse_entry_param param{};
482     memset(&param, 0, sizeof(fuse_entry_param));
483
484     this->buildParam(*inode, param);
485
486     fuse_reply_entry(req, &param);
487 }
488
489 void Mtfs::mkdir(fuse_req_t req, fuse_ino_t ino, const string &name, mode_t
mode) {
490     int ret;
491     internalInode_st *parentInode = this->getIntInode(ino);
492
493     // Create and write new inode
494     internalInode_st *inode = this->newInode(mode, fuse_req_ctx(req));
495     inode->inode.accessRight |= S_IFDIR;
496     vector<ident_t> inodeIds;
497     if (0 != (ret = this->insertInode(inode->inode, inodeIds))) {
498         delete inode;
499         fuse_reply_err(req, ret);
500         return;
501     }
502
503     pool mPool(SIMULT_UP);
504     this->initMetas(*parentInode, inodeIds, blockType::INODE, &mPool);
505
506     // Add entry in dir
507     if (0 != (ret = this->addEntry(parentInode, name, inodeIds))) {
508         delete inode;
509         fuse_reply_err(req, ret);

```

```

510     return;
511 }
512
513 //      reply to fuse
514 fuse_entry_param param{};
515 memset(&param, 0, sizeof(fuse_entry_param));
516
517 this->buildParam(*inode, param);
518
519 fuse_reply_entry(req, &param);
520 }
521
522 void Mtf::unlink(fuse_req_t req, fuse_ino_t parent, const string &name) {
523     internalNode_st *parentInode = this->getIntInode(parent);
524
525     int ret;
526     if (SUCCESS != (ret = this->doUnlink(parentInode, name)))
527         fuse_reply_err(req, ret);
528     else
529         fuse_reply_err(req, SUCCESS);
530 }
531
532 void Mtf::rmdir(fuse_req_t req, fuse_ino_t parent, const string &name) {
533     (void) parent, name;
534
535     fuse_reply_err(req, ENOSYS);
536 }
537
538 void Mtf::access(fuse_req_t req, fuse_ino_t ino, int mask) {
539
540     const struct fuse_ctx *context = fuse_req_ctx(req);
541
542     if (FUSE_ROOT_ID == ino) {
543         int ret = EACCES;
544         inode_t inode = instance->getRootInode();
545         if (context->uid == inode.uid) {
546             if ((mask << 6 & inode.accesRight) != 0)
547                 ret = 0;
548             else {
549                 cerr << "[MTFS]: bad user right: " << (inode.accesRight & (mask << 6))
550 ) << endl;
551             }
552         } else {
553
554 //      get all user groups;
555         auto bufsize = (size_t) sysconf(_SC_GETPW_R_SIZE_MAX);
556         if (bufsize == -1)
557             bufsize = 16384;
558
559         auto *buf = (char *) malloc(bufsize * sizeof(char));
560         if (nullptr == buf) {
561             fuse_reply_err(req, ENOMEM);
562             free(buf);
563             return;
564         }
565
566         struct passwd pwd{}, *result;
567         int s = getpwuid_r(context->uid, &pwd, buf, bufsize, &result);
568         if (nullptr == result) {

```

```

568     int err;
569     if (0 == s)
570         err = EAGAIN;
571     else
572         err = s;
573     fuse_reply_err(req, err);
574     free(buf);
575     return;
576 }
577
578 int ngroups = GROUPS_TO_SEARCH;
579 auto *groups = (gid_t *) malloc(ngroups * sizeof(gid_t));
580 if (getgrouplist(pwd.pw_name, pwd.pw_gid, groups, &ngroups) == -1) {
581     cerr << "get group list error. ngoups " << ngroups << endl;
582     free(buf);
583     free(groups);
584 }
585
586 if (find(groups, groups + ngroups, inode.gid) != groups + ngroups) {
587     if ((mask << 3 & inode.accesRight) != 0)
588         ret = 0;
589 } else {
590     if ((mask & inode.accesRight) != 0)
591         ret = 0;
592 }
593
594 free(buf);
595 free(groups);
596 }
597
598 fuse_reply_err(req, ret);
599
600
601 } else {
602     fuse_reply_err(req, SUCCESS);
603 }
604 }
605
606 void Mtf::open(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
607     (void) ino;
608     // internalInode_st *inode = getIntInode(ino);
609
610     // if (!inode->inode.dataBlocks.empty()) {
611     //     uint8_t *firstBlock = (uint8_t *) malloc(this->blockSize * sizeof(
612     //         uint8_t));
613     //     this->blocks->get(inode->inode.dataBlocks.front().front(), &firstBlock,
614     //         blockType::DATA_BLOCK);
615     //     fi->fh = (uint64_t) firstBlock;
616     // } else
617     //     fi->fh = 0;A
618
619     auto *fd = new fd_st();
620     fd->blockT = fd->DATA;
621     fd->blkDl.dlThPool = new pool(this->SIMULT_DL);
622     fd->blkDl.sem = new Semaphore();
623     fd->blkDl.fifoMu = new mutex();
624     fd->blkDl.endMu = new mutex();
625     fd->blkDl.fifo.blkQueue = new queue<uint8_t *>();

```

```

625     fd->inoDl.dlThPool = new pool(this->SIMULT_UP);
626
627     fi->fh = (uint64_t) fd;
628
629     fuse_reply_open(req, fi);
630 }
631
632 void Mtf::release(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
633     (void) ino;
634     // internalInode_st *intIno = this->getIntInode(ino);
635
636     delete (uint8_t *) fi->fh;
637     fi->fh = 0;
638
639     fuse_reply_err(req, SUCCESS);
640 }
641
642 void Mtf::opendir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {
643
644     internalInode_st *intInode = this->getIntInode(ino);
645
646     auto *fd = new fd_st();
647     fd->blockT = fd->DIR;
648     fd->blkDl.dlThPool = new pool(this->SIMULT_DL);
649     fd->blkDl.sem = new Semaphore();
650     fd->blkDl.fifoMu = new mutex();
651     fd->blkDl.endMu = new mutex();
652     fd->blkDl.fifo.dirQueue = new queue<dirBlock_t>();
653
654     int i = 0;
655     for (auto &&blks: intInode->inode.dataBlocks) {
656
657         // bind(&Mtf::dlDirBlocks, this, blks, &blkQueue, &queueMutex, &semaphore
658         // ));
659         fd->blkDl.dlThPool->schedule(
660             bind(&Mtf::dlDirBlocks, this, blks, fd->blkDl.fifo.dirQueue, fd->
661             blkDl.fifoMu, fd->blkDl.sem));
662
663         if (this->INIT_DL <= i)
664             break;
665     }
666
667     fi->fh = (uint64_t) fd;
668
669     fuse_reply_open(req, fi);
670     // fuse_reply_err(req, ENOSYS);
671 }
672
673 void Mtf::readdir(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
674     fuse_file_info *fi) {
675     this->doReaddir(req, ino, size, off, fi, false);
676 }
677
678 void Mtf::readdirplus(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
679     fuse_file_info *fi) {
680     this->doReaddir(req, ino, size, off, fi, true);
681 }
682
683 void Mtf::releasedir(fuse_req_t req, fuse_ino_t ino, fuse_file_info *fi) {

```

```

680     (void) ino;
681
682     auto *fd = (fd_st *) fi->fh;
683     delete fd;
684     fi->fh = 0;
685
686     fuse_reply_err(req, SUCCESS);
687 }
688
689 void Mtfs::write(fuse_req_t req, fuse_ino_t ino, const char *buf, size_t size
, off_t off, fuse_file_info *fi) {
690     (void) ino, buf, size, off, fi;
691     //     internalNode_st *inode = this->getIntNode(ino);
692     //
693     //     int startBlock = (int) (off / this->blockSize);
694     //
695     //     fuse_reply_write(req, size);
696     fuse_reply_err(req, ENOSYS);
697 }
698
699 void Mtfs::write_buf(fuse_req_t req, fuse_ino_t ino, fuse_bufvec *bufv, off_t
off, fuse_file_info *fi) {
700     int ret = 0;
701
702     auto *fd = (fd_st *) fi->fh;
703
704     internalNode_st *inode = this->getIntNode(ino);
705
706     size_t size = bufv->buf[0].size;
707     auto firstBlock = (int) (off / this->blockSize);
708     auto lastBlock = (int) ((off + size - 1) / this->blockSize);
709
710     ruleInfo_t newBlockInfo = ruleInfo_st();
711     newBlockInfo.gid = inode->inode.gid;
712     newBlockInfo.uid = inode->inode.uid;
713     newBlockInfo.lastAccess = this->now();
714
715     size_t rem = size;
716     size_t write = 0;
717
718     pool metasTp(SIMULT_UP);
719
720     vector<uint64_t> toFree;
721
722     for (int blockToWrite = firstBlock; blockToWrite <= lastBlock; ++
blockToWrite) {
723         auto *block = (uint8_t *) calloc(sizeof(uint8_t), this->blockSize);
724         memset(block, 0, this->blockSize);
725
726         vector<ident_t> blockIdents;
727         if (inode->inode.dataBlocks.empty() || inode->inode.dataBlocks.size() <=
blockToWrite) {
728             if (SUCCESS !=
729                 (ret = this->blocks->add(newBlockInfo, blockIdents, blockType::
DATA_BLOCK,
730                                         (const int) this->redundancy))) {
731                 fuse_reply_err(req, ret);
732                 return;
733             }

```

```

734         this->initMetas(*inode, blockIds, blockType::DATA_BLOCK, &metasTp);
735
736         inode->inode.dataBlocks.push_back(blockIds);
737     } else {
738         blockIds = inode->inode.dataBlocks[blockToWrite];
739
740         for (auto &&id: blockIds) {
741             ret = this->blocks->get(id, block, blockType::DATA_BLOCK);
742             if (SUCCESS == ret)
743                 break;
744         }
745         if (SUCCESS != ret)
746             return (void) fuse_reply_err(req, ret);
747     }
748
749     size_t startPos = blockToWrite == firstBlock ? off % this->blockSize : 0;
750     size_t endPos = (blockToWrite == lastBlock ? rem % (this->blockSize + 1)
751 : this->blockSize) - startPos;
752
753     memcpy(block + startPos, bufv->buf->mem, endPos);
754     write += endPos;
755     rem -= endPos;
756
757     for (auto &&id: blockIds) {
758         fd->blkDl.dlThPool->schedule(bind(&Acces::put, this->blocks, id, block,
759         blockType::DATA_BLOCK));
760     }
761
762     toFree.push_back((uint64_t) block);
763 }
764
765 auto *data = (uint8_t *) malloc(size * sizeof(uint8_t));
766 memcpy(data, bufv->buf[0].mem, bufv->buf[0].size);
767
768 inode->inode.size = off + size;
769
770 // TODO supprimer les block en trop.
771 vector<ident_t> toDel;
772 if (inode->inode.dataBlocks.size() > lastBlock + 1) {
773     for (auto &dataBlock : inode->inode.dataBlocks) {
774         for (auto &&id: dataBlock) {
775             toDel.push_back(id);
776         }
777     }
778
779     inode->inode.dataBlocks.erase(inode->inode.dataBlocks.begin() + lastBlock
780 + 1,
781                                 inode->inode.dataBlocks.end());
782 }
783
784 for (auto &&id: inode->idents) {
785     fd->inoDl.dlThPool->schedule(bind(&Acces::put, this->inodes, id, &inode->
786 inode, blockType::INODE));
787 }
788
789 fd->blkDl.dlThPool->wait();
790 fd->inoDl.dlThPool->wait();

```

```

789     fuse_reply_write(req, write);
790
791     pool_delPool(this->redundancy);
792     for (auto &&id: toDel) {
793         delPool.schedule(bind(&Acces::del, this->blocks, id, blockType::
DATA_BLOCK));
794     }
795
796     for (auto &&ptr: toFree) {
797         free((uint8_t *) ptr);
798     }
799 }
800
801 void Mtfs::read(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
fuse_file_info *fi) {
802     (void) fi;
803     internalInode_st *inode = this->getIntInode(ino);
804
805
806     size = min(size, inode->inode.size);
807     int firstBlock, lastBlock;
808     firstBlock = (int) (off / this->blockSize);
809     lastBlock = (int) ((off + size - 1) / this->blockSize);
810
811     size_t rem = size;
812     size_t read = 0;
813
814     char *buf;
815     buf = new char[size];
816     char *p = buf;
817
818     for (int blockToRead = firstBlock; blockToRead <= lastBlock; blockToRead++)
819     {
820         uint8_t block[this->blockSize];
821         if (SUCCESS != this->blocks->get(inode->inode.dataBlocks[blockToRead][0],
block, blockType::DATA_BLOCK)) {
822             delete[] buf;
823             fuse_reply_err(req, EAGAIN);
824
825             return;
826         }
827
828         size_t startPos = blockToRead == firstBlock ? off % this->blockSize : 0;
829         size_t endPos = blockToRead == lastBlock ? rem % (this->blockSize + 1) :
this->blockSize;
830
831         size_t toCopy = endPos - startPos;
832         memcpy(p, block + startPos, toCopy);
833
834         rem -= toCopy;
835         read += toCopy;
836         p += toCopy;
837     }
838
839     fuse_reply_buf(req, buf, read);
840     delete[] buf;
841 }

```

```

842 //
843 ///////////////////////////////////////////////////
844 //                                     PRIVATE                                     //
845 ///////////////////////////////////////////////////
846
847 MtfS::MtfS() : redundancy(1), maxEntryPerBlock(ENTRY_PER_BLOCK) {
848     this->inodes = nullptr;
849     this->dirBlocks = nullptr;
850     this->blocks = nullptr;
851     this->rootIn = new internalInode_st();
852 }
853
854 /**
855  * Get the root inode
856  *
857  * @return the inode
858  */
859 inode_t MtfS::getRootInode() {
860     return this->rootIn->inode;
861 }
862
863 /**
864  * Build MTFS system
865  *
866  * @param superblock
867  * @return
868  */
869 bool MtfS::build(const superblock_t &superblock) {
870
871     this->redundancy = superblock.redundancy;
872     this->blockSize = superblock.blockSz;
873     this->rootIn->idents = superblock.rootInodes;
874
875     pluginSystem::PluginManager *manager = pluginSystem::PluginManager::
876     getInstance();
877
878     PoolManager *poolManager;
879     poolManager = new PoolManager();
880     for (auto &poolSt: superblock.pools) {
881         Pool *pool;
882         pool = new Pool(this->blockSize);
883
884         for (auto volSt: poolSt.second.volumes) {
885             volSt.second.params.insert(make_pair("home", MTFS_PLUGIN_HOME));
886             volSt.second.params.insert(make_pair("blockSize", to_string(this->
887             blockSize)));
888
889             pluginSystem::Plugin *plugin = manager->getPlugin(volSt.second.
890             pluginName);
891
892             if (!plugin->attach(volSt.second.params)) {
893                 cerr << "Failed attach plugin " << volSt.second.pluginName << endl;
894             }
895
896             pool->addVolume(volSt.first, new Volume(plugin), volSt.second.rule);
897         }
898     }
899 }

```



```

894     }
895
896     poolManager->addPool(poolSt.first , pool , poolSt.second.rule);
897 }
898
899 this->inodes = poolManager;
900 this->dirBlocks = poolManager;
901 this->blocks = poolManager;
902
903 return true;
904 }
905
906 /**
907  * Read the root inode
908  */
909 void Mtfs::readRootInode() {
910     string filename = string(MTFS_INSTALL_DIR) + "/" + systemName + "/root.json
911 ";
912     ifstream file(filename);
913     if (!file.is_open())
914         return;
915
916     IStreamWrapper wrapper(file);
917
918     Document d(kObjectType);
919     d.ParseStream(wrapper);
920
921     assert(d.HasMember(IN_MODE));
922     this->rootIn->inode.accesRight = d[IN_MODE].GetUint();
923
924     assert(d.HasMember(IN_LINKS));
925     this->rootIn->inode.linkCount = (uint8_t) d[IN_LINKS].GetUint();
926
927     assert(d.HasMember(IN_UID));
928     this->rootIn->inode.uid = d[IN_UID].GetUint();
929
930     assert(d.HasMember(IN_GID));
931     this->rootIn->inode.gid = d[IN_GID].GetUint();
932
933     assert(d.HasMember(IN_SIZE));
934     this->rootIn->inode.size = d[IN_SIZE].GetUint64();
935
936     assert(d.HasMember(IN_ACCESS));
937     this->rootIn->inode.atime = d[IN_ACCESS].GetUint64();
938
939     assert(d.HasMember(IN_BLOCKS));
940     for (auto &&item: d[IN_BLOCKS].GetArray()) {
941         vector<ident_t> blocksRedundancy;
942         for (auto &&block: item.GetArray()) {
943             ident_t ident = ident_t();
944
945             assert(block.HasMember(ID_POOL));
946             ident.poolId = block[ID_POOL].GetUint();
947
948             assert(block.HasMember(ID_VOLUME));
949             ident.volumeId = block[ID_VOLUME].GetUint();
950
951             assert(block.HasMember(ID_ID));
952             ident.id = block[ID_ID].GetUint64();

```

```

952         blocksRedundancy.push_back(ident);
953     }
954     this->rootIn->inode.dataBlocks.push_back(blocksRedundancy);
955 }
956 }
957 }
958
959 /**
960  * Write the root inode
961  */
962 void Mtfs::writeRootInode() {
963     Document d;
964
965     this->rootIn->inode.toJson(d);
966
967     StringBuffer sb;
968     PrettyWriter<StringBuffer> pw(sb);
969     d.Accept(pw);
970
971     string filename = string(MTFS_INSTALL_DIR) + "/" + systemName + "/root.json";
972     ofstream rootFile(filename);
973     rootFile << sb.GetString() << endl;
974     rootFile.close();
975 }
976
977 /**
978  * @brief Add entry in directory.
979  *
980  * This function add a new entry in directory and update parent inode if
981  * necessary.
982  *
983  * @param parentInode Directory to add entry
984  * @param name Name of entry
985  * @param inodeIds Inodes of entry
986  *
987  * @return 0 if success else errno.
988  */
989 int Mtfs::addEntry(internalInode_st *parentInode, std::string name, vector<
990 ident_t> &inodeIds) {
991     int ret = 0;
992
993     dirBlock_t dirBlock = dirBlock_t();
994
995     vector<ident_t> blockIds;
996
997     pool iPool(parentInode->idents.size());
998     pool mPool(SIMULT_UP);
999
1000     // if directory is empty add one block
1001     if (parentInode->inode.dataBlocks.empty()) {
1002         if (0 !=
1003             (ret = this->dirBlocks->add(getRuleInfo(parentInode->inode),
1004                                     blockIds, blockType::DIR_BLOCK,
1005                                     this->redundancy))) {
1006             return ret;
1007         }
1008     }
1009
1010     dirBlock.entries.clear();

```

```

1007     this->initMetas(*parentInode, blockIds, blockType::DIR_BLOCK, &mPool);
1008
1009     parentInode->inode.atime = this->now();
1010     parentInode->inode.dataBlocks.push_back(blockIds);
1011     for (auto &&ident: parentInode->idents) {
1012         iPool.schedule(bind(&Acces::put, this->inodes, ident, &parentInode->
1013 inode, blockType::INODE));
1014     }
1015     if (parentInode == this->rootIn)
1016         this->writeRootInode();
1017
1018 } else {
1019 //     else get the first block.
1020     blockIds = parentInode->inode.dataBlocks.back();
1021
1022     for (const auto &blockIdent : blockIds) {
1023         if (0 == (ret = this->dirBlocks->get(blockIdent, &dirBlock, blockType::
1024 DIR_BLOCK)))
1025             break;
1026     }
1027
1028     if (ret != SUCCESS)
1029         return ret;
1030 }
1031
1032 //     if block is full, allocate new block.
1033 if (this->maxEntryPerBlock == dirBlock.entries.size()) {
1034     ruleInfo_t info = getRuleInfo(parentInode->inode);
1035     info.lastAccess = this->now();
1036     if (0 != (ret = this->dirBlocks->add(info, blockIds, blockType::
1037 DIR_BLOCK, this->redundancy))) {
1038         return ret;
1039     }
1040
1041     this->initMetas(*parentInode, blockIds, blockType::DIR_BLOCK, &mPool);
1042
1043     dirBlock.entries.clear();
1044
1045     parentInode->inode.atime = this->now();
1046     parentInode->inode.dataBlocks.push_back(blockIds);
1047     for (auto &&ident: parentInode->idents) {
1048         iPool.schedule(bind(&Acces::put, this->inodes, ident, &parentInode->
1049 inode, blockType::INODE));
1050     }
1051     if (parentInode == this->rootIn)
1052         this->writeRootInode();
1053 }
1054
1055 //     write block
1056 dirBlock.entries.insert(make_pair(name, inodeIds));
1057 pool wpool(blockIds.size());
1058 for (auto &&ident: blockIds) {
1059 //     TODO try again if block not write.
1060     wpool.schedule(bind(&Acces::put, this->dirBlocks, ident, &dirBlock,
1061 blockType::DIR_BLOCK));
1062 }

```

```

1061     return 0;
1062 }
1063
1064 /**
1065  * @brief Insert inode in system.
1066  *
1067  * This function get free inode Id and put inode in volumes
1068  *
1069  * @param[in] inode Inode to put
1070  * @param[out] idents Ids to new inode
1071  *
1072  * @return 0 if SUCCESS else @see static const vars.
1073  */
1074 int Mtfs::insertInode(const inode_t &inode, vector<ident_t> &idents) {
1075     int ret;
1076
1077     // get new inode idents
1078     if (0 != (ret = this->inodes->add(this->getRuleInfo(inode), idents,
1079         blockType::INODE, this->redundancy))) {
1080         return ret;
1081     }
1082
1083     // write inode in volumes
1084     {
1085         pool thPool(idents.size());
1086         for (auto &ident: idents) {
1087             thPool.schedule(bind(&Acces::put, this->inodes, ident, &inode,
1088                 blockType::INODE));
1089         }
1090     }
1091     return ret;
1092 }
1093
1094 /**
1095  * Download all blocks in inode
1096  *
1097  * @param inode Inode who contains the blocks
1098  * @param dlSt Download struct
1099  * @param type Type of block to download DIR_BLOCK | DATA_BLOCK
1100  * @param firstBlockIdx Index of first block to download
1101  */
1102 void Mtfs::dlBlocks(const inode_t &inode, dl_st *dlSt, const blockType type,
1103     const int firstBlockIdx) {
1104     // TODO implements
1105
1106     if (firstBlockIdx < inode.dataBlocks.size()) {
1107         for (auto idents = inode.dataBlocks.begin() + firstBlockIdx; idents !=
1108             inode.dataBlocks.end(); idents++) {
1109             switch (type) {
1110                 case DIR_BLOCK:
1111                     dlSt->dlThPool->schedule(
1112                         bind(&Mtfs::dlDirBlocks, this, *idents, dlSt->fifo.dirQueue,
1113                             dlSt->fifoMu, dlSt->sem));
1114                     break;
1115                 case DATA_BLOCK:
1116                     break;
1117                 default:

```

```

1115 //      TODO Log error
1116         break;
1117     }
1118 }
1119 }
1120
1121 dlSt->dlThPool->wait();
1122
1123 unique_lock<mutex> endLk(*dlSt->endMu);
1124 dlSt->end = true;
1125 endLk.unlock();
1126 dlSt->sem->notify();
1127 }
1128
1129 /**
1130  * @brief download or get a directory block.
1131  *
1132  * This function is create for work in a other thread than worker.
1133  *
1134  * @param[in] ids    Block ids
1135  * @param[out] q    Queue
1136  * @param[in] queueMutex
1137  * @param[in] sem
1138  */
1139 void Mtf::dlDirBlocks(vector<ident_t> &ids, queue<dirBlock_t> *q, mutex *
queueMutex, Semaphore *sem) {
1140 // TODO gerer le cas ou aucun bloc n'a pu etre dl.
1141 int ret = 1;
1142 dirBlock_t db = dirBlock_t();
1143
1144 for (auto &&id: ids) {
1145     ret = this->dlDirBlocks->get(id, &db, blockType::DIR_BLOCK);
1146     if (SUCCESS == ret)
1147         break;
1148 }
1149
1150 if (SUCCESS == ret) {
1151     unique_lock<mutex> lk(*queueMutex);
1152     q->push(db);
1153     sem->notify();
1154 }
1155 }
1156
1157 /**
1158  * Download all inodes from dir entries.
1159  *
1160  * @param src dl dir entries struct
1161  * @param dst dl inode struct
1162  */
1163 void Mtf::dlInodes(dl_st *src, dl_st *dst) {
1164     pool inodesPool(this->SIMULT_DL);
1165     unique_lock<mutex> srcEndLk(*src->endMu, defer_lock);
1166     unique_lock<mutex> srcFifoLk(*src->fifoMu, defer_lock);
1167     std::lock(srcEndLk, srcFifoLk);
1168     while (!(src->end && src->fifo.dirQueue->empty())) {
1169         srcEndLk.unlock();
1170         srcFifoLk.unlock();
1171
1172         src->sem->wait();

```

```

1173     std::lock(srcEndLk, srcFifoLk);
1174     if (src->end && src->fifo.dirQueue->empty())
1175         break;
1176     srcEndLk.unlock();
1177
1178     dirBlock_t dirBlock = src->fifo.dirQueue->front();
1179     src->fifo.dirQueue->pop();
1180     srcFifoLk.unlock();
1181
1182     for (auto &&entry: dirBlock.entries) {
1183         inodesPool.schedule(
1184             bind(&Mtfs::dlInode, this, entry.second, dst->fifo.inodeQueue, dst
1185                 ->fifoMu, dst->sem,
1186                 entry.first));
1187     }
1188
1189     std::lock(srcEndLk, srcFifoLk);
1190 }
1191 srcEndLk.unlock();
1192 srcFifoLk.unlock();
1193
1194 inodesPool.wait();
1195
1196 unique_lock<mutex> dstEndLock(*dst->endMu);
1197 dst->end = true;
1198 dstEndLock.unlock();
1199
1200 dst->sem->notify();
1201 }
1202
1203 /**
1204  * Download inodes
1205  *
1206  * @param ids
1207  * @param queue
1208  * @param queueMutex
1209  * @param sem
1210  * @param key
1211  */
1212 void
1213 Mtfs::dlInode(vector<ident_t> &ids, queue<pair<string, inode_t>> *queue,
1214               mutex *queueMutex, Semaphore *sem,
1215               string &key) {
1216     int ret = 1;
1217     inode_t in = inode_t();
1218
1219     for (auto &&id: ids) {
1220         ret = this->inodes->get(id, &in, blockType::INODE);
1221         if (SUCCESS == ret)
1222             break;
1223     }
1224
1225     if (SUCCESS == ret) {
1226         unique_lock<mutex> lk(*queueMutex);
1227         queue->push(make_pair(key, in));
1228         sem->notify();
1229     }
1230 }

```

```

1230
1231 internalInode_st *Mtfs::getIntInode(fuse_ino_t ino) {
1232     if (FUSE_ROOT_ID == ino)
1233         return this->rootIn;
1234
1235     return (internalInode_st *) ino;
1236 }
1237
1238 void Mtfs::doReaddir(fuse_req_t req, fuse_ino_t ino, size_t size, off_t off,
1239     fuse_file_info *fi, const bool &plus) {
1240     (void) off;
1241     internalInode_st *inode = this->getIntInode(ino);
1242     fd_st *fd;
1243     fd = (fd_st *) fi->fh;
1244
1245     char *buf, *p;
1246     size_t rem, currSize;
1247
1248     buf = (char *) calloc(size, 1);
1249     if (nullptr == buf)
1250         return (void) fuse_reply_err(req, ENOMEM);
1251
1252     p = buf;
1253     rem = size;
1254     currSize = 0;
1255
1256     if (fd->firstCall) {
1257         fd->firstCall = false;
1258         size_t entrySize = this->dirBufAdd(req, p, currSize, ".", *inode, plus);
1259         p += entrySize;
1260         rem -= entrySize;
1261         currSize += entrySize;
1262
1263         // Dl other blocks and inodes
1264         int idx = this->INIT_DL;
1265         fd->blkThr = new boost::thread(
1266             bind(&Mtfs::dlBlocks, this, inode->inode, &fd->blkDl, blockType::
1267                 DIR_BLOCK, idx));
1268
1269         fd->inoDl.fifoMu = new mutex();
1270         fd->inoDl.endMu = new mutex();
1271         fd->inoDl.sem = new Semaphore();
1272         fd->inoDl.fifo.inodeQueue = new queue<pair<string, inode_t>>();
1273
1274         fd->inoThr = new boost::thread(bind(&Mtfs::dlInodes, this, &fd->blkDl, &
1275             fd->inoDl));
1276     }
1277
1278     unique_lock<mutex> inoEndLock(*fd->inoDl.endMu, defer_lock);
1279     unique_lock<mutex> inoFifoLock(*fd->inoDl.fifoMu, defer_lock);
1280     lock(inoEndLock, inoFifoLock);
1281     while (!(fd->inoDl.end && fd->inoDl.fifo.inodeQueue->empty())) {
1282         inoEndLock.unlock();
1283         inoFifoLock.unlock();
1284
1285         fd->inoDl.sem->wait();
1286
1287         internalInode_st *intInode;
1288         intInode = new internalInode_st();
1289     }

```

```

1286 //      shared_ptr<internalInode_st> intInode(new internalInode_st());
1287 lock(inoEndLock, inoFifoLock);
1288 if (fd->inoDl.end && fd->inoDl.fifo.inodeQueue->empty())
1289     break;
1290 inoEndLock.unlock();
1291
1292 string entryName = fd->inoDl.fifo.inodeQueue->front().first;
1293 intInode->inode = fd->inoDl.fifo.inodeQueue->front().second;
1294 fd->inoDl.fifo.inodeQueue->pop();
1295 inoFifoLock.unlock();
1296
1297
1298 size_t entrySize = this->dirBufAdd(req, p, currSize, entryName, *intInode
, plus);
1299
1300 if (entrySize > rem) {
1301     lock(inoEndLock, inoFifoLock);
1302     break;
1303 }
1304
1305 //      p += entrySize;
1306 rem -= entrySize;
1307 currSize += entrySize;
1308
1309 lock(inoEndLock, inoFifoLock);
1310 break;
1311 }
1312 inoEndLock.unlock();
1313 inoFifoLock.unlock();
1314 fuse_reply_buf(req, buf, size - rem);
1315 }
1316
1317 size_t Mtf::dirBufAdd(fuse_req_t &req, char *buf, size_t &currentSize, std::
string name, internalInode_st &inode,
1318                     const bool &plus) {
1319     size_t entSize = 0;
1320
1321     if (plus) {
1322         fuse_entry_param *p;
1323         p = new fuse_entry_param();
1324         buildParam(inode, *p);
1325
1326         entSize += fuse_add_dirent_plus(req, nullptr, 0, name.c_str(), nullptr,
0);
1327         fuse_add_dirent_plus(req, buf, entSize, name.c_str(), p, currentSize +
entSize);
1328     } else {
1329         struct stat stbuf{};
1330         memset(&stbuf, 0, sizeof(stbuf));
1331         buildStat(inode, stbuf);
1332
1333         entSize += fuse_add_dirent(req, nullptr, 0, name.c_str(), nullptr, 0);
1334         fuse_add_dirent(req, buf, entSize, name.c_str(), &stbuf, currentSize +
entSize);
1335     }
1336     return entSize;
1337 }
1338

```



```

1339 void Mtfs::buildParam(const internalInode_st &inode, fuse_entry_param &param)
1340 {
1341     param.ino = (fuse_ino_t) &inode;
1342     this->buildStat(inode, param.attr);
1343
1344     param.generation = 1;
1345     param.attr_timeout = 0.0;
1346     param.entry_timeout = 0.0;
1347 }
1348
1349 void Mtfs::buildStat(const internalInode_st &inode, struct stat &st) {
1350     st.st_dev = 0;
1351     st.st_ino = (__ino_t) &inode;
1352
1353     st.st_mode = inode.inode.accesRight;
1354     st.st_nlink = inode.inode.linkCount;
1355     st.st_uid = inode.inode.uid;
1356     st.st_gid = inode.inode.gid;
1357     st.st_size = inode.inode.size;
1358     time_t time = inode.inode.atime;
1359     st.st_atim.tv_sec = time;
1360     st.st_ctim.tv_sec = time;
1361     st.st_mtim.tv_sec = time;
1362     st.st_blksize = this->blockSize;
1363     st.st_blocks = inode.inode.dataBlocks.size();
1364 }
1365
1366 ruleInfo_t Mtfs::getRuleInfo(const inode_t &inode) {
1367     return {inode.uid, inode.gid, inode.atime};
1368 }
1369
1370 internalInode_st *Mtfs::newInode(const mode_t &mode, const fuse_ctx *ctx) {
1371     internalInode_st *inode;
1372     inode = new internalInode_st();
1373
1374     inode->inode.accesRight = mode;
1375     inode->inode.uid = ctx->uid;
1376     inode->inode.gid = ctx->gid;
1377     if ((mode & S_IFDIR) != 0)
1378         inode->inode.linkCount = 2;
1379
1380     return inode;
1381 }
1382
1383 uint64_t Mtfs::now() {
1384     return (uint64_t) time(nullptr);
1385 }
1386
1387 /**
1388  * Init metas blocs
1389  *
1390  * @param parentInode
1391  * @param ids
1392  * @param type
1393  * @param thPool
1394  */
1395 void Mtfs::initMetas(const internalInode_st &parentInode, const vector<
    ident_t> &ids, const blockType &type,

```

```

1396         boost::threadpool::pool *thPool) {
1397     blockInfo_t metas = blockInfo_t();
1398     metas.lastAccess = (uint64_t) time(nullptr);
1399     metas.referenceId = parentInode.idents;
1400
1401     for (auto &&id :ids) {
1402         metas.id = id;
1403
1404         switch (type) {
1405             case INODE:
1406                 if (nullptr != thPool)
1407                     thPool->schedule(bind(&Acces::putMetas, this->inodes, id, metas,
1408 type));
1409                 else
1410                     this->inodes->putMetas(id, metas, type);
1411                 break;
1412             case DIR_BLOCK:
1413                 if (nullptr != thPool)
1414                     thPool->schedule(bind(&Acces::putMetas, this->dirBlocks, id, metas,
1415 type));
1416                 else
1417                     this->dirBlocks->putMetas(id, metas, type);
1418                 break;
1419             case DATA_BLOCK:
1420                 if (nullptr != thPool)
1421                     thPool->schedule(bind(&Acces::putMetas, this->blocks, id, metas,
1422 type));
1423                 else
1424                     this->blocks->putMetas(id, metas, type);
1425                 break;
1426             default:
1427                 break;
1428         }
1429     }
1430 }
1431
1432 int Mtf::doUnlink(internalInode_st *parent, const std::string name) {
1433     for (auto &&dBlkIds :parent->inode.dataBlocks) {
1434         dirBlock_t dBlk = dirBlock_t();
1435
1436         this->dirBlocks->get(dBlkIds.front(), &dBlk, blockType::DIR_BLOCK);
1437
1438         if (dBlk.entries.end() != dBlk.entries.find(name))
1439             return this->delEntry(dBlkIds, dBlk, name);
1440     }
1441     return ENOENT;
1442 }
1443
1444 int Mtf::delEntry(std::vector<ident_t> &ids, dirBlock_t &blk, const std::
1445 string &name) {
1446     vector<ident_t> inodeIds = blk.entries.find(name)->second;
1447
1448     inode_t inode = inode_t();
1449     this->inodes->get(inodeIds.front(), &inode, blockType::INODE);
1450
1451     if (0 < (inode.accesRight & S_IFDIR))
1452         return EISDIR;

```

```

1451     blk.entries.erase(name);
1452     pool upPool(SIMULT_UP);
1453
1454     for (auto &&id :ids) {
1455         upPool.schedule(bind(&Acces::put, this->dirBlocks, id, &blk, blockType::
1456         DIR_BLOCK));
1457     }
1458
1459     for (auto &&dataBlkIds :inode.dataBlocks) {
1460         for (auto &&blkId :dataBlkIds) {
1461             this->blocks->del(blkId, blockType::DATA_BLOCK);
1462         }
1463     }
1464
1465     for (auto &&inodeId :inodeIds) {
1466         this->inodes->del(inodeId, blockType::INODE);
1467     }
1468
1469     return 0;
1470 }
1471
1472 } // namespace mtf

```

Listing 3.19 – "core/Mtfs.cpp"

3.4.4 Pool

Pool.h

```

1  /**
2   * \file Pool.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFs.
9
10  MTFs is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24  #ifndef FILESTORAGE_POOL_H
25  #define FILESTORAGE_POOL_H
26
27  #include <string>
28  #include <vector>

```

```

29 #include <mtfs/structs.h>
30 // #include <mtfs/Volume.h>
31 #include <mtfs/Rule.h>
32
33 namespace mtfs {
34     class Volume;
35
36     class Pool {
37
38     public:
39         static constexpr const char *POOLS = "pools";
40
41         // STATUS CODE
42         static const int SUCCESS = 0;
43         static const int VOLUME_ID_EXIST = 1;
44         static const int NO_VALID_VOLUME = 2;
45
46     public:
47         Pool(const size_t &blkSize);
48
49         virtual ~Pool();
50
51         static bool validate(const rapidjson::Value &pool);
52
53         static void
54         structToJson(const pool_t &pool, rapidjson::Value &dest, rapidjson::
55         Document::AllocatorType &allocator);
56
57         static void jsonToStruct(rapidjson::Value &src, pool_t &pool);
58
59         int addVolume(uint32_t volumeId, Volume *volume, Rule *rule);
60
61         int add(const ruleInfo_t &info, std::vector<ident_t> &idents, const
62         blockType &type, const int &nb = 1);
63
64         int del(const uint32_t &volumeId, const uint64_t &id, const blockType &type
65         );
66
67         int get(const uint32_t &volumeId, const uint64_t &id, void *data, const
68         blockType &type);
69
70         int put(const uint32_t &volumeId, const uint64_t &id, const void *data,
71         const blockType &type);
72
73         int getMetas(const uint32_t &volumeId, const uint64_t &id, blockInfo_t &
74         metas, const blockType &type);
75
76         int putMetas(const uint32_t &volumeId, const uint64_t &id, const
77         blockInfo_t &metas, const blockType &type);
78
79         void
80         doMigration(std::map<ident_t, ident_t> &movedBlk, std::vector<ident_t> &
81         unsatisfyBlk, const blockType &type);
82
83     private:
84         const size_t blockSize;
85
86     };
87
88 };

```

```

80     std::map<uint32_t, Volume *> volumes;
81     std::map<uint32_t, Rule *> rules;
82
83
84     int getValidVolumes(const ruleInfo_t &info, std::vector<uint32_t> &
85     volumeIds);
86 };
87 } // namespace mtfs
88 #endif

```

Listing 3.20 – "core/Pool.h"

Pool.cpp

```

1  /**
2   * \file Pool.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #include <utils/Logger.h>
25 #include "mtfs/Pool.h"
26
27 using namespace std;
28
29 namespace mtfs {
30     Pool::Pool(const size_t &blkSize) : blockSize(blkSize) {}
31
32     Pool::~~Pool() {
33         for (auto &&vol: this->volumes) {
34             delete vol.second;
35         }
36         for (auto &&rule: this->rules) {
37             delete rule.second;
38         }
39     }
40
41     bool Pool::validate(const rapidjson::Value &pool) {
42         if (!pool.HasMember(mtf::Volume::VOLUMES))
43             throw invalid_argument("Volumes missing!");
44         if (!pool[mtf::Volume::VOLUMES].IsObject())

```

```

45     throw invalid_argument("Volumes is not a object");
46
47     int migration = -1;
48     if (pool[mtfs::Volume::VOLUMES].MemberCount() <= 0)
49         throw invalid_argument("Number of volumes invalid!");
50     else if (pool[mtfs::Volume::VOLUMES].MemberCount() != 1) {
51         if (!pool.HasMember(Rule::MIGRATION))
52             throw invalid_argument("Migration missing!");
53
54         migration = pool[Rule::MIGRATION].GetInt();
55     }
56
57     for (auto &m: pool[mtfs::Volume::VOLUMES].GetObject()) {
58         if (Rule::rulesAreValid(migration, m.value) != Rule::VALID_RULES)
59             throw invalid_argument(string("Rules invalid for volume ") + m.name.
60 GetString() + " '!");
61
62         if (!Volume::validate(m.value))
63             throw invalid_argument(string("Volume ") + m.name.GetString() + "
64 invalid!");
65     }
66
67     return true;
68 }
69
70 void Pool::structToJson(const pool_t &pool, rapidjson::Value &dest, rapidjson
71 ::Document::AllocatorType &allocator) {
72     dest.AddMember(rapidjson::StringRef(Rule::MIGRATION), rapidjson::Value(pool
73 .migration), allocator);
74
75     pool.rule->toJson(dest, allocator);
76
77     rapidjson::Value volumes(rapidjson::kObjectType);
78     for (auto &&item : pool.volumes) {
79         rapidjson::Value volume(rapidjson::kObjectType);
80
81         Volume::structToJson(item.second, volume, allocator);
82
83         string id = to_string(item.first);
84         rapidjson::Value index(id.c_str(), (rapidjson::SizeType) id.size(),
85 allocator);
86         volumes.AddMember(index, volume, allocator);
87     }
88
89     dest.AddMember(rapidjson::StringRef(Volume::VOLUMES), volumes, allocator);
90 }
91
92 void Pool::jsonToStruct(rapidjson::Value &src, pool_t &pool) {
93     assert(src.HasMember(Rule::MIGRATION));
94     pool.migration = src[Rule::MIGRATION].GetInt();
95
96     assert(src.HasMember(Volume::VOLUMES));
97     for (auto &&item : src[Volume::VOLUMES].GetObject()) {
98         uint32_t id = (uint32_t) stoul(item.name.GetString());
99         volume_t volume;
100         memset(&volume, 0, sizeof(volume_t));
101         volume.params.clear();
102
103         volume.rule = Rule::buildRule(pool.migration, item.value);

```

```

99     Volume::jsonToStruct(item.value , volume);
100
101     pool.volumes.insert(make_pair(id , volume));
102 }
103
104 }
105
106
107 int Pool::addVolume(uint32_t volumeId , Volume *volume , Rule *rule) {
108     if (this->volumes.find(volumeId) != this->volumes.end() && this->rules.find
109         (volumeId) != this->rules.end())
110         return VOLUME_ID_EXIST;
111
112     this->volumes[volumeId] = volume;
113     this->rules[volumeId] = rule;
114
115     rule->configureStorage(volume);
116
117     return 0;
118 }
119
120 int Pool::add(const ruleInfo_t &info , std::vector<ident_t> &idents , const
121     blockType &type , const int &nb) {
122     int ret;
123     vector<uint32_t> volumeIds;
124     if (0 != (ret = this->getValidVolumes(info , volumeIds))) {
125         return ret;
126     }
127
128     if (0 == volumeIds.size())
129         return NO_VALID_VOLUME;
130
131     if (nb <= volumeIds.size()) {
132         for (int i = 0; i < nb; i++) {
133             uint64_t id = 0;
134             const uint32_t vid = volumeIds[i];
135
136             this->volumes[vid]->add(id , type);
137
138             idents.push_back(ident_t(id , vid));
139         }
140     } else {
141         int blkPerPool = (int) (nb / volumeIds.size());
142         int remainder = (int) (nb % volumeIds.size());
143         for (auto &&vid: volumeIds) {
144             vector<uint64_t> tmpIds;
145             int nbToAllocate = blkPerPool;
146
147             if (0 < remainder) {
148                 nbToAllocate++;
149                 remainder--;
150             }
151
152             this->volumes[vid]->add(tmpIds , nbToAllocate , type);
153
154             for (auto &&tmpId: tmpIds) {
155                 idents.push_back(ident_t(tmpId , vid));
156             }
157         }
158     }
159 }

```

```

156     }
157
158     return ret;
159 }
160
161 int Pool::del(const uint32_t &volumeId, const uint64_t &id, const blockType &
162   type) {
163     return this->volumes[volumeId]->del(id, type);
164 }
165
166 int Pool::get(const uint32_t &volumeId, const uint64_t &id, void *data, const
167   blockType &type) {
168     return this->volumes[volumeId]->get(id, data, type);
169 }
170
171 int Pool::put(const uint32_t &volumeId, const uint64_t &id, const void *data,
172   const blockType &type) {
173     return this->volumes[volumeId]->put(id, data, type);
174 }
175
176 int Pool::getMetas(const uint32_t &volumeId, const uint64_t &id, blockInfo_t
177   &metas, const blockType &type) {
178     return this->volumes[volumeId]->getMetas(id, metas, type);
179 }
180
181 int Pool::putMetas(const uint32_t &volumeId, const uint64_t &id, const
182   blockInfo_t &metas, const blockType &type) {
183     return this->volumes[volumeId]->putMetas(id, metas, type);
184 }
185
186 int Pool::getValidVolumes(const ruleInfo_t &info, vector<uint32_t> &volumeIds
187   ) {
188     for (auto &&item: this->rules) {
189         if (item.second->satisfyRules(info))
190             volumeIds.push_back(item.first);
191     }
192     return 0;
193 }
194
195 void
196 Pool::doMigration(std::map<ident_t, ident_t> &movedBlk, std::vector<ident_t>
197   &unsatisfyBlk, const blockType &type) {
198     for (auto &&volume : this->volumes) {
199         Logger::getInstance()->log("Pool.doMigration", "do migration for volume:
200 " + to_string(volume.first),
201           Logger::L_DEBUG);
202         vector<blockInfo_t> unsatisfy;
203         unsatisfy.clear();
204         volume.second->getUnsatisfy(unsatisfy, type);
205
206         for (auto &&blkInfos : unsatisfy) {
207             ident_t oldIdent(blkInfos.id.id, volume.first);
208             blkInfos.id.volumeId = volume.first;
209
210             ruleInfo_t info = ruleInfo_t();
211             info.lastAccess = blkInfos.lastAccess;
212
213             vector<uint32_t> newVolumes;
214             this->getValidVolumes(info, newVolumes);

```



```

207     ident_t newIdent(0, newVolumes.front());
208     this->volumes[newIdent.volumeId]->add(newIdent.id, type);
209
210     if (oldIdent.volumeId == newIdent.volumeId) {
211         continue;
212     }
213
214     void *datas = nullptr;
215     switch (type) {
216         case INODE:
217             datas = new inode_t();
218             break;
219         case DIR_BLOCK:
220             datas = new dirBlock_t();
221             break;
222         case DATA_BLOCK:
223             datas = new uint8_t[this->blockSize];
224             break;
225         default:
226             continue;
227     }
228     this->volumes[oldIdent.volumeId]->get(oldIdent.id, datas, type);
229
230     this->volumes[newIdent.volumeId]->put(newIdent.id, datas, type);
231     this->volumes[newIdent.volumeId]->putMetas(newIdent.id, blkInfos, type)
232 ;
233
234     movedBlk.emplace(oldIdent, newIdent);
235
236 //     this->volumes[oldIdent.volumeId]->del(oldIdent.id, type);
237
238     switch (type) {
239         case INODE:
240             delete ((inode_t *) datas);
241             break;
242         case DIR_BLOCK:
243             delete ((dirBlock_t *) datas);
244             break;
245         case DATA_BLOCK:
246             delete []((uint8_t *) datas);
247             break;
248         default:
249             break;
250     }
251 }
252
253 unsigned long nb = unsatisfy.size();
254
255 if (0 != nb)
256     Logger::getInstance()->log("Pool.doMigration", "Block need to move: " +
257 to_string(nb), Logger::L_DEBUG);
258 }
259
260 } // namespace mtf
261

```

Listing 3.21 – "core/Pool.cpp"

3.4.5 PoolManager

PoolManager.h

```

1  /**
2   * \file PoolManager.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24  #ifndef FILESTORAGE_POOL_MANAGER_H
25  #define FILESTORAGE_POOL_MANAGER_H
26
27  #include <string>
28  #include <vector>
29  #include <list>
30  #include <set>
31  #include <iostream>
32  #include <mutex>
33  #include <assert.h>
34
35  #include "mtfs/Acces.h"
36  #include "mtfs/Pool.h"
37
38  namespace mtfs {
39      class Pool;
40
41      class Acces;
42
43      class PoolManager : public Acces {
44
45      public:
46          static const int SUCCESS = 0;
47          static const int POOL_ID_EXIST = 1;
48          static const int NO_VALID_POOL = 2;
49          static const int IS_LOCKED = 3;
50
51      private:
52
53          std::map<uint32_t, Pool *> pools;
54          std::map<uint32_t, Rule *> rules;
55
56          std::mutex inodeMutex;

```

```

57     std::set<ident_t> lockedInodes;
58     std::recursive_mutex inodeTransMutex;
59     std::map<ident_t, ident_t> inodeTranslateMap;
60
61     std::mutex dirMutex;
62     std::set<ident_t> lockedDirBlock;
63     std::recursive_mutex dirTransMutex;
64     std::map<ident_t, ident_t> dirBlockTranslateMap;
65
66     std::mutex blockMutex;
67     std::set<ident_t> lockedBlocks;
68     std::recursive_mutex blockTransMutex;
69     std::map<ident_t, ident_t> blockTranslateMap;
70
71
72 public:
73     virtual ~PoolManager();
74
75     int addPool(uint32_t poolId, Pool *pool, Rule *rule);
76
77     int add(const ruleInfo_t &info, std::vector<ident_t> &ids, const blockType
type, const size_t nb) override;
78
79     int del(const ident_t &id, const blockType type) override;
80
81     int get(const ident_t &id, void *data, const blockType type) override;
82
83     int put(const ident_t &id, const void *data, const blockType type) override
;
84
85     int getMetas(const ident_t &id, blockInfo_t &metas, const blockType type)
override;
86
87     int putMetas(const ident_t &id, const blockInfo_t &metas, const blockType
type) override;
88
89     void doMigration(const blockType type);
90
91 private:
92
93     bool isLocked(const ident_t &id, const blockType &type);
94
95     bool lock(const ident_t &id, const blockType &type);
96
97     bool unlock(const ident_t &id, const blockType &type);
98
99     bool hasMoved(const ident_t &id, ident_t &newId, const blockType &type);
100
101     int getValidPools(const ruleInfo_t &info, std::vector<uint32_t> &poolIds);
102
103     void dumpTranslateMap(const int &nb, const blockType &type);
104
105     const int moveBlk(const ident_t &old, const ident_t &cur, const blockType &
type);
106 };
107
108 } // namespace mtfs
109 #endif

```

Listing 3.22 – "core/PoolManager.h"

PoolManager.cpp

```

1  /**
2   * \file PoolManager.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #include <mtfs/PoolManager.h>
25 #include <utils/Logger.h>
26
27 namespace mtfs {
28     using namespace std;
29
30     PoolManager::~PoolManager() {
31         // dumpTranslateMap(INT_MAX, INODE);
32         // dumpTranslateMap(INT_MAX, DIR_BLOCK);
33         // dumpTranslateMap(INT_MAX, DATA_BLOCK);
34         for (auto &&pool: this->pools) {
35             delete pool.second;
36         }
37
38         for (auto &&rule: this->rules) {
39             delete rule.second;
40         }
41     }
42
43     int PoolManager::addPool(uint32_t poolId, Pool *pool, Rule *rule) {
44         if (pools.find(poolId) != pools.end())
45             return POOL_ID_EXIST;
46
47         pools[poolId] = pool;
48         rules[poolId] = rule;
49
50         return SUCCESS;
51     }
52

```

```

53 int PoolManager::add(const ruleInfo_t &info, std::vector<ident_t> &ids, const
    blockType type,
54         const size_t nb) {
55     int ret;
56     vector<uint32_t> poolIds;
57     if (0 != (ret = this->getValidPools(info, poolIds))) {
58         return ret;
59     }
60     vector<ident_t> tmpIdent;
61
62     if (nb <= poolIds.size()) {
63         for (auto &&id: poolIds) {
64             tmpIdent.clear();
65
66             this->pools[id]->add(info, tmpIdent, type);
67
68             for (auto &&ident: tmpIdent) {
69                 ident.poolId = id;
70                 ids.push_back(ident);
71             }
72         }
73     } else {
74         int blkPerPool = (int) (nb / poolIds.size());
75         int remainder = (int) (nb % poolIds.size());
76         for (auto &&id: poolIds) {
77             tmpIdent.clear();
78
79             int nbToAlloc = blkPerPool;
80             if (0 < remainder) {
81                 nbToAlloc++;
82                 remainder--;
83             }
84
85             this->pools[id]->add(info, tmpIdent, type, nbToAlloc);
86
87             for (auto &&ident: tmpIdent) {
88                 ident.poolId = id;
89                 ids.push_back(ident);
90             }
91         }
92     }
93
94     return ret;
95 }
96
97 int PoolManager::del(const ident_t &id, const blockType type) {
98     int ret = IS_LOCKED;
99
100     ident_t newId = id;
101     this->hasMoved(id, newId, type);
102
103     if (!this->isLocked(newId, type))
104         ret = this->pools[newId.poolId]->del(newId.volumeId, newId.id, type);
105
106     return ret;
107 }
108
109 int PoolManager::get(const ident_t &id, void *data, const blockType type) {
110     int ret = IS_LOCKED;

```

```

111     ident_t newId = id;
112     this->hasMoved(id, newId, type);
113
114     if (!this->isLocked(newId, type)) {
115         string message = "get ";
116
117         switch (type) {
118             case INODE:
119                 message += "inode ";
120                 break;
121             case DIR_BLOCK:
122                 message += "dir block ";
123                 break;
124             case DATA_BLOCK:
125                 message += "dat block ";
126                 break;
127             case SUPERBLOCK:
128                 message += "superblock ";
129                 break;
130         }
131
132         message += newId.toString();
133
134         Logger::getInstance()->log("POOL_MANAGER", message, Logger::L_DEBUG);
135         ret = this->pools[newId.poolId]->get(newId.volumeId, newId.id, data, type
136     );
137     }
138
139     return ret;
140 }
141
142 int PoolManager::put(const ident_t &id, const void *data, const blockType
143 type) {
144     int ret = IS_LOCKED;
145
146     ident_t newId = id;
147     this->hasMoved(id, newId, type);
148
149     if (this->pools.end() == this->pools.find(newId.poolId))
150         throw out_of_range("Invalid pool id " + newId.poolId);
151
152     if (this->lock(newId, type)) {
153         ret = this->pools[newId.poolId]->put(newId.volumeId, newId.id, data, type
154     );
155         this->unlock(newId, type);
156     }
157
158     return ret;
159 }
160
161 int PoolManager::getMetas(const ident_t &id, blockInfo_t &metas, const
162 blockType type) {
163     int ret = IS_LOCKED;
164
165     ident_t newId = id;
166     this->hasMoved(id, newId, type);
167
168     if (!this->isLocked(newId, type))

```

```

166     ret = this->pools[newId.poolId]->getMetas(newId.volumeId, newId.id, metas
167     , type);
168     return ret;
169 }
170
171 int PoolManager::putMetas(const ident_t &id, const blockInfo_t &metas, const
172     blockType type) {
173     int ret = IS_LOCKED;
174
175     ident_t newId = id;
176     this->hasMoved(id, newId, type);
177
178     if (this->pools.end() == this->pools.find(newId.poolId))
179         throw out_of_range("Invalid pool id " + newId.poolId);
180
181     if (this->lock(newId, type)) {
182         ret = this->pools[newId.poolId]->putMetas(newId.volumeId, newId.id, metas
183         , type);
184         this->unlock(newId, type);
185     }
186
187     return ret;
188 }
189
190 void PoolManager::doMigration(const blockType type) {
191     vector<ident_t> unsatisfyBlk;
192
193     // this->dumpTranslateMap(10, type);
194
195     for (auto &&pool: this->pools) {
196         map<ident_t, ident_t> tmpMovedBlk;
197         pool.second->doMigration(tmpMovedBlk, unsatisfyBlk, type);
198         for (auto &&item :tmpMovedBlk) {
199             ident_t key = item.first;
200             ident_t val = item.second;
201
202             key.poolId = pool.first;
203             val.poolId = pool.first;
204             recursive_mutex *mu = nullptr;
205             map<ident_t, ident_t> *map = nullptr;
206             switch (type) {
207                 case INODE:
208                     mu = &this->inodeTransMutex;
209                     map = &this->inodeTranslateMap;
210                     break;
211                 case DIR_BLOCK:
212                     mu = &this->dirTransMutex;
213                     map = &this->dirBlockTranslateMap;
214                     break;
215                 case DATA_BLOCK:
216                     mu = &this->blockTransMutex;
217                     map = &this->blockTranslateMap;
218                     break;
219                 case SUPERBLOCK:
220                     continue;
221                     break;
222             }
223             unique_lock<recursive_mutex> lk(*mu);

```

```

222     map->emplace(key, val);
223 }
224
225     Logger::getInstance()->log("Poolmanager", "endMig", Logger::level::
L_DEBUG);
226 }
227 }
228
229 bool PoolManager::isLocked(const ident_t &id, const blockType &type) {
230     mutex *mu = nullptr;
231     set<ident_t> *lSet = nullptr;
232
233     switch (type) {
234     case INODE:
235         mu = &this->inodeMutex;
236         lSet = &this->lockedInodes;
237         break;
238     case DIR_BLOCK:
239         mu = &this->dirMutex;
240         lSet = &this->lockedDirBlock;
241         break;
242     case DATA_BLOCK:
243         mu = &this->blockMutex;
244         lSet = &this->lockedBlocks;
245         break;
246     }
247
248     if (nullptr == mu || nullptr == lSet)
249         return false;
250
251     unique_lock<mutex> lk(*mu);
252     return lSet->find(id) != lSet->end();
253 }
254
255 bool PoolManager::lock(const ident_t &id, const blockType &type) {
256     mutex *mu = nullptr;
257     set<ident_t> *lSet = nullptr;
258
259     switch (type) {
260     case INODE:
261         mu = &this->inodeMutex;
262         lSet = &this->lockedInodes;
263         break;
264     case DIR_BLOCK:
265         mu = &this->dirMutex;
266         lSet = &this->lockedDirBlock;
267         break;
268     case DATA_BLOCK:
269         mu = &this->blockMutex;
270         lSet = &this->lockedBlocks;
271         break;
272     default:
273         // TODO log error
274         return false;
275         break;
276     }
277
278     unique_lock<mutex> lk(*mu);
279     if (lSet->find(id) != lSet->end())

```



```

280     return false;
281
282     lSet->insert(id);
283
284     return true;
285 }
286
287 bool PoolManager::unlock(const ident_t &id, const blockType &type) {
288     mutex *mu = nullptr;
289     set<ident_t> *lSet = nullptr;
290
291     switch (type) {
292     case INODE:
293         mu = &this->inodeMutex;
294         lSet = &this->lockedInodes;
295         break;
296     case DIR_BLOCK:
297         mu = &this->dirMutex;
298         lSet = &this->lockedDirBlock;
299         break;
300     case DATA_BLOCK:
301         mu = &this->blockMutex;
302         lSet = &this->lockedBlocks;
303         break;
304     default:
305         // TODO log error
306         return false;
307         break;
308     }
309
310     unique_lock<mutex> lk(*mu);
311     if (lSet->find(id) == lSet->end())
312         return false;
313
314     lSet->erase(id);
315
316     return true;
317 }
318
319 bool PoolManager::hasMoved(const ident_t &id, ident_t &newId, const blockType
&type) {
320     recursive_mutex *mu = nullptr;
321     map<ident_t, ident_t> *transMap = nullptr;
322
323     switch (type) {
324     case INODE:
325         mu = &this->inodeTransMutex;
326         transMap = &this->inodeTranslateMap;
327         break;
328     case DIR_BLOCK:
329         mu = &this->dirTransMutex;
330         transMap = &this->dirBlockTranslateMap;
331         break;
332     case DATA_BLOCK:
333         mu = &this->blockTransMutex;
334         transMap = &this->blockTranslateMap;
335         break;
336     default:
337         return false;

```

```

338     }
339
340     assert(mu != nullptr);
341     assert(transMap != nullptr);
342
343     unique_lock<recursive_mutex> lk(*mu);
344     if (transMap->find(id) != transMap->end()) {
345         newId = (*transMap)[id];
346         return true;
347     }
348
349     return false;
350 }
351
352 int PoolManager::getValidPools(const ruleInfo_t &info, std::vector<uint32_t>
    &poolIds) {
353     for (auto &&item: this->rules) {
354         if (item.second->satisfyRules(info))
355             poolIds.push_back(item.first);
356     }
357
358     return SUCCESS;
359 }
360
361 void PoolManager::dumpTranslateMap(const int &nb, const blockType &type) {
362
363     recursive_mutex *mu;
364     map<ident_t, ident_t> *tm;
365     switch (type) {
366     case INODE:
367         mu = &this->inodeTransMutex;
368         tm = &this->inodeTranslateMap;
369         break;
370     case DIR_BLOCK:
371         mu = &this->dirTransMutex;
372         tm = &this->dirBlockTranslateMap;
373         break;
374     case DATA_BLOCK:
375         mu = &this->blockTransMutex;
376         tm = &this->blockTranslateMap;
377         break;
378     default:
379         return;
380     }
381
382     unique_lock<recursive_mutex> lk(*mu);
383     int i = 0;
384     for (auto &&item : *tm) {
385         if (nb == i)
386             break;
387         if (SUCCESS == this->moveBlk(item.first, item.second, type)) {
388             tm->erase(item.first);
389         }
390         i++;
391     }
392 }
393
394 const int PoolManager::moveBlk(const ident_t &old, const ident_t &cur, const
    blockType &type) {

```

```

395     blockInfo_t metas{};
396     this->pools[cur.poolId]->getMetas(cur.volumeId, cur.id, metas, type);
397     dirBlock_t dirBlock{};
398
399     switch (type) {
400     case INODE:
401         this->get(metas.referenceId.front(), &dirBlock, DIR_BLOCK);
402         for (auto &&item : dirBlock.entries) {
403             for (auto &i : item.second) {
404                 if (old == i)
405                     i = cur;
406             }
407         }
408         for (auto &&ref : metas.referenceId) {
409             this->put(ref, &dirBlock, DIR_BLOCK);
410         }
411         break;
412     case DIR_BLOCK:
413     case DATA_BLOCK:
414         inode_t inode{};
415         this->get(metas.referenceId.front(), &inode, INODE);
416
417         for (auto &&blks : inode.dataBlocks) {
418             for (auto &blk : blks) {
419                 if (old == blk)
420                     blk = cur;
421             }
422         }
423         for (auto &&ref : metas.referenceId) {
424             this->put(ref, &inode, INODE);
425         }
426         break;
427     }
428
429     return SUCCESS;
430 }
431
432 } // namespace mtfs

```

Listing 3.23 – "core/PoolManager.cpp"

3.4.6 Rule

Rule.h

```

1  /**
2   * \file Rule.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MIFS.
9
10  MIFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.

```

```

14
15     FooBar is distributed in the hope that it will be useful ,
16     but WITHOUT ANY WARRANTY; without even the implied warranty of
17     MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.  See the
18     GNU General Public License for more details .
19
20     You should have received a copy of the GNU General Public License
21     along with FooBar.  If not, see <http://www.gnu.org/licenses/>.
22 */
23
24 #ifndef FILESTORAGE_RULE_H
25 #define FILESTORAGE_RULE_H
26
27 #include <string>
28 #include <vector>
29 #include <list>
30 #include <iostream>
31 #include <assert.h>
32 #include <rapidjson/schema.h>
33 #include <mtfs/structs.h>
34 #include <mtfs/Volume.h>
35
36 /**
37  *
38  */
39 namespace mtfs {
40     class Pool;
41
42     class Volume;
43
44     /**
45      * @interface
46      */
47     class Rule {
48     public:
49         static constexpr const char *MIGRATION = "migration";    ///< Migration
50                             name in the config file
51
52         static const int NO_MIGRATION = -1;    ///< No migration value
53         static const int TIME_MIGRATION = 0;    ///< Time migration value
54         static const int RIGHT_MIGRATION = 1;    ///< User right migration
55
56         static const int SUCCESS = 0;    ///< Success code.
57         static const int VALID_RULES = 0;    ///< Valid rule code.
58         static const int INVALID_RULES = -1;    ///< Invalid rule code.
59         static const int UNKNOW_MIGRATION = -2;    ///< Unknow migration code.
60     public:
61
62         /**
63          * @brief Copy JSON config .
64          *
65          * @param migration    Migration type.
66          * @param source        JSON source config.
67          * @param destination    JSON destination config.
68          * @param allocator    Allocator for destination.
69          *
70          * @return SUCCESS if no error or corresponding code.
71          */

```

```

71     static int copyConfig(int migration, rapidjson::Document &source, rapidjson
::Value &destination,
72         rapidjson::Document::AllocatorType &allocator);
73
74     /**
75      * @brief Check if rule config is valid.
76      *
77      * @param migration Migration type.
78      * @param value Json config.
79      * @return VALID_RULES or INVALID_RULES
80      */
81     static int rulesAreValid(int migration, const rapidjson::Value &value);
82
83     /**
84      * @brief Build a Rule object.
85      *
86      * @param migration Migration type.
87      * @param value Json Config.
88      *
89      * @return Pointer on new Rule or nullptr if fail.
90      */
91     static Rule *buildRule(int migration, const rapidjson::Value &value);
92
93     /**
94      * Dump object to Json.
95      *
96      * @param json
97      * @param allocator
98      * @return
99      */
100    virtual bool toJson(rapidjson::Value &json, rapidjson::Document::
AllocatorType &allocator)=0;
101
102    /**
103     * @brief Check if the block or inode satisfy rules.
104     * @param info Info of Block/Inode
105     *
106     * @return true or false.
107     */
108    virtual bool satisfyRules(mtrfs::ruleInfo_t info)=0;
109
110    /**
111     * Configure volume storage
112     * @param volume
113     * @return
114     */
115    virtual int configureStorage(Volume *volume);
116
117    /**
118     * Configure pool storage
119     * @param volume
120     * @return
121     */
122    virtual int configureStorage(Pool *pool);
123 };
124
125 } // namespace mtrfs
126 #endif

```

Listing 3.24 – "core/Rule.h"

Rule.cpp

```

1  /**
2   * \file Rule.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #include <mtfs/TimeRule.h>
25 #include <mtfs/UserRightRule.h>
26 #include <pwd.h>
27
28 namespace mtfs {
29
30  int Rule::copyConfig(int migration, rapidjson::Document &source, rapidjson::
31  Value &destination,
32  rapidjson::Document::AllocatorType &allocator) {
33  switch (migration) {
34  case TIME_MIGRATION:
35  return TimeRule::copyConfig(source, destination, allocator);
36  case RIGHT_MIGRATION:
37  return UserRightRule::copyConfig(source, destination, allocator);
38  default:
39  return UNKNOWN_MIGRATION;
40  }
41
42  int Rule::rulesAreValid(int migration, const rapidjson::Value &value) {
43  switch (migration) {
44  case NO_MIGRATION:
45  return SUCCESS;
46  case TIME_MIGRATION:
47  return TimeRule::rulesAreValid(value);
48  case RIGHT_MIGRATION:
49  return UserRightRule::rulesAreValid(value);
50  default:
51  return UNKNOWN_MIGRATION;
52  }

```

```

53 }
54
55 Rule *Rule::buildRule(int migration, const rapidjson::Value &value) {
56     Rule *rule = nullptr;
57     if (migration == NO_MIGRATION)
58         rule = nullptr;
59     else if (migration == TIME_MIGRATION) {
60         uint64_t ll = 0, hl = 0;
61
62         if (value.HasMember(TimeRule::TIME_LOW_LIMIT))
63             ll = value[TimeRule::TIME_LOW_LIMIT].GetUint64();
64
65         if (value.HasMember(TimeRule::TIME_HIGH_LIMIT))
66             hl = value[TimeRule::TIME_HIGH_LIMIT].GetUint64();
67
68         rule = new TimeRule(ll, hl);
69     } else if (migration == RIGHT_MIGRATION) {
70         UserRightRule *uRule = new UserRightRule();
71
72         if (value.HasMember(UserRightRule::ALLOW_USER))
73             for (auto &ua: value[UserRightRule::ALLOW_USER].GetArray())
74                 uRule->addAllowUid(getpwnam(ua.GetString())->pw_uid);
75
76         if (value.HasMember(UserRightRule::DENY_USER))
77             for (auto &ud: value[UserRightRule::DENY_USER].GetArray())
78                 uRule->addDenyUid(getpwnam(ud.GetString())->pw_uid);
79
80         if (value.HasMember(UserRightRule::ALLOW_GROUP))
81             for (auto &ga: value[UserRightRule::ALLOW_GROUP].GetArray())
82                 uRule->addAllowGid(getpwnam(ga.GetString())->pw_gid);
83
84         if (value.HasMember(UserRightRule::DENY_GROUP))
85             for (auto &ga: value[UserRightRule::DENY_GROUP].GetArray())
86                 uRule->addDenyGid(getpwnam(ga.GetString())->pw_gid);
87
88         rule = uRule;
89     }
90     return rule;
91 }
92
93 int Rule::configureStorage(Volume *volume) {
94     volume->setIsTimeVolume(false);
95
96     return ENOSYS;
97 }
98
99 int Rule::configureStorage(Pool *pool) {
100     (void) pool;
101
102     return ENOSYS;
103 }
104
105 }

```

Listing 3.25 – "core/Rule.cpp"

3.4.7 TimeRule

TimeRule.h

```

1  /**
2   * \file TimeRule.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  FooBar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with FooBar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24  #ifndef FILESTORAGE_TIME_RULE_H
25  #define FILESTORAGE_TIME_RULE_H
26
27  #include <string>
28  #include <vector>
29  #include <list>
30  #include <iostream>
31  #include <assert.h>
32
33  #include <mtfs/Rule.h>
34
35  namespace mtfs {
36  class TimeRule : public Rule {
37  public:
38      static constexpr const char *TIME_LOW_LIMIT = "lowLimit";
39      static constexpr const char *TIME_HIGH_LIMIT = "highLimit";
40
41  private:
42      uint64_t lowerLimit;
43      uint64_t higerLimit;
44
45  public:
46      static int copyConfig(rapidjson::Document &source, rapidjson::Value &
47          destination,
48          rapidjson::Document::AllocatorType &allocator);
49
50      static int rulesAreValid(const rapidjson::Value &value);
51
52      TimeRule(uint64_t lowerLimit, uint64_t higerLimit);
53
54      bool satisfyRules(ruleInfo_st info) override;

```



```

55     bool toJson(rapidjson::Value &json, rapidjson::Document::AllocatorType &
56     allocator) override;
57
58     int configureStorage(Volume *volume) override;
59
60     int configureStorage(Pool *pool) override;
61
62 private:
63     uint64_t now();
64 };
65 } // namespace mtfs
66 #endif

```

Listing 3.26 – "core/TimeRule.h"

TimeRule.cpp

```

1  /**
2   * \file TimeRule.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #include "mtfs/TimeRule.h"
25
26 namespace mtfs {
27     TimeRule::TimeRule(uint64_t lowerLimit, uint64_t higerLimit) : lowerLimit(
28         lowerLimit),
29
30         higerLimit(higerLimit) {}
31
32     bool TimeRule::satisfyRules(ruleInfo_st info) {
33         uint64_t now = this->now();
34         uint64_t delay = now - info.lastAccess;
35
36         if (0 == this->lowerLimit && 0 == this->higerLimit)
37             return true;
38         else if (0 == this->lowerLimit)
39             return delay <= this->higerLimit;
40         else if (0 == this->higerLimit)
41             return delay >= this->lowerLimit;
42         else

```

```

41     return this->lowerLimit <= delay <= this->higerLimit;
42 }
43
44 int TimeRule::copyConfig(rapidjson::Document &source, rapidjson::Value &
45     destination,
46     rapidjson::Document::AllocatorType &allocator) {
47     if (source.HasMember(TIME_LOW_LIMIT))
48         destination.AddMember(rapidjson::StringRef(TIME_LOW_LIMIT), source[
49             TIME_LOW_LIMIT], allocator);
50
51     if (source.HasMember(TIME_HIGH_LIMIT))
52         destination.AddMember(rapidjson::StringRef(TIME_HIGH_LIMIT), source[
53             TIME_HIGH_LIMIT], allocator);
54
55     return SUCCESS;
56 }
57
58 int TimeRule::rulesAreValid(const rapidjson::Value &value) {
59     if (value.HasMember(TIME_LOW_LIMIT) || value.HasMember(TIME_HIGH_LIMIT)) {
60         return VALID_RULES;
61     } else {
62         return INVALID_RULES;
63     }
64 }
65
66 bool TimeRule::toJson(rapidjson::Value &json, rapidjson::Document::
67     AllocatorType &allocator) {
68
69     rapidjson::Value v;
70
71     if (this->lowerLimit != 0) {
72         v.SetUint64(this->lowerLimit);
73         json.AddMember(rapidjson::StringRef(TIME_LOW_LIMIT), v, allocator);
74     }
75
76     if (this->higerLimit != 0) {
77         v.SetUint64(this->higerLimit);
78         json.AddMember(rapidjson::StringRef(TIME_HIGH_LIMIT), v, allocator);
79     }
80
81     return true;
82 }
83
84 uint64_t TimeRule::now() {
85     return (uint64_t) time(NULL);
86 }
87
88 int TimeRule::configureStorage(Volume *volume) {
89     volume->setMinDelay(this->lowerLimit);
90     volume->setMaxDelay(this->higerLimit);
91     volume->setIsTimeVolume(true);
92
93     return 0;
94 }
95
96 int TimeRule::configureStorage(Pool *pool) {
97     return 0;
98 }

```

```

96 }
97 // namespace mtfs

```

Listing 3.27 – "core/TimeRule.cpp"

3.4.8 UserRightRule

UserRightRule.h

```

1  /**
2   * \file UserRightRule.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #ifndef FILESTORAGE_USER_GROUP_RULE_H
25 #define FILESTORAGE_USER_GROUP_RULE_H
26
27 #include <string>
28 #include <vector>
29 #include <list>
30 #include <iostream>
31 #include <assert.h>
32
33 #include "mtfs/Rule.h"
34
35 namespace mtfs {
36 class UserRightRule : public Rule {
37 public:
38     static constexpr const char *ALLOW_USER = "allowUsers";
39     static constexpr const char *DENY_USER = "denyUsers";
40     static constexpr const char *ALLOW_GROUP = "allowGroups";
41     static constexpr const char *DENY_GROUP = "denyGroups";
42
43 private:
44     std::vector<uid_t> uidAllowed;
45
46     std::vector<uid_t> uidDenied;
47
48     std::vector<gid_t> gidAllowed;
49

```

```

50     std::vector<gid_t> gidDenied;
51
52 public:
53     static int copyConfig(rapidjson::Document &source, rapidjson::Value &
54         destination,
55         rapidjson::Document::AllocatorType &allocator);
56
57     static int rulesAreValid(const rapidjson::Value &value);
58
59     void addAllowUid(uid_t uid);
60
61     void addDenyUid(uid_t uid);
62
63     void addAllowGid(gid_t gid);
64
65     void addDenyGid(gid_t gid);
66
67     bool toJson(rapidjson::Value &json, rapidjson::Document::AllocatorType &
68         allocator) override;
69
70     bool satisfyRules(ruleInfo_st info) override;
71
72 };
73 // namespace mtfs
74 #endif

```

Listing 3.28 – "core/UserRightRule.h"

UserRightRule.cpp

```

1  /**
2   * \file UserRightRule.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFs.
9
10  MTFs is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  FooBar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with FooBar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #include <pwd.h>
25 #include "mtfs/UserRightRule.h"
26
27 namespace mtfs {

```

```

28 int UserRightRule::copyConfig(rapidjson::Document &source, rapidjson::Value &
    destination,
29                             rapidjson::Document::AllocatorType &allocator) {
30 //     TODO copy config
31     if (source.HasMember(ALLOW_USER))
32         destination.AddMember(rapidjson::StringRef(ALLOW_USER), source[ALLOW_USER],
    allocator);
33
34     if (source.HasMember(DENY_USER))
35         destination.AddMember(rapidjson::StringRef(DENY_USER), source[DENY_USER],
    allocator);
36
37     if (source.HasMember(ALLOW_GROUP))
38         destination.AddMember(rapidjson::StringRef(ALLOW_GROUP), source[
    ALLOW_GROUP], allocator);
39
40     if (source.HasMember(DENY_GROUP))
41         destination.AddMember(rapidjson::StringRef(DENY_GROUP), source[DENY_GROUP],
    allocator);
42
43     return SUCCESS;
44 }
45
46 int UserRightRule::rulesAreValid(const rapidjson::Value &value) {
47     if (value.HasMember(ALLOW_USER) || value.HasMember(ALLOW_GROUP) ||
48         value.HasMember(DENY_USER) || value.HasMember(DENY_GROUP)) {
49         return VALID_RULES;
50     } else {
51         return INVALID_RULES;
52     }
53 }
54
55 void UserRightRule::addAllowUid(uid_t uid) {
56     uidAllowed.push_back(uid);
57 }
58
59 void UserRightRule::addDenyUid(uid_t uid) {
60     uidDenied.push_back(uid);
61 }
62
63 void UserRightRule::addAllowGid(gid_t gid) {
64     gidAllowed.push_back(gid);
65 }
66
67 void UserRightRule::addDenyGid(gid_t gid) {
68     gidDenied.push_back(gid);
69 }
70
71 bool UserRightRule::satisfyRules(ruleInfo_st info) {
72     return true;
73 }
74
75 bool UserRightRule::toJson(rapidjson::Value &json, rapidjson::Document::
    AllocatorType &allocator) {
76
77     rapidjson::Value v;
78     rapidjson::Value a(rapidjson::kArrayType);
79
80     struct passwd *pw;

```

```

81 char buffer[50];
82 memset(buffer, 0, 50 * sizeof(char));
83
84 if (this->uidAllowed.size() != 0) {
85     for (auto &&allowed : this->uidAllowed) {
86         pw = getpwuid(allowed);
87         int len = sprintf(buffer, "%s", pw->pw_name);
88         v.SetString(buffer, (rapidjson::SizeType) len, allocator);
89         a.PushBack(v, allocator);
90     }
91     json.AddMember(rapidjson::StringRef(ALLOW_USER), a, allocator);
92 }
93
94 if (this->uidDenied.size() != 0) {
95     for (auto &&denied : this->uidDenied) {
96         pw = getpwuid(denied);
97         int len = sprintf(buffer, "%s", pw->pw_name);
98         v.SetString(buffer, (rapidjson::SizeType) len, allocator);
99         a.PushBack(v, allocator);
100     }
101     json.AddMember(rapidjson::StringRef(DENY_USER), a, allocator);
102 }
103
104 if (this->gidAllowed.size() != 0) {
105     for (auto &&allowed : this->gidAllowed) {
106         pw = getpwuid(allowed);
107         int len = sprintf(buffer, "%s", pw->pw_name);
108         v.SetString(buffer, (rapidjson::SizeType) len, allocator);
109         a.PushBack(v, allocator);
110     }
111     json.AddMember(rapidjson::StringRef(ALLOW_GROUP), a, allocator);
112 }
113
114 if (this->gidDenied.size() != 0) {
115     for (auto &&denied : this->gidDenied) {
116         pw = getpwuid(denied);
117         int len = sprintf(buffer, "%s", pw->pw_name);
118         v.SetString(buffer, (rapidjson::SizeType) len, allocator);
119         a.PushBack(v, allocator);
120     }
121     json.AddMember(rapidjson::StringRef(DENY_GROUP), a, allocator);
122 }
123
124 return true;
125 }
126
127 } // namespace mtfs

```

Listing 3.29 – "core/UserRightRule.cpp"

3.4.9 Volume

Volume.h

```

1 /**
2  * \file Volume.h
3  * \brief
4  * \author David Wittwer

```

```

5  * \version 0.0.1
6  * \copyright GNU Publis License V3
7  *
8  * This file is part of MTFS.
9
10 MTFS is free software: you can redistribute it and/or modify
11 it under the terms of the GNU General Public License as published by
12 the Free Software Foundation, either version 3 of the License, or
13 (at your option) any later version.
14
15 Fooobar is distributed in the hope that it will be useful,
16 but WITHOUT ANY WARRANTY; without even the implied warranty of
17 MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18 GNU General Public License for more details.
19
20 You should have received a copy of the GNU General Public License
21 along with Fooobar. If not, see <http://www.gnu.org/licenses/>.
22 */
23
24 #ifndef FILESTORAGE_VOLUME_H
25 #define FILESTORAGE_VOLUME_H
26
27 #include <string>
28 #include <vector>
29 #include <climits>
30
31 #include <rapidjson/document.h>
32 #include <mtfs/structs.h>
33 #include <pluginSystem/Plugin.h>
34 #include <mutex>
35
36 namespace mtfs {
37     class Volume {
38
39     public:
40         static constexpr const char *VOLUMES = "volumes";
41
42     private:
43         static const uint64_t NEW_DELAY = 50;
44
45         pluginSystem::Plugin *plugin;
46
47         bool isTimeVolume;
48         uint64_t minDelay;
49         uint64_t maxDelay;
50
51         std::mutex iaMutex;
52         std::map<uint64_t, uint64_t> inodesAccess;
53         std::mutex daMutex;
54         std::map<uint64_t, uint64_t> dirBlockAccess;
55         std::mutex baMutex;
56         std::map<uint64_t, uint64_t> blocksAccess;
57
58         std::mutex niMutex;
59         std::map<uint64_t, uint64_t> newInode;
60         std::mutex ndMutex;
61         std::map<uint64_t, uint64_t> newDir;
62         std::mutex nbMutex;
63         std::map<uint64_t, uint64_t> newData;

```

```

64
65 public:
66     virtual ~Volume();
67
68     void setMinDelay(uint64_t minDelay);
69
70     void setMaxDelay(uint64_t maxDelay);
71
72     void setIsTimeVolume(bool b);
73
74     static bool validate(const rapidjson::Value &volume);
75
76     static void
77     structToJson(const volume_t &volume, rapidjson::Value &dest, rapidjson::
Document::AllocatorType &allocator);
78
79     static void jsonToStruct(rapidjson::Value &src, volume_t &volume);
80
81     explicit Volume(pluginSystem::Plugin *plugin);
82
83     int add(uint64_t &id, const blockType &type);
84
85     int add(std::vector<uint64_t> &ids, const int &nb, const blockType &type);
86
87     int del(const uint64_t &id, const blockType &type);
88
89     int get(const uint64_t &id, void *data, const blockType &type);
90
91     int put(const uint64_t &id, const void *data, const blockType &type);
92
93     int getMetas(const uint64_t &id, blockInfo_t &metas, const blockType &type)
;
94
95     int putMetas(const uint64_t &id, const blockInfo_t &metas, const blockType
&type);
96
97     int getUnsatisfy(std::vector<blockInfo_t> &unsatisfy, const blockType &type
, int limit = INT_MAX);
98
99 private:
100     int getOutOfTime(std::vector<uint64_t> &blocks, const blockType &type);
101
102     bool updateLastAccess(const uint64_t &id, const blockType &type);
103
104     void purgeNewMap(const blockType &type);
105 };
106
107 } // namespace mtfs
108 #endif

```

Listing 3.30 – "core/Volume.h"

Volume.cpp

```

1 /**
2  * \file Volume.cpp
3  * \brief
4  * \author David Wittwer
5  * \version 0.0.1

```



```

6  * \copyright GNU Publis License V3
7  *
8  * This file is part of MTFs.
9
10 MTFs is free software: you can redistribute it and/or modify
11 it under the terms of the GNU General Public License as published by
12 the Free Software Foundation, either version 3 of the License, or
13 (at your option) any later version.
14
15 FooBar is distributed in the hope that it will be useful,
16 but WITHOUT ANY WARRANTY; without even the implied warranty of
17 MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18 GNU General Public License for more details.
19
20 You should have received a copy of the GNU General Public License
21 along with FooBar. If not, see <http://www.gnu.org/licenses/>.
22 */
23
24 #include <pluginSystem/PluginManager.h>
25 #include <thread>
26 #include <boost/thread.hpp>
27 #include <boost/threadpool/pool.hpp>
28 #include "mtfs/Volume.h"
29 #include <mtfs/Rule.h>
30
31 using namespace std;
32
33 namespace mtfs {
34
35     Volume::Volume(pluginSystem::Plugin *plugin) : plugin(plugin), minDelay(0),
36         maxDelay(0) {}
37
38     Volume::~Volume() {
39         plugin->detach();
40
41         pluginSystem::PluginManager::getInstance()->freePlugin(this->plugin->
42             getName(), this->plugin);
43     }
44
45     bool Volume::validate(const rapidjson::Value &volume) {
46         if (!volume.HasMember(pluginSystem::Plugin::TYPE))
47             return false;
48
49         if (!volume.HasMember(pluginSystem::Plugin::PARAMS))
50             return false;
51
52         string volType = volume[pluginSystem::Plugin::TYPE].GetString();
53
54         pluginSystem::PluginManager *manager = pluginSystem::PluginManager::
55             getInstance();
56
57         pluginSystem::Plugin *plugin = manager->getPlugin(volType);
58         if (nullptr == plugin)
59             return false;
60
61         pluginSystem::pluginInfo_t info;
62
63         if (manager->getInfo(volType, info) != pluginSystem::PluginManager::SUCCESS
64         )

```

```

61     return false;
62
63     for (auto &&inf : info.params) {
64         if (!volume[pluginSystem::Plugin::PARAMS].HasMember(inf.c_str()))
65             return false;
66     }
67
68     //     vector<string> infos = plugin->getInfos();
69     //     infos.erase(infos.begin());
70     //     for (auto i:infos)
71     //         if (!volume.HasMember(i.c_str()))
72     //             return false;
73
74     return true;
75 }
76
77 void Volume::structToJson(const volume_t &volume, rapidjson::Value &dest,
78                          rapidjson::Document::AllocatorType &allocator) {
79     volume.rule->toJson(dest, allocator);
80
81     rapidjson::Value v;
82
83     v.SetString(rapidjson::StringRef(volume.pluginName.c_str()));
84     dest.AddMember(rapidjson::StringRef(pluginSystem::Plugin::TYPE), v,
85                   allocator);
86
87     rapidjson::Value p(rapidjson::kObjectType);
88     for (auto &&item : volume.params) {
89         v.SetString(rapidjson::StringRef(item.second.c_str()));
90         p.AddMember(rapidjson::StringRef(item.first.c_str()), v, allocator);
91     }
92     dest.AddMember(rapidjson::StringRef(pluginSystem::Plugin::PARAMS), p,
93                   allocator);
94 }
95
96 void Volume::jsonToStruct(rapidjson::Value &src, volume_t &volume) {
97     assert(src.HasMember(pluginSystem::Plugin::TYPE));
98     volume.pluginName = src[pluginSystem::Plugin::TYPE].GetString();
99
100     assert(src.HasMember(pluginSystem::Plugin::PARAMS));
101     for (auto &&item : src[pluginSystem::Plugin::PARAMS].GetObject()) {
102         volume.params.insert(make_pair(item.name.GetString(), item.value.
103                                     GetString()));
104     }
105 }
106
107 int Volume::add(uint64_t &id, const blockType &type) {
108     int ret;
109     ret = this->plugin->add(&id, type);
110
111     if (ENOSYS == ret) {
112         //         TODO log noimplemented
113     }
114     mutex *mu = nullptr;
115     map<uint64_t, uint64_t> *ma = nullptr;
116
117     switch (type) {
118         case INODE:
119             mu = &this->niMutex;

```

```

117     ma = &this->newInode;
118     break;
119     case DIR_BLOCK:
120         mu = &this->ndMutex;
121         ma = &this->newData;
122         break;
123     case DATA_BLOCK:
124         mu = &this->nbMutex;
125         ma = &this->newData;
126         break;
127     default:
128         return ENOSYS;
129 }
130
131 //      add id in newBlock map
132 unique_lock<mutex> lk(*mu);
133 ma->emplace(id, time(nullptr));
134 lk.unlock();
135
136 return ret;
137 }
138
139 int Volume::add(std::vector<uint64_t> &ids, const int &nb, const blockType &
type) {
140     boost::threadpool::pool thPool((size_t) nb);
141
142     vector<uint64_t *> tmp;
143     for (int i = 0; i < nb; ++i) {
144         uint64_t *id;
145         id = new uint64_t();
146         tmp.push_back(id);
147         thPool.schedule(bind(&pluginSystem::Plugin::add, this->plugin, id, type))
;
148     }
149
150     thPool.wait();
151
152     mutex *mu = nullptr;
153     map<uint64_t, uint64_t> *ma = nullptr;
154
155     switch (type) {
156     case INODE:
157         mu = &this->niMutex;
158         ma = &this->newInode;
159         break;
160     case DIR_BLOCK:
161         mu = &this->ndMutex;
162         ma = &this->newData;
163         break;
164     case DATA_BLOCK:
165         mu = &this->nbMutex;
166         ma = &this->newData;
167         break;
168     default:
169         return ENOSYS;
170     }
171
172     for (auto &&item: tmp) {
173 //      add id in newBlock map

```

```

174     unique_lock<mutex> lk(*mu);
175     ma->emplace(*item, time(nullptr));
176     lk.unlock();
177
178     ids.push_back(*item);
179     delete item;
180 }
181
182 return 0;
183 }
184
185 int Volume::del(const uint64_t &id, const blockType &type) {
186     int ret;
187     ret = this->plugin->del(id, type);
188
189     if (ENOSYS == ret) {
190         // TODO log noimplemented
191     }
192
193     map<uint64_t, uint64_t> *access = nullptr;
194     mutex *mu;
195     switch (type) {
196     case INODE:
197         mu = &this->iaMutex;
198         access = &this->inodesAccess;
199         break;
200     case DIR_BLOCK:
201         mu = &this->daMutex;
202         access = &this->dirBlockAccess;
203         break;
204     case DATA_BLOCK:
205         mu = &this->baMutex;
206         access = &this->blocksAccess;
207         break;
208     default:
209         return ENOSYS;
210     }
211
212     unique_lock<mutex> lk(*mu);
213     access->erase(id);
214     lk.unlock();
215
216     return ret;
217 }
218
219 int Volume::get(const uint64_t &id, void *data, const blockType &type) {
220     int ret;
221     ret = this->plugin->get(id, data, type, false);
222
223     if (ENOSYS == ret) {
224         // TODO log noimplemented
225     }
226
227     this->updateLastAccess(id, type);
228
229     return ret;
230 }
231

```

```

232 int Volume::put(const uint64_t &id, const void *data, const blockType &type)
233 {
234     int ret;
235     ret = this->plugin->put(id, data, type, false);
236
237     if (ENOSYS == ret) {
238         // TODO log noimplemented
239     }
240
241     this->updateLastAccess(id, type);
242
243     return ret;
244 }
245
246 int Volume::getMetas(const uint64_t &id, blockInfo_t &metas, const blockType
247 &type) {
248     int ret;
249     if (EXIT_SUCCESS != (ret = this->plugin->get(id, &metas, type, true))) {
250         return ret;
251     }
252
253     map<uint64_t, uint64_t> *access = nullptr;
254     mutex *mu;
255     switch (type) {
256     case INODE:
257         mu = &this->iaMutex;
258         access = &this->inodesAccess;
259         break;
260     case DIR_BLOCK:
261         mu = &this->daMutex;
262         access = &this->dirBlockAccess;
263         break;
264     case DATA_BLOCK:
265         mu = &this->baMutex;
266         access = &this->blocksAccess;
267         break;
268     default:
269         return ENOSYS;
270     }
271
272     unique_lock<mutex> lk(*mu);
273     if (access->end() == access->find(id))
274         (*access)[id] = metas.lastAccess;
275
276     return ret;
277 }
278
279 int Volume::putMetas(const uint64_t &id, const blockInfo_t &metas, const
280 blockType &type) {
281     if (0 == id && INODE == type)
282         return EXIT_SUCCESS;
283
284     int ret;
285     if (EXIT_SUCCESS != (ret = this->plugin->put(id, &metas, type, true))) {
286         return ret;
287     }
288
289     map<uint64_t, uint64_t> *access = nullptr;
290     mutex *mu;

```

```

288     switch (type) {
289         case INODE:
290             mu = &this->iaMutex;
291             access = &this->inodesAccess;
292             break;
293         case DIR_BLOCK:
294             mu = &this->daMutex;
295             access = &this->dirBlockAccess;
296             break;
297         case DATA_BLOCK:
298             mu = &this->baMutex;
299             access = &this->blocksAccess;
300             break;
301         default:
302             return ENOSYS;
303     }
304
305     unique_lock<mutex> lk(*mu);
306     (*access)[id] = metas.lastAccess;
307
308     return ret;
309 }
310
311 bool Volume::updateLastAccess(const uint64_t &id, const blockType &type) {
312     blockInfo_t metas = blockInfo_t();
313     this->getMetas(id, metas, type);
314
315     metas.lastAccess = static_cast<uint64_t>(time(nullptr));
316     this->putMetas(id, metas, type);
317
318     return true;
319 }
320
321 void Volume::setMinDelay(uint64_t minDelay) {
322     this->minDelay = minDelay;
323 }
324
325 void Volume::setMaxDelay(uint64_t maxDelay) {
326     this->maxDelay = maxDelay;
327 }
328
329 void Volume::setIsTimeVolume(bool b) {
330     this->isTimeVolume = b;
331 }
332
333 int Volume::getUnsatisfy(vector<blockInfo_t> &unsatisfy, const blockType &
334     type, const int limit) {
335     this->purgeNewMap(type);
336
337     int nb = 0;
338     if (this->isTimeVolume) {
339         vector<uint64_t> under;
340         this->getOutOfTime(under, type);
341
342         for (auto &&blk : under) {
343             blockInfo_t info = blockInfo_t();
344             this->getMetas(blk, info, type);
345             info.id.id = blk;

```

```

346         unsatisfy.push_back(info);
347
348         nb++;
349         if (limit == nb)
350             return 0;
351     }
352 }
353
354 return 0;
355 }
356
357 int Volume::getOutOfTime(vector<uint64_t> &blocks, const blockType &type) {
358     uint64_t now;
359     now = (uint64_t) time(nullptr);
360     const uint64_t maxTimestamp = now - this->minDelay;
361     const uint64_t minTimestamp = 0 == this->maxDelay ? 0 : now - this->
maxDelay;
362
363     mutex *mu;
364     map<uint64_t, uint64_t> *mp;
365     mutex *newMut;
366     map<uint64_t, uint64_t> *newMap;
367     switch (type) {
368     case INODE:
369         mu = &this->iaMutex;
370         mp = &this->inodesAccess;
371         newMut = &this->niMutex;
372         newMap = &this->newInode;
373         break;
374     case DIR_BLOCK:
375         mu = &this->daMutex;
376         mp = &this->dirBlockAccess;
377         newMut = &this->ndMutex;
378         newMap = &this->newDir;
379         break;
380     case DATA_BLOCK:
381         mu = &this->baMutex;
382         mp = &this->blocksAccess;
383         newMut = &this->ndMutex;
384         newMap = &this->newData;
385         break;
386     default:
387         return ENOSYS;
388     }
389
390     unique_lock<mutex> lk(*mu);
391     for (auto &&item : *mp) {
392         if (0 == item.first && INODE == type)
393             continue;
394         unique_lock<mutex> nlk(*newMut);
395         if (newMap->end() != newMap->find(item.first)) {
396             nlk.unlock();
397             continue;
398         }
399         nlk.unlock();
400
401         if (!(maxTimestamp > item.second && minTimestamp < item.second))
402             blocks.push_back(item.first);
403     }

```

```

404     return 0;
405 }
406
407 void Volume::purgeNewMap(const blockType &type) {
408     mutex *mu = nullptr;
409     map<uint64_t, uint64_t> *ma = nullptr;
410
411     switch (type) {
412     case INODE:
413         mu = &this->niMutex;
414         ma = &this->newInode;
415         break;
416     case DIR_BLOCK:
417         mu = &this->ndMutex;
418         ma = &this->newData;
419         break;
420     case DATA_BLOCK:
421         mu = &this->nbMutex;
422         ma = &this->newData;
423         break;
424     default:
425         return;
426     }
427
428     uint64_t newTimestamp = time(nullptr) - NEW_DELAY;
429
430     unique_lock<mutex> lk(*mu);
431     for (auto &&item : *ma) {
432         if (newTimestamp > item.second)
433             ma->erase(item.first);
434     }
435 }
436
437 } // namespace mtf
438

```

Listing 3.31 – "core/Volume.cpp"

3.5 Migrator

3.5.1 Migrator.h

```

1  /**
2   * \file Migrator.h
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Public License V3
7   *
8   * This file is part of MTFs.
9
10  MTFs is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  FooBar is distributed in the hope that it will be useful,

```



```

16 but WITHOUT ANY WARRANTY; without even the implied warranty of
17 MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18 GNU General Public License for more details.
19
20 You should have received a copy of the GNU General Public License
21 along with FooBar. If not, see <http://www.gnu.org/licenses/>.
22 */
23
24
25 #ifndef MTFS_MIGRATOR_H
26 #define MTFS_MIGRATOR_H
27
28 #include <mutex>
29 #include <condition_variable>
30 #include "PoolManager.h"
31
32 namespace mtfs {
33     class Migrator {
34     private:
35         static const int MIGRATION_DELAY = 10;
36
37     public:
38         struct info_st {
39             std::mutex *endMutex;
40             std::condition_variable condV;
41             bool end;
42             PoolManager *poolManager;
43
44             info_st() : endMutex(new std::mutex), end(false) {};
45
46             ~info_st() {
47                 delete endMutex;
48             }
49         };
50
51         static void main(info_st *infos);
52     };
53 }
54
55
56 #endif //MTFS_MIGRATOR_H

```

Listing 3.32 – "migrator/Migrator.h"

3.5.2 Migrator.cpp

```

1 /**
2  * \file Migrator.cpp
3  * \brief
4  * \author David Wittwer
5  * \version 0.0.1
6  * \copyright GNU Publis License V3
7  *
8  * This file is part of MTFS.
9
10 MTFS is free software: you can redistribute it and/or modify
11 it under the terms of the GNU General Public License as published by
12 the Free Software Foundation, either version 3 of the License, or

```

```

13     (at your option) any later version.
14
15     Fooobar is distributed in the hope that it will be useful,
16     but WITHOUT ANY WARRANTY; without even the implied warranty of
17     MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18     GNU General Public License for more details.
19
20     You should have received a copy of the GNU General Public License
21     along with Fooobar. If not, see <http://www.gnu.org/licenses/>.
22 */
23
24 #include <zconf.h>
25 #include <iostream>
26 #include <utils/Logger.h>
27 #include "mtfs/Migrator.h"
28
29 #define MIGRATION_DELAY chrono::seconds(30)
30
31 using namespace std;
32
33 void mtfs::Migrator::main(info_st *infos) {
34
35     unique_lock<mutex> lk(*infos->endMutex);
36     while (!infos->end) {
37         lk.unlock();
38
39         lk.lock();
40         infos->condV.wait_for(lk, MIGRATION_DELAY);
41         if (infos->end)
42             break;
43         lk.unlock();
44
45         // Do migration;
46         Logger::getInstance()->log("MIGRATOR", "start inode migration", Logger::
L_DEBUG);
47         infos->poolManager->doMigration(blockType::INODE);
48         Logger::getInstance()->log("MIGRATOR", "start dirblock migration", Logger::
L_DEBUG);
49         infos->poolManager->doMigration(blockType::DIR_BLOCK);
50         Logger::getInstance()->log("MIGRATOR", "start datablock migration", Logger
::L_DEBUG);
51         infos->poolManager->doMigration(blockType::DATA_BLOCK);
52         Logger::getInstance()->log("MIGRATOR", "end migration", Logger::L_DEBUG);
53
54         lk.lock();
55     }
56     lk.unlock();
57
58 }

```

Listing 3.33 – "migrator/Migrator.cpp"

3.6 Tests

3.7 Migration test

Listing 3.34 – "migrationTests.py"

3.8 test de vitesse

```

1  /**
2   * \file main.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #include <iostream>
25 #include <boost/filesystem.hpp>
26 #include <fcntl.h>
27
28 #define FIRST_SIZE 512
29 #define LAST_SIZE 131072
30 #define SAMPLE 5
31 #define LOOP 50
32
33 #define FILENAME "/tmp/mtfs/test24"
34
35 using namespace std;
36
37 int main(int argc, char **argv) {
38     cout << "Start Mtfs perfs tests" << endl;
39
40     ofstream file("sample.csv");
41
42     creat(FILENAME, S_IRUSR | S_IWUSR | S_IRGRP | S_IWGRP | S_IROTH | S_IWOTH);
43
44     file << "octets ";
45     for (int i = 1; i <= LOOP; ++i) {
46         file << to_string(i) << " ";
47     }
48     file << endl;
49
50     string readTable, writeTable;
51
52     for (size_t i = FIRST_SIZE; i <= LAST_SIZE; i += 2) {

```

```

53 string readRow, writeRow;
54 readRow = to_string(i);
55 writeRow = to_string(i);
56
57 for (int j = 0; j < LOOP; ++j) {
58     char buffer[i + 1];
59     char *rbuffer = static_cast<char *>(calloc(sizeof(char), i + 1));
60     for (int k = 0; k < i; ++k) {
61         buffer[k] = 'a';
62     }
63
64     clock_t writeSum = 0;
65     clock_t readSum = 0;
66
67     for (int l = 0; l < SAMPLE; ++l) {
68         int fd = open(FILENAME, O_SYNC | O_RDWR);
69         if (fd <= 0) {
70             cerr << "open error " << strerror(errno) << endl;
71             break;
72         }
73         ssize_t nb = 0;
74         clock_t start;
75
76         start = clock();
77         do {
78             nb += write(fd, buffer, i);
79         } while (nb < i);
80         writeSum += clock() - start;
81
82         if (nb < 0) {
83             cerr << "write error " << strerror(errno) << endl;
84             break;
85         }
86
87         lseek(fd, 0, SEEK_SET);
88
89         start = clock();
90         do {
91             nb += read(fd, buffer, i);
92         } while (nb < i);
93         readSum += clock() - start;
94         close(fd);
95     }
96     float wm = (float) writeSum / SAMPLE;
97     float rm = (float) readSum / SAMPLE;
98     // cout << to_string(m) << " " << to_string((m / CLOCKS_PER_SEC) * 1000)
99     << endl;
100     writeRow += " " + to_string(wm);
101     readRow += " " + to_string(rm);
102     // file << to_string(m) << " ";
103 }
104
105 writeTable += writeRow + "\n";
106 readTable += readRow + "\n";
107 cout << "size " << to_string(i) << endl;
108 }
109
110 file << readTable << endl << writeTable << endl;

```

111
112 }

Listing 3.35 – "speedTest.cpp"

3.9 Plugin tests

3.9.1 block device

```

1  /**
2   * \file blockDeviceTests.cpp
3   * \brief
4   * \author David Wittwer
5   * \version 0.0.1
6   * \copyright GNU Publis License V3
7   *
8   * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22  */
23
24 #include <gtest/gtest.h>
25 #include <BlockDevice/BlockDevice.h>
26 #include <mtfs/Mtfs.h>
27 // #include <mtfs/structs.h>
28
29
30 using namespace std;
31 using namespace pluginSystem;
32
33 #define HOME "/home/david/Cours/4eme/Travail_bachelor/Home/Plugins/"
34 #define BLOCK_SIZE 4096
35 #define DEVICE "/dev/tests"
36 #define FS_TYPE "ext4"
37
38 TEST(BlockDevice, attachDetach) {
39 #ifndef DEBUG
40     if (setuid(0) != 0)
41         cout << "fail setuid" << endl;
42 #endif
43
44     BlockDevice blockDevice;
45     map<string, string> params;
46     params.insert(make_pair("home", HOME));
47     params.insert(make_pair("blockSize", to_string(BLOCK_SIZE)));
48     params.insert(make_pair("devicePath", DEVICE));

```

```

49     params.insert(make_pair("fsType", FS_TYPE));
50     ASSERT_TRUE(blockDevice.attach(params));
51     ASSERT_TRUE(blockDevice.detach());
52 }
53
54 class BlockDeviceFixture : public ::testing::Test {
55 public:
56     BlockDeviceFixture() {
57         rootIdent.poolId = 0;
58         rootIdent.volumeId = 0;
59         rootIdent.id = 0;
60     }
61
62     virtual void SetUp() {
63         map<string, string> params;
64         params.insert(make_pair("home", HOME));
65         params.insert(make_pair("blockSize", "4096"));
66         params.insert(make_pair("devicePath", DEVICE));
67         params.insert(make_pair("fsType", FS_TYPE));
68         blockDevice.attach(params);
69     }
70
71     virtual void TearDown() {
72         blockDevice.detach();
73     }
74
75     ~BlockDeviceFixture() {
76     }
77
78     BlockDevice blockDevice;
79     mtfs::ident_t rootIdent;
80 };
81
82 TEST_F(BlockDeviceFixture, addInode) {
83     uint64_t inode = 0;
84     ASSERT_EQ(0, blockDevice.add(&inode, mtfs::blockType::INODE));
85     ASSERT_NE(0, inode);
86
87     uint64_t inode2 = 0;
88     ASSERT_EQ(0, blockDevice.add(&inode2, mtfs::blockType::INODE));
89     ASSERT_NE(0, inode2);
90     ASSERT_GT(inode2, inode);
91
92     uint64_t inode3 = 0;
93     ASSERT_EQ(0, blockDevice.add(&inode3, mtfs::blockType::INODE));
94 }
95
96 TEST_F(BlockDeviceFixture, delInode) {
97     uint64_t inode;
98     blockDevice.add(&inode, mtfs::blockType::INODE);
99     ASSERT_EQ(0, blockDevice.del(inode, mtfs::blockType::INODE));
100
101     uint64_t inode2;
102     blockDevice.add(&inode2, mtfs::blockType::INODE);
103     blockDevice.add(&inode2, mtfs::blockType::INODE);
104     blockDevice.add(&inode2, mtfs::blockType::INODE);
105     blockDevice.del(inode, mtfs::blockType::INODE);
106 }
107

```

```

108 TEST_F(BlockDeviceFixture, writeInode) {
109     uint64_t inodeId;
110     blockDevice.add(&inodeId, mtfs::blockType::INODE);
111
112     mtfs::ident_t oIdent;
113     oIdent.id = 42;
114     oIdent.volumeId = 1;
115     oIdent.poolId = 1;
116
117     mtfs::inode_t ino;
118     ino.accesRight = 0666;
119     ino.uid = 0;
120     ino.gid = 0;
121     ino.size = 1024;
122     ino.linkCount = 1;
123     ino.atime = (unsigned long &&) time(nullptr);
124
125     for (int i = 0; i < 4; ++i) {
126         vector<mtfs::ident_t> blocks;
127
128         for (uint j = 0; j < 3; ++j) {
129             mtfs::ident_t ident = mtfs::ident_t(1, j, i * 3 + j);
130
131             blocks.push_back(ident);
132         }
133         ino.dataBlocks.push_back(blocks);
134     }
135
136     ASSERT_EQ(0, blockDevice.put(inodeId, &ino, mtfs::blockType::INODE, false));
137 }
138
139 TEST_F(BlockDeviceFixture, readInode) {
140     mtfs::inode_t original, inode;
141     memset(&original, 0, sizeof(mtfs::inode_t));
142     memset(&inode, 0, sizeof(mtfs::inode_t));
143
144     original.accesRight = 0644;
145     original.uid = 1;
146     original.gid = 1;
147     original.size = 1024;
148     original.linkCount = 1;
149     original.atime = (unsigned long &&) time(nullptr);
150
151     for (int i = 0; i < 4; ++i) {
152         vector<mtfs::ident_t> blocks;
153
154         for (uint j = 0; j < 3; ++j) {
155             mtfs::ident_t ident = mtfs::ident_t(1, j, i * 3 + j);
156
157             blocks.push_back(ident);
158         }
159         original.dataBlocks.push_back(blocks);
160     }
161
162     uint64_t inodeId;
163     blockDevice.add(&inodeId, mtfs::blockType::INODE);
164     blockDevice.put(inodeId, &original, mtfs::blockType::INODE, false);
165

```

```

166 ASSERT_EQ(0, blockDevice.get(inodeId, &inode, mfts::blockType::INODE, false))
167 ;
168 // cout << original.accesRight << " " << inode.accesRight << endl;
169 ASSERT_EQ(original, inode);
170 original.size = 2048;
171 blockDevice.put(inodeId, &original, mfts::blockType::INODE, false);
172 ASSERT_EQ(0, blockDevice.get(inodeId, &inode, mfts::blockType::INODE, false))
173 ;
174 }
175 TEST_F(BlockDeviceFixture, addBlock) {
176     uint64_t block = 0;
177     ASSERT_EQ(0, blockDevice.add(&block, mfts::blockType::DATA_BLOCK));
178     ASSERT_NE(0, block);
179
180     uint64_t block2 = 0;
181     ASSERT_EQ(0, blockDevice.add(&block2, mfts::blockType::DATA_BLOCK));
182     ASSERT_NE(0, block2);
183     ASSERT_GT(block2, block);
184
185     uint64_t block3 = 0;
186     ASSERT_EQ(0, blockDevice.add(&block3, mfts::blockType::DATA_BLOCK));
187 }
188
189 TEST_F(BlockDeviceFixture, delBlock) {
190     uint64_t block;
191     blockDevice.add(&block, mfts::blockType::DATA_BLOCK);
192     ASSERT_EQ(0, blockDevice.del(block, mfts::blockType::DATA_BLOCK));
193
194     uint64_t block2;
195     blockDevice.add(&block2, mfts::blockType::DATA_BLOCK);
196     blockDevice.add(&block2, mfts::blockType::DATA_BLOCK);
197     blockDevice.add(&block2, mfts::blockType::DATA_BLOCK);
198     blockDevice.del(block, mfts::blockType::DATA_BLOCK);
199 }
200
201 TEST_F(BlockDeviceFixture, writeBlock) {
202     uint8_t block[BLOCK_SIZE];
203     memset(block, 0, BLOCK_SIZE);
204     for (uint8_t i = 0; i < 50; ++i) {
205         block[i] = 'a';
206     }
207
208     uint64_t blockId;
209     blockDevice.add(&blockId, mfts::blockType::DATA_BLOCK);
210     ASSERT_EQ(0, blockDevice.put(blockId, &block, mfts::blockType::DATA_BLOCK,
211         false));
212 }
213
214 TEST_F(BlockDeviceFixture, readBlock) {
215     uint8_t block[BLOCK_SIZE];
216     memset(block, 0, BLOCK_SIZE);
217
218     uint8_t readBlock[BLOCK_SIZE];
219     memset(readBlock, 0, BLOCK_SIZE);
220
221     for (int i = 0; i < 500; ++i) {
222         block[i] = (uint8_t) to_string(i)[0];

```



```

222 }
223
224 uint64_t blockId;
225 blockDevice.add(&blockId, mtfs::blockType::DATA_BLOCK);
226 blockDevice.put(blockId, &block, mtfs::blockType::DATA_BLOCK, false);
227
228 ASSERT_EQ(0, blockDevice.get(blockId, &readBlock, mtfs::blockType::DATA_BLOCK
229 , false));
230 ASSERT_TRUE(0 == memcmp(block, readBlock, BLOCK_SIZE));
231 }
232
233 TEST_F(BlockDeviceFixture, rootInode) {
234     mtfs::inode_t inode;
235     inode.accessRight = 0444;
236     inode.uid = 1000;
237     inode.gid = 1000;
238     inode.size = 1024;
239     inode.linkCount = 2;
240     inode.atime = (uint64_t) time(nullptr);
241     inode.dataBlocks.clear();
242
243     mtfs::inode_t readInode;
244
245     ASSERT_EQ(0, blockDevice.put(0, &inode, mtfs::blockType::INODE, false));
246     ASSERT_EQ(0, blockDevice.get(0, &readInode, mtfs::blockType::INODE, false));
247     ASSERT_EQ(inode, readInode);
248 }
249
250 TEST_F(BlockDeviceFixture, superblock) {
251     mtfs::superblock_t superblock;
252     superblock.iCacheSz = superblock.dCacheSz = superblock.bCacheSz = superblock.
253         blockSz = 4096;
254     superblock.migration = superblock.redundancy = 1;
255     superblock.pools.clear();
256     for (int i = 0; i < 5; ++i) {
257         mtfs::pool_t pool;
258         pool.migration = 0;
259         pool.rule = nullptr;
260         pool.volumes.clear();
261         superblock.pools[i] = pool;
262     }
263
264     ASSERT_TRUE(blockDevice.putSuperblock(superblock));
265 }
266
267 TEST_F(BlockDeviceFixture, putMetas) {
268     mtfs::blockInfo_t blockInfo;
269
270     mtfs::ident_t ident1(1, 1, 1);
271     mtfs::ident_t ident2(2, 2, 2);
272
273     blockInfo.referenceId.push_back(ident1);
274     blockInfo.referenceId.push_back(ident2);
275     blockInfo.lastAccess = (uint64_t) time(nullptr);
276
277     ASSERT_EQ(0, blockDevice.put(1, &blockInfo, mtfs::blockType::INODE, true));
278     ASSERT_EQ(0, blockDevice.put(1, &blockInfo, mtfs::blockType::DIR_BLOCK, true)
279 );

```

```

278 ASSERT_EQ(0, blockDevice.put(1, &blockInfo, mtfs::blockType::DATA_BLOCK, true
279 ));
280 }
281 TEST_F(BlockDeviceFixture, getMetas) {
282     mtfs::blockInfo_t blockInfo = mtfs::blockInfo_t(), receiveInfo = mtfs::
283         blockInfo_t();
284     // memset(&blockInfo, 0, sizeof(mtrfs::blockInfo_t));
285     // memset(&receiveInfo, 0, sizeof(mtrfs::blockInfo_t));
286
287     mtfs::ident_t ident1(1, 1, 1);
288     mtfs::ident_t ident2(2, 2, 2);
289
290     blockInfo.referenceId.push_back(ident1);
291     blockInfo.referenceId.push_back(ident2);
292     blockInfo.lastAccess = (uint64_t) time(nullptr);
293
294     ASSERT_EQ(0, blockDevice.put(2, &blockInfo, mtfs::blockType::INODE, true));
295     ASSERT_EQ(0, blockDevice.get(2, &receiveInfo, mtfs::blockType::INODE, true));
296     ASSERT_EQ(blockInfo, receiveInfo);
297
298     receiveInfo = mtfs::blockInfo_t();
299     ASSERT_EQ(0, blockDevice.put(2, &blockInfo, mtfs::blockType::DIR_BLOCK, true
300 ));
301     ASSERT_EQ(0, blockDevice.get(2, &receiveInfo, mtfs::blockType::DIR_BLOCK,
302 true));
303     ASSERT_EQ(blockInfo, receiveInfo);
304
305     receiveInfo = mtfs::blockInfo_t();
306     ASSERT_EQ(0, blockDevice.put(2, &blockInfo, mtfs::blockType::DATA_BLOCK, true
307 ));
308     ASSERT_EQ(0, blockDevice.get(2, &receiveInfo, mtfs::blockType::DATA_BLOCK,
309 true));
310     ASSERT_EQ(blockInfo, receiveInfo);
311 }
312
313 TEST_F(BlockDeviceFixture, putDirBlock) {
314     mtfs::ident_t id1, id2;
315     id2.poolId = id1.poolId = 1;
316     id1.volumeId = 1;
317     id1.id = 2;
318     id2.volumeId = 2;
319     id2.id = 42;
320
321     vector<mtfs::ident_t> ids;
322     ids.push_back(id1);
323     ids.push_back(id2);
324
325     mtfs::dirBlock_t block = mtfs::dirBlock_t();
326     block.entries.clear();
327     block.entries.insert(make_pair("baz", ids));
328     block.entries.insert(make_pair("test", ids));
329
330     uint64_t dId = 0;
331
332     ASSERT_EQ(0, blockDevice.add(&dId, mtfs::blockType::DIR_BLOCK));
333     EXPECT_LE(0, dId);
334     ASSERT_EQ(0, blockDevice.put(dId, &block, mtfs::blockType::DIR_BLOCK, false))
335 ;

```

```

330
331 mtfs::dirBlock_t readBlock;
332 ASSERT_EQ(0, blockDevice.get(dId, &readBlock, mtfs::blockType::DIR_BLOCK,
333     false));
334 ASSERT_EQ(block.entries.size(), readBlock.entries.size());
335 for (auto &&item: block.entries) {
336     ASSERT_NE(readBlock.entries.end(), readBlock.entries.find(item.first));
337     int i = 0;
338     for (auto &&ident: item.second) {
339         ASSERT_EQ(ident, readBlock.entries[item.first][i]);
340         i++;
341     }
342 }
343 ASSERT_EQ(0, blockDevice.del(dId, mtfs::blockType::DIR_BLOCK));
344 }
345
346 int main(int argc, char **argv) {
347     ::testing::InitGoogleTest(&argc, argv);
348     return RUN_ALL_TESTS();
349 }

```

Listing 3.36 – "blockDeviceTests.cpp"

3.9.2 S3

```

1 /**
2  * \file s3Tests.cpp
3  * \brief
4  * \author David Wittwer
5  * \version 0.0.1
6  * \copyright GNU Publis License V3
7  *
8  * This file is part of MTFS.
9
10  MTFS is free software: you can redistribute it and/or modify
11  it under the terms of the GNU General Public License as published by
12  the Free Software Foundation, either version 3 of the License, or
13  (at your option) any later version.
14
15  Foobar is distributed in the hope that it will be useful,
16  but WITHOUT ANY WARRANTY; without even the implied warranty of
17  MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the
18  GNU General Public License for more details.
19
20  You should have received a copy of the GNU General Public License
21  along with Foobar. If not, see <http://www.gnu.org/licenses/>.
22 */
23
24 #include <gtest/gtest.h>
25 #include <S3/S3.h>
26 #include <mtfs/Mtfs.h>
27 // #include <mtfs/structs.h>
28
29 #define BLOCK_SIZE 4096
30
31 using namespace std;
32 using namespace pluginSystem;

```

```

33
34 #define HOME "/home/david/Cours/4eme/Travail_bachelor/Home/Plugins/"
35
36 TEST(S3, attachDetach) {
37     #ifndef DEBUG
38         if (setuid(0) != 0)
39             cout << "fail setuid" << endl;
40     #endif
41
42     S3 device;
43     map<string, string> params;
44     params.insert(make_pair("home", HOME));
45     params.insert(make_pair("blockSize", to_string(BLOCK_SIZE)));
46     params.emplace("region", "eu-central-1");
47     params.emplace("bucket", "mtfs");
48     ASSERT_TRUE(device.attach(params));
49     ASSERT_TRUE(device.detach());
50 }
51
52 class S3Fixture : public ::testing::Test {
53 public:
54     S3Fixture() {
55         rootIdent.poolId = 0;
56         rootIdent.volumeId = 0;
57         rootIdent.id = 0;
58     }
59
60     virtual void SetUp() {
61         map<string, string> params;
62         params.emplace("home", HOME);
63         params.emplace("blockSize", to_string(BLOCK_SIZE));
64         params.emplace("region", "eu-central-1");
65         params.emplace("bucket", "mtfs");
66         s3.attach(params);
67     }
68
69     virtual void TearDown() {
70         s3.detach();
71     }
72
73     ~S3Fixture() {
74     }
75
76     S3 s3;
77     mtfs::ident_t rootIdent;
78 };
79
80 TEST_F(S3Fixture, addInode) {
81     uint64_t inode = 0;
82     ASSERT_EQ(0, s3.add(&inode, mtfs::blockType::INODE));
83     ASSERT_NE(0, inode);
84
85     uint64_t inode2 = 0;
86     ASSERT_EQ(0, s3.add(&inode2, mtfs::blockType::INODE));
87     ASSERT_NE(0, inode2);
88     ASSERT_GT(inode2, inode);
89
90     uint64_t inode3 = 0;

```

```

92  ASSERT_EQ(0, s3.add(&inode3, mtfs::blockType::INODE));
93  }
94
95  TEST_F(S3Fixture, delInode) {
96      uint64_t inode;
97      s3.add(&inode, mtfs::blockType::INODE);
98      ASSERT_EQ(0, s3.del(inode, mtfs::blockType::INODE));
99
100     uint64_t inode2;
101     s3.add(&inode2, mtfs::blockType::INODE);
102     s3.add(&inode2, mtfs::blockType::INODE);
103     s3.add(&inode2, mtfs::blockType::INODE);
104     s3.del(inode, mtfs::blockType::INODE);
105 }
106
107 TEST_F(S3Fixture, writeInode) {
108     uint64_t inodeId = 0;
109     s3.add(&inodeId, mtfs::blockType::INODE);
110
111     mtfs::ident_t oIdent;
112     oIdent.id = 42;
113     oIdent.volumeId = 1;
114     oIdent.poolId = 1;
115
116     mtfs::inode_t ino;
117     ino.accesRight = 0666;
118     ino.uid = 0;
119     ino.gid = 0;
120     ino.size = 1024;
121     ino.linkCount = 1;
122     ino.atime = (unsigned long &&) time(nullptr);
123
124     for (int i = 0; i < 4; ++i) {
125         vector<mtfs::ident_t> blocks;
126
127         for (uint j = 0; j < 3; ++j) {
128             mtfs::ident_t ident = mtfs::ident_t(1, j, i * 3 + j);
129
130             blocks.push_back(ident);
131         }
132         ino.dataBlocks.push_back(blocks);
133     }
134
135     ASSERT_EQ(0, s3.put(inodeId, &ino, mtfs::blockType::INODE, false));
136 }
137
138 TEST_F(S3Fixture, readInode) {
139     mtfs::inode_t original, inode;
140     memset(&original, 0, sizeof(mtfs::inode_t));
141     memset(&inode, 0, sizeof(mtfs::inode_t));
142
143     original.accesRight = 0644;
144     original.uid = 1;
145     original.gid = 1;
146     original.size = 1024;
147     original.linkCount = 1;
148     original.atime = (unsigned long &&) time(nullptr);
149
150     for (int i = 0; i < 4; ++i) {

```

```

151     vector<mtfs::ident_t> blocks;
152
153     for (uint j = 0; j < 3; ++j) {
154         mtfs::ident_t ident = mtfs::ident_t(1, j, i * 3 + j);
155
156         blocks.push_back(ident);
157     }
158     original.dataBlocks.push_back(blocks);
159 }
160
161 uint64_t inodeId;
162 s3.add(&inodeId, mtfs::blockType::INODE);
163 s3.put(inodeId, &original, mtfs::blockType::INODE, false);
164 ASSERT_EQ(0, s3.get(inodeId, &inode, mtfs::blockType::INODE, false));
165
166 // cout << original.accesRight << " " << inode.accesRight << endl;
167 ASSERT_EQ(original, inode);
168 original.size = 2048;
169 s3.put(inodeId, &original, mtfs::blockType::INODE, false);
170 ASSERT_EQ(0, s3.get(inodeId, &inode, mtfs::blockType::INODE, false));
171 }
172
173 TEST_F(S3Fixture, addBlock) {
174     uint64_t block = 0;
175     ASSERT_EQ(0, s3.add(&block, mtfs::blockType::DATA_BLOCK));
176     // EXPECT_NE(0, block);
177
178     uint64_t block2 = 0;
179     ASSERT_EQ(0, s3.add(&block2, mtfs::blockType::DATA_BLOCK));
180     ASSERT_NE(0, block2);
181     ASSERT_GT(block2, block);
182
183     uint64_t block3 = 0;
184     ASSERT_EQ(0, s3.add(&block3, mtfs::blockType::DATA_BLOCK));
185 }
186
187 TEST_F(S3Fixture, delBlock) {
188     uint64_t block;
189     s3.add(&block, mtfs::blockType::DATA_BLOCK);
190     ASSERT_EQ(0, s3.del(block, mtfs::blockType::DATA_BLOCK));
191
192     uint64_t block2;
193     s3.add(&block2, mtfs::blockType::DATA_BLOCK);
194     s3.add(&block2, mtfs::blockType::DATA_BLOCK);
195     s3.add(&block2, mtfs::blockType::DATA_BLOCK);
196     s3.del(block, mtfs::blockType::DATA_BLOCK);
197 }
198
199 TEST_F(S3Fixture, writeBlock) {
200     uint8_t block[BLOCK_SIZE];
201     memset(block, 0, BLOCK_SIZE);
202     for (uint8_t i = 0; i < 50; ++i) {
203         block[i] = 'a';
204     }
205
206     uint64_t blockId = 0;
207     s3.add(&blockId, mtfs::blockType::DATA_BLOCK);
208     ASSERT_EQ(0, s3.put(blockId, &block, mtfs::blockType::DATA_BLOCK, false));
209 }

```

```

210 TEST_F(S3Fixture, readBlock) {
211     uint8_t block[BLOCK_SIZE];
212     memset(block, 0, BLOCK_SIZE);
213
214     uint8_t readBlock[BLOCK_SIZE];
215     memset(readBlock, 0, BLOCK_SIZE);
216
217     for (int i = 0; i < 500; ++i) {
218         block[i] = (uint8_t) to_string(i)[0];
219     }
220
221     uint64_t blockId = 0;
222     s3.add(&blockId, mtfs::blockType::DATA_BLOCK);
223     s3.put(blockId, &block, mtfs::blockType::DATA_BLOCK, false);
224
225     ASSERT_EQ(0, s3.get(blockId, &readBlock, mtfs::blockType::DATA_BLOCK, false));
226     ;
227     ASSERT_TRUE(0 == memcmp(block, readBlock, BLOCK_SIZE));
228 }
229
230 TEST_F(S3Fixture, rootInode) {
231     mtfs::inode_t inode;
232     inode.accessRight = 0444;
233     inode.uid = 1000;
234     inode.gid = 1000;
235     inode.size = 1024;
236     inode.linkCount = 2;
237     inode.atime = (uint64_t) time(nullptr);
238     inode.dataBlocks.clear();
239
240     mtfs::inode_t readInode;
241
242     ASSERT_EQ(0, s3.put(0, &inode, mtfs::blockType::INODE, false));
243     ASSERT_EQ(0, s3.get(0, &readInode, mtfs::blockType::INODE, false));
244     ASSERT_EQ(inode, readInode);
245 }
246
247 TEST_F(S3Fixture, superblock) {
248     mtfs::superblock_t superblock;
249     superblock.iCacheSz = superblock.dCacheSz = superblock.bCacheSz = superblock.
        blockSz = 4096;
250     superblock.migration = superblock.redundancy = 1;
251     superblock.pools.clear();
252     for (int i = 0; i < 5; ++i) {
253         mtfs::pool_t pool;
254         pool.migration = 0;
255         pool.rule = nullptr;
256         pool.volumes.clear();
257         superblock.pools[i] = pool;
258     }
259
260     ASSERT_TRUE(s3.putSuperblock(superblock));
261 }
262
263 TEST_F(S3Fixture, putMetas) {
264     mtfs::blockInfo_t blockInfo;
265
266

```

```

267 mtfs::ident_t ident1(1, 1, 1);
268 mtfs::ident_t ident2(2, 2, 2);
269
270 blockInfo.referenceId.push_back(ident1);
271 blockInfo.referenceId.push_back(ident2);
272 blockInfo.lastAccess = (uint64_t) time(nullptr);
273
274 ASSERT_EQ(0, s3.put(1, &blockInfo, mtfs::blockType::INODE, true));
275 ASSERT_EQ(0, s3.put(1, &blockInfo, mtfs::blockType::DIR_BLOCK, true));
276 ASSERT_EQ(0, s3.put(1, &blockInfo, mtfs::blockType::DATA_BLOCK, true));
277 }
278
279 //TEST_F(S3Fixture, getMetas) {
280 // mtfs::blockInfo_t blockInfo = mtfs::blockInfo_t(), receiveInfo = mtfs::
    blockInfo_t();
281 // memset(&blockInfo, 0, sizeof(mtfs::blockInfo_t));
282 // memset(&receiveInfo, 0, sizeof(mtfs::blockInfo_t));
283 //
284 // mtfs::ident_t ident1(1, 1, 1);
285 // mtfs::ident_t ident2(2, 2, 2);
286 //
287 // blockInfo.referenceId.push_back(ident1);
288 // blockInfo.referenceId.push_back(ident2);
289 // blockInfo.lastAccess = (uint64_t) time(nullptr);
290 //
291 // ASSERT_EQ(0, s3.put(2, &blockInfo, mtfs::blockType::INODE, true));
292 // ASSERT_EQ(0, s3.get(2, &receiveInfo, mtfs::blockType::INODE, true));
293 // ASSERT_EQ(blockInfo, receiveInfo);
294 //
295 // receiveInfo = mtfs::blockInfo_t();
296 // ASSERT_EQ(0, s3.put(2, &blockInfo, mtfs::blockType::DIR_BLOCK, true));
297 // ASSERT_EQ(0, s3.get(2, &receiveInfo, mtfs::blockType::DIR_BLOCK, true));
298 // ASSERT_EQ(blockInfo, receiveInfo);
299 //
300 // receiveInfo = mtfs::blockInfo_t();
301 // ASSERT_EQ(0, s3.put(2, &blockInfo, mtfs::blockType::DATA_BLOCK, true));
302 // ASSERT_EQ(0, s3.get(2, &receiveInfo, mtfs::blockType::DATA_BLOCK, true));
303 // ASSERT_EQ(blockInfo, receiveInfo);
304 //}
305
306 TEST_F(S3Fixture, putDirBlock) {
307 mtfs::ident_t id1, id2;
308 id2.poolId = id1.poolId = 1;
309 id1.volumeId = 1;
310 id1.id = 2;
311 id2.volumeId = 2;
312 id2.id = 42;
313
314 vector<mtfs::ident_t> ids;
315 ids.push_back(id1);
316 ids.push_back(id2);
317
318 mtfs::dirBlock_t block = mtfs::dirBlock_t();
319 block.entries.clear();
320 block.entries.insert(make_pair("baz", ids));
321 block.entries.insert(make_pair("test", ids));
322
323 uint64_t dId = 0;
324

```



```
325 ASSERT_EQ(0, s3.add(&dId, mtf::blockType::DIR_BLOCK));
326 EXPECT_LE(0, dId);
327 ASSERT_EQ(0, s3.put(dId, &block, mtf::blockType::DIR_BLOCK, false));
328
329 mtf::dirBlock_t readBlock;
330 ASSERT_EQ(0, s3.get(dId, &readBlock, mtf::blockType::DIR_BLOCK, false));
331 ASSERT_EQ(block.entries.size(), readBlock.entries.size());
332 for (auto &&item: block.entries) {
333     ASSERT_NE(readBlock.entries.end(), readBlock.entries.find(item.first));
334     int i = 0;
335     for (auto &&ident: item.second) {
336         ASSERT_EQ(ident, readBlock.entries[item.first][i]);
337         i++;
338     }
339 }
340
341 ASSERT_EQ(0, s3.del(dId, mtf::blockType::DIR_BLOCK));
342 }
343
344 int main(int argc, char **argv) {
345     ::testing::InitGoogleTest(&argc, argv);
346     return RUN_ALL_TESTS();
347 }
```

Listing 3.37 – "s3Tests.cpp"

Chapitre 4

CMakeLists

4.1 CMakeLists général

```
1 cmake_minimum_required(VERSION 3.5)
2 project(mtFS)
3
4 set(CMAKE_CXX_STANDARD 14)
5
6 include_directories(library)
7 include_directories(mtFS/include)
8
9 add_subdirectory(mtFS)
10 add_subdirectory(mtFS_Tests)
11 add_subdirectory(Plugin)
```

Listing 4.1 – "CMakeLists.txt"

4.2 CMakeLists mtfs

```
1 if (NOT DEFINED MTFS_HOME_DIR)
2     set(MTFS_HOME_DIR /home/mtfs)
3 endif ()
4 add_definitions(-DMTFS_HOME_DIR="${MTFS_HOME_DIR}")
5 add_definitions(-DMTFS_PLUGIN_HOME="${MTFS_HOME_DIR}/Plugins/")
6 add_definitions(-DMTFS_CONFIG_DIR="${MTFS_HOME_DIR}/Configs/")
7 add_definitions(-DMTFS_INSTALL_DIR="${MTFS_HOME_DIR}/Systems/")
8
9 set(FS_HEAD
10     include/mtfs/Cache.h
11     include/mtfs/Access.h
12     include/mtfs/Mtfs.h
13     include/mtfs/Pool.h
14     include/mtfs/PoolManager.h
15     include/mtfs/Rule.h
16     include/mtfs/structs.h
17     include/mtfs/TimeRule.h
18     include/mtfs/UserRightRule.h
19     include/mtfs/Volume.h
20     include/mtfs/Migrator.h)
21
22 set(FS_SRC
23     src/mtfs/Cache.cpp
```

```

24     src/mtfs/Mtfs.cpp
25     src/mtfs/Pool.cpp
26     src/mtfs/PoolManager.cpp
27     src/mtfs/Rule.cpp
28     src/mtfs/TimeRule.cpp
29     src/mtfs/UserRightRule.cpp
30     src/mtfs/Volume.cpp
31     src/mtfs/Migrator.cpp)
32
33 set(MTFUSE_HEAD
34     include/wrapper/FuseBase.h
35     include/wrapper/FuseCallback.h
36     include/wrapper/MtfsFuse.h
37 )
38
39 set(MTFUSE_SRC
40     src/wrapper/FuseBase.cpp
41     src/wrapper/FuseCallback.cpp
42     src/wrapper/MtfsFuse.cpp
43 )
44
45 set(CONF_HEAD
46     include/pluginSystem/PluginManager.h
47     include/pluginSystem/Plugin.h
48     src/utls/Semaphore.cpp include/utls/Semaphore.h src/mtfs/Migrator.cpp
49     include/mtfs/Migrator.h)
50
51 set(CONF_SRC
52     src/pluginSystem/PluginManager.cpp
53 )
54
55 set(UTILS
56     include/utls/Fs.h
57     src/utls/Fs.cpp
58     src/utls/Semaphore.cpp
59     include/utls/Semaphore.h
60     src/utls/Logger.cpp
61     include/utls/Logger.h
62 )
63
64 set(CMAKE_PREFIX_PATH /opt/aws-build)
65 find_package(aws-sdk-cpp)
66 add_definitions(-DUSE_IMPORT_EXPORT)
67
68 add_executable(mtfscCreate src/mtfsCreate.cpp
69     ${FS_SRC} ${CONF_SRC} ${UTILS})
70
71 target_link_libraries(mtfscCreate fuse3 dl boost_system boost_filesystem
72     boost_thread pthread)
73
74 #add_definitions(-DDEBUG)
75
76 add_executable(mtfsmount src/mtfsMount.cpp
77     ${CONF_HEAD} ${CONF_SRC}
78     ${FS_HEAD} ${FS_SRC}
79     ${MTFUSE_HEAD} ${MTFUSE_SRC}
80     ${UTILS}
81 )

```

```
81 target_link_libraries(mtfsMount fuse3 dl boost_system boost_filesystem
82 boost_thread pthread)
```

Listing 4.2 – "mtFS/CMakeLists.txt"

4.3 CMakeLists plugin

```
1 add_subdirectory(BlockDevice)
2 add_subdirectory(S3)
3
4 set(PLUGIN_H ../mtFS/include/pluginSystem/Plugin.h)
5
6 #set(BLOCK_SRCS BlockDevice/BlockDevice.h BlockDevice/BlockDevice.cpp ${
7   PLUGIN_H})
8 #add_executable(Block ${BLOCK_SRCS})
9 #
10 #set(S3_SRCS S3/S3.h S3/S3.cpp ${PLUGIN_H})
11 #add_executable(S3exe ${S3_SRCS})
12 #target_link_libraries(S3exe aws-cpp-sdk-s3)
```

Listing 4.3 – "Plugin/CMakeLists.txt"

4.3.1 CMakeLists block

```
1 project(blockDevice)
2
3 set(CMAKE_BUILD_TYPE Release)
4
5 add_library(block SHARED BlockDevice.cpp ../../mtFS/src/Utils/Logger.cpp)
```

Listing 4.4 – "Plugin/BlockDevice/CMakeLists.txt"

4.3.2 CMakeLists S3

```
1 project(S3)
2
3 set(CMAKE_BUILD_TYPE Release)
4
5 add_library(s3 SHARED S3.cpp ../../mtFS/src/Utils/Logger.cpp)
6 target_link_libraries(s3 aws-cpp-sdk-s3)
```

Listing 4.5 – "Plugin/S3/CMakeLists.txt"

4.4 CMakeLists Tests

```
1 project(mtFS_Tests)
2
3 add_subdirectory(PluginTests)
4 add_subdirectory(PerformanceTests)
```

Listing 4.6 – "mtFS_Tests/CMakeLists.txt"

4.4.1 CMakeLists Performance tests

```
1 add_executable(perfTests main.cpp)
2
3 target_link_libraries(perfTests boost_system boost_filesystem boost_thread)
```

Listing 4.7 – "mtFS_Tests/PerformanceTests/CMakeLists.txt"

4.4.2 CMakeLists PluginTests

```
1 include_directories(${gtest_SOURCE_DIR}/include ${gtest_SOURCE_DIR})
2 include_directories ../../mtFS/include ../../Plugin/
3
4 set(BLOCK_SRCS ../../Plugin/BlockDevice/BlockDevice.h ../../Plugin/BlockDevice/
   BlockDevice.cpp)
5
6 add_definitions(-DDEBUG)
7 add_executable(blockTests blockDeviceTests.cpp ${BLOCK_SRCS} ../../mtFS/src/
   utils/Logger.cpp)
8
9 target_link_libraries(blockTests gtest gtest_main boost_system)
10
11 set(CMAKE_PREFIX_PATH /opt/aws-build)
12
13 find_package(aws-sdk-cpp)
14 add_definitions(-DUSE_IMPORT_EXPORT)
15 set(S3_SRCS ../../Plugin/S3/S3.h ../../Plugin/S3/S3.cpp ../../mtFS/src/utils/
   Logger.cpp)
16 add_executable(s3Tests s3Tests.cpp ${S3_SRCS})
17
18 target_link_libraries(s3Tests gtest gtest_main boost_system aws-cpp-sdk-s3)
```

Listing 4.8 – "mtFS_Tests/PluginTests/CMakeLists.txt"

Bibliographie

- [1] CMake. Compilateur c/c++. <https://cmake.org/download/>, 2017. [En ligne ; Page disponible le 16-août-2017].